
OLinear: A Linear Model for Time Series Forecasting in Orthogonally Transformed Domain

**Wenzhen Yue¹, Yong Liu², Hao Wang³, Haoxuan Li^{4,*}, Xianghua Ying^{1,*},
Ruohao Guo¹, Bowei Xing¹, Ji Shi¹**

¹State Key Laboratory of General Artificial Intelligence, School of Intelligence Science and Technology, Peking University, ²School of Software, BNRist, Tsinghua University,

³College of Control Science and Engineering, Zhejiang University,

⁴Center for Data Science, Peking University

{yuewenzhen, hxli}@stu.pku.edu.cn, xhying@pku.edu.cn

Abstract

This paper presents **OLinear**, a linear-based multivariate time series forecasting model that operates in an orthogonally transformed domain. Recent forecasting models typically adopt the temporal forecast (TF) paradigm, which directly encode and decode time series in the time domain. However, the entangled step-wise dependencies in series data can hinder the performance of TF. To address this, some forecasters conduct encoding and decoding in the transformed domain using fixed, dataset-independent bases (e.g., sine and cosine signals in the Fourier transform). In contrast, we utilize **OrthoTrans**, a data-adaptive transformation based on an orthogonal matrix that diagonalizes the series' temporal Pearson correlation matrix. This approach enables more effective encoding and decoding in the decorrelated feature domain and can serve as a plug-in module to enhance existing forecasters. To enhance the representation learning for multivariate time series, we introduce a customized linear layer, **NormLin**, which employs a normalized weight matrix to capture multivariate dependencies. Empirically, the NormLin module shows a surprising performance advantage over multi-head self-attention, while requiring nearly half the FLOPs. Extensive experiments on 24 benchmarks and 140 forecasting tasks demonstrate that OLinear consistently achieves state-of-the-art performance with high efficiency. Notably, as a plug-in replacement for self-attention, the NormLin module consistently enhances Transformer-based forecasters. The code and datasets are available at <https://github.com/jackyue1994/OLinear>.

1 Introduction

Multivariate time series forecasting is critical in fields such as weather [Wu et al., 2023a], transportation [Ma et al., 2021], energy [Zhou et al., 2021], and finance [Chen et al., 2023]. Time series forecasters typically adopt the temporal forecast (TF) paradigm [Wang et al., 2025a, Liu et al., 2024a, Nie et al., 2023, Wang et al., 2024a], which encodes time series into latent representations and decodes them back, all within the time domain. However, this paradigm struggles to fully exploit the forecasting potential in the presence of entangled intra-series dependencies [Yi et al., 2023, Yue et al., 2025]. To mitigate this issue, recent studies apply Fourier [Yi et al., 2024a, 2023, Yue et al., 2025] or wavelet [Masserano et al., 2024] transforms to obtain the decorrelated feature sequence and perform encoding and decoding in the transformed domain. Nevertheless, these methods rely on dataset-independent bases, which fail to exploit the dataset-specific temporal correlation information.

*Haoxuan Li and Xianghua Ying are the corresponding authors.

In this paper, we introduce **OrthoTrans**, a dataset-adaptive transformation scheme that constructs an orthogonal basis via eigenvalue decomposition of the temporal Pearson correlation matrix [Gray and Davisson, 2004]. Projecting the series onto this basis obtains the decorrelated feature domain, providing a disentangled input for linear encoding and empirically improving forecasting performance. Notably, OrthoTrans is modular and can be integrated into existing forecasters to enhance their performance. In-depth ablation studies reveal that OrthoTrans promotes representation diversity and increases the rank of attention matrices in Transformer-based models.

As OrthoTrans transforms complex temporal variations into decorrelated features, the representation learning process can be effectively handled by linear layers [Yi et al., 2023]. Specifically, we employ a linear-based Cross-Series Learner (CSL) and Intra-Series Learner (ISL) to model multivariate correlations and sequential dynamics, respectively. To motivate our design, we note that in the classic self-attention mechanism, the attention entries are all positive with the row-wise L1 norm fixed as 1. Inspired by this, we design a new linear layer in CSL, called **NormLin**, where the weight matrix entries are made positive via the Softplus function and then row-wise **normalized**. Surprisingly, the NormLin module consistently outperforms the self-attention mechanism while improving computational efficiency in the field of time series forecasting. The overall model is referred to as **OLinear**, and our contributions can be summarized as follows:

- In contrast to the commonly used TF paradigm, we introduce **OrthoTrans**, a dataset-adaptive transformation scheme which leverages the orthogonal matrix derived from the eigenvalue decomposition of the temporal Pearson correlation matrix. As a plug-in, it consistently improves the performance of existing forecasters.
- For better representation learning, we present the **NormLin** layer, which employs a row-normalized weight matrix to capture multivariate correlations. Notably, as a plug-in, the NormLin module improves both the accuracy and efficiency of Transformer-based forecasters. It also adapts well to decoder architectures and large-scale time series models.
- Extensive experiments on 24 benchmarks and 140 forecasting tasks (covering various datasets and prediction settings) demonstrate that **OLinear** consistently achieves state-of-the-art performance with competitive computational efficiency.

2 Related work

2.1 Transformed domain in time series forecasting

Deep learning-based forecasters typically adopt the TF paradigm [Wang et al., 2025a, 2024a, Yu et al., 2024, Wang et al., 2024b, Liu et al., 2024a, Nie et al., 2023, Wu et al., 2023b, Zeng et al., 2023], with the entire process performed in the time domain. However, this paradigm may underperform in the presence of strong intra-series correlations [Yi et al., 2023, Wang et al., 2025b]. To mitigate this, recent models propose to forecast in the frequency domain [Yi et al., 2024a, Xu et al., 2024, Yi et al., 2023] or the wavelet domain [Masserano et al., 2024] where step-wise correlations are reduced and performance improvements are observed. FreDF [Wang et al., 2025b] incorporates a frequency-domain regularization term into the loss function to address the issue of biased forecast. However, classical Fourier and wavelet transforms employ fixed, dataset-agnostic bases that do not explicitly use the specific statistical characteristics of the dataset. In this work, we leverage orthogonal matrices derived from the temporal Pearson correlation matrix to decorrelate the series data. This approach provides a more suitable input for linear encoding and improves forecasting performance.

2.2 Time series forecasters

Deep learning based time series forecasters can generally be categorized into Transformer-based [Wang et al., 2025a, Nie et al., 2023, Liu et al., 2024a, Wang et al., 2024b, Yu et al., 2024], linear-based [Wang et al., 2024a, Zeng et al., 2023, Yi et al., 2023, Xu et al., 2024, Yi et al., 2024a], TCN-based [Wu et al., 2023b, Luo and Wang, 2024], RNN-based [Lai et al., 2018, Rangapuram et al., 2018], and GNN-based [Huang et al., 2023a, Yi et al., 2024b] models. Recently, research interests increasingly focus on Transformer-based and linear-based methods, each offering distinct advantages. Transformer-based forecasters typically exhibit strong expressiveness, whereas linear-based models offer better computational efficiency. In this paper, we aim to achieve state-of-the-art performance using an efficient linear-based model. Specifically, we employ the linear-based NormLin module to

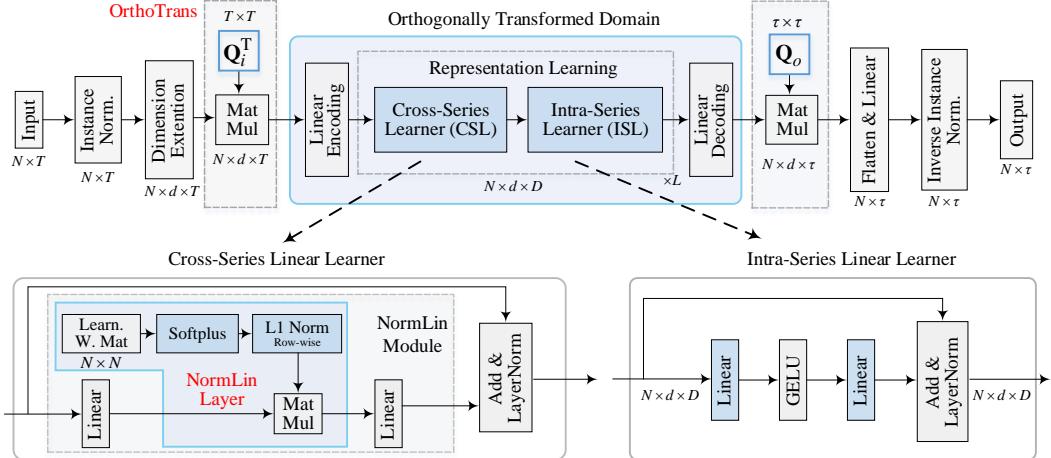


Figure 1: Overall structure of OLinear. The orthogonal matrix derived from the series’ temporal Pearson correlation matrix transforms the series into the new feature domain with removed sequential correlations. The cross-series learners and intra-series learners are then employed for robust representation learning. Specifically, we introduce the NormLin layer, whose weight matrix is processed by the Softplus function and row-wise L1 normalization, to capture multivariate correlations.

robustly model inter-variate correlations. Remarkably, the NormLin module consistently outperforms the self-attention mechanism in Transformer-based forecasters, highlighting the strong capability of linear layers for time series forecasting.

3 Preliminaries

Real-world time series often exhibit temporal dependencies, where each time step is influenced by its predecessors. In this section, we derive the expected value of a future step given its past, assuming a multivariate Gaussian distribution. Without loss of generality, we consider a series $\mathbf{x} \in \mathbb{R}^t$, and the following time step $y \in \mathbb{R}$. We now present the following theorem.

Theorem 1 (Expected value of y). *We assume that $\mathbf{x} \sim \mathcal{N}(\mu_{\mathbf{x}}, \Sigma_{\mathbf{x}})$, $y \sim \mathcal{N}(\mu_y, \sigma_y^2)$, and the joint distribution $[\mathbf{x}, y] \sim \mathcal{N}\left([\mu_{\mathbf{x}}, \mu_y], \begin{bmatrix} \Sigma_{\mathbf{x}} & \Sigma_{\mathbf{x}y} \\ \Sigma_{\mathbf{x}y}^T & \sigma_y^2 \end{bmatrix}\right)$. Here, $\mathcal{N}(\cdot, \cdot)$ denotes a Gaussian distribution with the specified mean vector and covariance matrix. $\Sigma_{\mathbf{x}} \in \mathbb{R}^{t \times t}$ is symmetric and positive definite. $\Sigma_{\mathbf{x}y} \in \mathbb{R}^{t \times 1}$ denotes the covariance between \mathbf{x} and y . Then, the expected value of y given \mathbf{x} is*

$$\mu_{y|\mathbf{x}} = \mu_y + \Sigma_{\mathbf{x}y}^T \Sigma_{\mathbf{x}}^{-1} (\mathbf{x} - \mu_{\mathbf{x}}) \quad (1)$$

The proof of Theorem 1 is provided in Appendix A. Equation 1 shows that for a temporally correlated series $[\mathbf{x}, y]$, the expected value of y depends not only on its own mean μ_y , but also on its past observations \mathbf{x} . The second term in Equation 1 introduces additional difficulty to forecasting under the TF paradigm. In contrast, decorrelating the series simplifies the forecasting task. In this work, we introduce OrthoTrans to transform the original series into a decorrelated transformed domain, converting the temporal forecasting problem into the *feature prediction* task. Experiments in Section 5.2 validate the effectiveness and generality of this transformation strategy.

4 Method

Problem formulation For multivariate time series forecasting, given a historical sequence $\mathbf{X} = \{\mathbf{x}_1, \dots, \mathbf{x}_T\} \in \mathbb{R}^{N \times T}$ with T time steps and N variates, the task is to predict the future τ time steps $\mathbf{Y} = \{\mathbf{x}_{T+1}, \dots, \mathbf{x}_{T+\tau}\} \in \mathbb{R}^{N \times \tau}$. Our goal is to approximate the ground truth \mathbf{Y} as closely as possible with predictions $\hat{\mathbf{Y}}$.

4.1 Overall architecture

As shown in Figure 1, OLinear adopts a simple yet effective architecture. Given the input series \mathbf{X} , a RevIN [Kim et al., 2021] layer first performs instance normalization to mitigate non-stationarity. A dimension extension module then enhances expressiveness [Yi et al., 2023] by computing the outer product with a learnable vector $\phi_d \in \mathbb{R}^d$, where d is the embedding size. Next, the time domain is decorrelated by multiplying with a transposed orthogonal matrix $\mathbf{Q}_i^\top \in \mathbb{R}^{T \times T}$, which will be detailed in Section 4.2. This process can be formulated as $\mathbf{Z} = (\text{RevIN}_{\text{Norm}}(\mathbf{X}) \otimes \phi_d) \mathbf{Q}_i^\top \in \mathbb{R}^{N \times d \times T}$, where \otimes denotes the outer product.

We then perform encoding and forecasting in the transformed domain. Specifically, a linear layer first encodes the decorrelated features \mathbf{Z} to the model dimension D . The cross-series learner (CSL) and intra-series learner (ISL) subsequently capture multivariate correlations and model intra-series dynamics, respectively. After passing through L stacked blocks, the representation is decoded to the desired prediction length τ , and then mapped back to the time domain via multiplication with the orthogonal matrix \mathbf{Q}_o . The overall process is summarized as:

$$\begin{aligned}\tilde{\mathbf{H}}^0 &= \text{LinearEncode}(\mathbf{Z}) \in \mathbb{R}^{N \times d \times D}, \\ \tilde{\mathbf{H}}^{l+1} &= \text{ISL}(\text{CSL}(\tilde{\mathbf{H}}^l)), \quad l = 0, \dots, L-1, \\ \tilde{\mathbf{Y}} &= \text{LinearDecode}(\tilde{\mathbf{H}}^L) \cdot \mathbf{Q}_o \in \mathbb{R}^{N \times d \times \tau}.\end{aligned}\quad (2)$$

Finally, the flattened output is mapped to shape $N \times \tau$ via a linear layer, then de-normalized to yield the final prediction: $\hat{\mathbf{Y}} = \text{RevIN}_{\text{DeNorm}}(\text{FlattenLinear}(\tilde{\mathbf{Y}}))$.

4.2 Orthogonal transformation (OrthoTrans)

One effective approach for decorrelating the series is based on the Pearson correlation matrix. Let $\mathbf{X}^{\text{train}} \in \mathbb{R}^{N \times M}$ denote the training set, where M is the length of the training series. For each variate j , we generate T lagged series with temporal offsets from 0 to $T - 1$: $s_i^j = \mathbf{X}[j, i : M - T + i]$, $i = 0, \dots, T - 1$, where s_i^j denotes the i -th temporally lagged series of variate j . We then compute the Pearson correlation matrix CorrMat_t^j of $\{s_i^j\}$, whose (p, q) -th entry is $\frac{\text{Cov}(s_p, s_q)}{\sqrt{\text{Var}(s_p) \cdot \text{Var}(s_q)}}$, $p, q \in \{0, \dots, T - 1\}$.

Here, $\text{Cov}(\cdot)$ denotes the covariance of two series, and $\text{Var}(\cdot)$ denotes the variance. The final Pearson correlation matrix is then obtained by averaging over all variates: $\text{CorrMat}_t = \frac{1}{N} \sum_{j=0}^{N-1} \text{CorrMat}_t^j$.

Following the above procedure, the resulting CorrMat_t is symmetric with all diagonal entries equal to 1. Based on the properties of symmetric matrices, we perform eigenvalue decomposition as $\text{CorrMat}_t = \mathbf{Q}_i \Lambda \mathbf{Q}_i^\top$, where \mathbf{Q}_i is an orthogonal matrix [Horn and Johnson, 2012] whose columns are the eigenvectors of CorrMat_t [Bishop and Nasrabadi, 2006], and Λ is a diagonal matrix containing the corresponding eigenvalues.

For the input series \mathbf{x} with temporal correlation matrix CorrMat_t and unit variance, the covariance matrix of the transformed vector $\mathbf{Q}_i^\top \mathbf{x}$ is $\text{CovMat}(\mathbf{Q}_i^\top \mathbf{x}) = \mathbf{Q}_i^\top \text{CorrMat}_t \mathbf{Q}_i = \mathbf{Q}_i^\top \mathbf{Q}_i \Lambda \mathbf{Q}_i^\top \mathbf{Q}_i = \Lambda$, which is a diagonal matrix. Therefore, the entries of $\mathbf{Q}_i^\top \mathbf{x}$ are linearly uncorrelated, removing sequential dependencies. Similarly, we can compute $\mathbf{Q}_o \in \mathbb{R}^{\tau \times \tau}$, and recover the temporal correlations

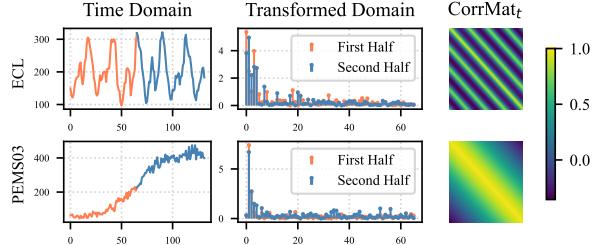


Figure 2: Comparison between the time and transformed domains. (1) The series exhibits strong consistency in the transformed domain, which is desirable for forecasting. (2) Temporal correlations are effectively reduced in the transformed domain. (3) Periodicity is clearly reflected in the correlation matrix, which facilitates modeling. (4) The new domain typically exhibits energy compaction, with only a few dominant components. More cases are shown in Figure 7.

by multiplying with \mathbf{Q}_o . Note that \mathbf{Q}_i and \mathbf{Q}_o are pre-computed and used throughout training and inference. Ablation studies on these two matrices are discussed in Appendix J.3. Moreover, as shown in Appendix I.10, OLinear remains robust even when the orthogonal matrices \mathbf{Q}_i and \mathbf{Q}_o are computed with limited training data.

Discussion (1) Mathematically, $\mathbf{Q}_i^T \mathbf{x}$ projects \mathbf{x} onto the eigenvectors of the temporal correlation matrix CorrMat_t . The transformed domain exhibits energy compaction, and **noise is suppressed** in the primary components [Bishop and Nasrabadi, 2006]. (2) Compared to DFT and wavelet transforms, or to the TF paradigm without any transformation, OrthoTrans produces **higher-rank attention matrices** for Transformer-based forecasters (see Figures 10–14), indicating greater representation diversity and enhanced model expressiveness. This partly explains why OrthoTrans enhances other forecasters when used as a plug-in module.

4.3 Representation learning

CSL and NormLin Inspired by the properties of attention matrices in self-attention [Vaswani et al., 2017], we impose two constraints on the weight matrix of the linear layer for robust multivariate representation learning: (1) all entries must be positive, and (2) each row must sum to 1. To enforce these constraints, we apply the **Softplus** function followed by row-wise L1 normalization to the learnable weight matrix \mathbf{W} . The resulting layer is referred to as NormLin, and is defined as follows:

$$\text{NormLin}(\mathbf{x}) = \text{RowNorm}_{\text{L1}}(\text{Softplus}(\mathbf{W})) \mathbf{x}, \mathbf{W} \in \mathbb{R}^{N \times N}, \mathbf{x} \in \mathbb{R}^N. \quad (3)$$

The variants of NormLin are discussed in Appendix J.1. Incorporating the other components in Figure 1(b), the CSL process is formulated as:

$$\text{CSL}(\tilde{\mathbf{H}}) = \text{LayerNorm}\left(\tilde{\mathbf{H}} + \text{Linear}\left(\text{NormLin}\left(\text{Linear}\left(\tilde{\mathbf{H}}\right)\right)\right)\right), \quad (4)$$

where $\text{Linear}(\cdot)$ and $\text{NormLin}(\cdot)$ operate on the sequence and variate dimensions, respectively. We define **the NormLin module** as $\text{NormLinModule}(\cdot) \triangleq \text{Linear}(\text{NormLin}(\text{Linear}(\cdot)))$. Ablation studies on the two linear layers in $\text{NormLinModule}(\cdot)$ are presented in Appendix J.2.

ISL We adopt two linear layers separated by the GELU activation function as a powerful predictive representation learner [Li et al., 2023]. It has been well established that MLPs are highly effective for encoding sequential dynamics and decoding future series [Zeng et al., 2023]. Similar to CSL, residual connections and LayerNorm are also applied. The complete ISL process is defined as:

$$\text{ISL}(\tilde{\mathbf{H}}) = \text{LayerNorm}\left(\tilde{\mathbf{H}} + \text{Linear}\left(\text{GELU}\left(\text{Linear}\left(\tilde{\mathbf{H}}\right)\right)\right)\right). \quad (5)$$

Discussion Compared to the self-attention mechanism, the NormLin module offers the following advantages: (1) As shown in Table 1, it reduces **computational complexity** by half and decreases memory footprint by a factor of $1/h$, where h is the number of attention heads; (2) The learned weight matrix in NormLin naturally exhibits **high rank**, in contrast to the low-rank nature of self-attention (see Appendix E.2). The low-rank issue in self-attention could arise from the sharp value concentration of the Softmax function. A higher-rank weight

(or attention) matrix often better preserves the rank of the representation space and thus improves the model’s expressiveness [Han et al., 2023]. (3) From the perspective of **gradient flow**, the NormLin layer provides a more direct backpropagation path for optimizing weight entries. Appendix B shows that the Jacobian matrices of $\text{Softmax}(\cdot)$ in self-attention and $\text{NormLin}(\text{Softplus}(\cdot))$ in the NormLin layer share a similar structure, but the latter offers greater flexibility. (4) The query-key mechanism of self-attention induces correlations in attention matrix entries, restricting the **optimization space**. In contrast, NormLin decouples these dependencies, allowing for more flexible parameter updates.

Table 1: Comparison of FLOPs and memory usage between the NormLin module and multi-head self-attention (MHSA). Here, h denotes the number of attention heads, typically set to 8.

Module	NormLin Module	MHSA
FLOPs	$\mathcal{O}(N^2D + 2ND^2)$	$\mathcal{O}(2N^2D + 4ND^2)$
Memory	$\mathcal{O}(N^2 + ND)$	$\mathcal{O}(hN^2 + ND)$

Table 2: Long-term forecasting results with prediction lengths $\tau \in \{12, 24, 48, 96\}$ for PEMS and $\tau \in \{96, 192, 336, 720\}$ for others. Lookback horizon $T = 96$. Results are averaged over four prediction lengths. Avg denotes further averaging over subsets. Full results are shown in Table 17.

Model	OLinear (Ours)	TimeMix. 2024a	FilterNet 2024a	DLinear 2023	TimeMix++. 2025a	Leddam 2024	CARD 2024b	Fredformer 2024	iTrans. 2024a	PatchTST 2023	TimesNet 2023b	
Metric	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE
ETT(Avg)	<u>0.359</u> 0.376	0.367	0.388	0.375	0.394	0.442	0.444	<u>0.349</u> <u>0.377</u>	0.367	0.387	0.366	0.385
ECL	<u>0.159</u> 0.248	0.182	0.273	0.173	0.268	0.212	0.300	<u>0.165</u> <u>0.253</u>	0.169	0.263	0.168	0.258
Exchange	0.355	<u>0.399</u>	0.387	0.416	0.388	0.419	<u>0.354</u>	0.414	0.357	0.409	<u>0.354</u>	0.402
Traffic	0.451	0.247	0.485	0.298	0.463	0.310	0.625	0.383	0.416 <u>0.264</u>	0.467	0.294	0.453
Weather	<u>0.237</u> 0.260	0.240	0.272	0.245	0.272	0.265	0.317	<u>0.226</u> <u>0.262</u>	0.242	0.272	0.239	0.265
Solar	<u>0.215</u> 0.217	0.216	0.280	0.235	0.266	0.330	0.401	0.203	0.258	0.230	0.264	0.237
PEMS03	<u>0.095</u> 0.199	0.167	0.267	0.145	0.251	0.278	0.375	0.165	0.263	<u>0.107</u> <u>0.210</u>	0.174	0.275
PEMS04	0.091 <u>0.190</u>	0.185	0.287	0.146	0.258	0.295	0.388	0.136	0.251	<u>0.103</u> <u>0.210</u>	0.206	0.299
PEMS07	0.077 <u>0.164</u>	0.181	0.271	0.123	0.229	0.329	0.395	0.152	0.258	<u>0.084</u> <u>0.180</u>	0.149	0.247
PEMS08	0.113 <u>0.194</u>	0.226	0.299	0.172	0.260	0.379	0.416	0.200	0.279	<u>0.122</u> <u>0.211</u>	0.201	0.280
1 st Count	5	9	0	0	0	0	0	4	0	0	0	0

Table 3: Short-term forecasting results. Two settings are applied: S1 (Input-12, Predict-{3, 6, 9, 12}) and S2 (Input-36, Predict-{24, 36, 48, 60}). Average results across eight prediction lengths are reported. Full results are provided in Tables 18 and 19.

Model	OLinear (Ours)	TimeMix. 2024a	FilterNet 2024a	DLinear 2023	TimeMix++. 2025a	Leddam 2024	CARD 2024b	Fredformer 2024	iTrans. 2024a	PatchTST 2023	TimesNet 2023b	
Metric	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE
ILI	1.429 <u>0.690</u>	1.864	0.806	1.793	0.791	2.742	1.126	1.805	0.793	1.725	0.777	1.959
COVID-19	5.187 <u>1.211</u>	5.919	1.350	5.607	1.322	8.279	1.601	5.974	1.369	<u>5.251</u> <u>1.285</u>	5.536	1.314
METR-LA	0.587	0.311	0.608	0.372	0.603	0.366	<u>0.580</u>	0.422	0.567	0.363	0.603	0.367
NASDAQ	<u>0.121</u> 0.201 <u>0.120</u>	<u>0.204</u>	0.127	0.211	0.150	0.251	0.125	0.210	0.128	0.212	0.127	0.210
Wiki	<u>6.395</u> 0.415	6.443	0.439	6.457	0.439	6.420	0.510	6.430	0.443	6.417	0.433	6.419
SP500	0.146 <u>0.250</u>	0.153	0.265	0.164	0.279	0.178	0.298	0.157	0.270	0.163	0.282	<u>0.147</u> <u>0.252</u>
DowJones	7.686 <u>0.619</u>	8.499	0.633	8.283	0.633	7.893	0.626	8.895	0.643	8.257	0.633	<u>7.699</u> 0.619
CarSales	0.330	0.305	0.333	0.322	0.328	0.319	0.387	0.376	0.337	0.321	0.335	0.322
Power	<u>1.248</u> 0.835 <u>1.234</u> <u>0.840</u>	1.309	0.870	1.278	0.870	1.234	0.841	1.295	0.868	1.288	0.847	1.302
Website	<u>0.225</u> 0.311	0.279	0.358	0.297	0.367	0.302	0.389	0.260	0.344	0.264	0.348	0.303
Unemp	0.729 <u>0.461</u>	1.581	0.708	1.286	0.627	<u>0.565</u> <u>0.509</u>	1.506	0.678	1.502	0.689	1.163	0.596
1 st Count	4	10	2	0	0	0	1	0	2	0	0	0

5 Experiments

Datasets and implementation details We extensively evaluate OLinear using 24 diverse real-world datasets: ETT (four subsets), Weather, ECL, Traffic, Exchange, Solar-Energy, PEMS (four subsets), ILI, COVID-19, METR-LA, NASDAQ, Wiki, SP500, DowJones, CarSales, Power, Website, Unemp. The weighted L1 loss function from CARD [Wang et al., 2024b] is adopted. The embedding size d is set as 16. Dataset description and more implementation details are presented in Appendices C and D, respectively.

5.1 Forecasting performance

Baselines We carefully choose 11 well-acknowledged state-of-the-art forecasting models as our baselines, including (1) Linear-based models: TimeMixer [Wang et al., 2024a], FilterNet [Yi et al., 2024a], FITS [Xu et al., 2024], DLinear [Zeng et al., 2023]; (2) Transformer-based models: TimeMixer++ [Wang et al., 2025a], Leddam [Yu et al., 2024], CARD [Wang et al., 2024b], Fredformer [Piao et al., 2024], iTransformer [Liu et al., 2024a], PatchTST [Nie et al., 2023]; (3) TCN-based model: TimesNet [Wu et al., 2023b].

Main results Comprehensive long-term and short-term forecasting results are presented in Tables 2 and 3, respectively, with the best results highlighted in **bold** and the second-best underlined. Lower MSE/MAE values indicate more accurate predictions. Across a wide range of benchmarks, OLinear

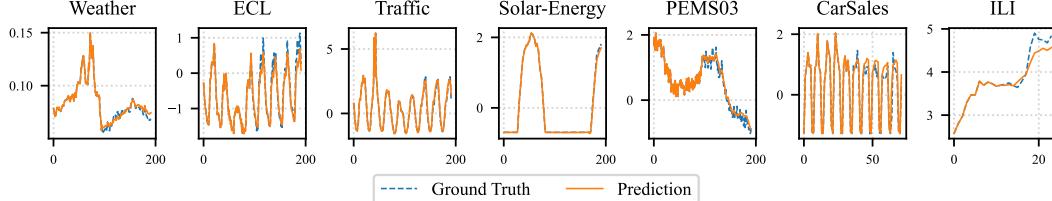


Figure 4: Visualization of the forecasting results of OLinear, demonstrating great accuracy.

Table 4: Ablation study on transformation bases. $S1$ and $S2$ represent ‘Input-12, Predict-{3, 6, 9, 12}’ and ‘Input-36, Predict-{24, 36, 48, 60}’, respectively. Wavelet1 and Wavelet2 use the Haar and discrete Meyer wavelets. Identity denotes no transformation. Full results are in Table 20.

Dataset	ECL		Solar-Energy		PEMS03		Power (S2)		ILI (S1)		COVID (S2)		METR-LA (S2)	
	Metric	MSE	MAE	MSE										
Ours	0.159	0.248	0.215	0.217	0.095	0.199	1.487	0.922	1.094	0.578	8.467	1.754	0.838	0.402
Fourier	0.161	0.250	<u>0.219</u>	<u>0.219</u>	<u>0.101</u>	<u>0.204</u>	1.614	0.967	1.268	0.584	9.165	1.839	0.843	<u>0.403</u>
Wavelet1	<u>0.160</u>	<u>0.249</u>	0.221	0.221	0.107	0.210	1.663	0.987	<u>0.116</u>	<u>0.580</u>	8.666	1.799	<u>0.840</u>	0.404
Wavelet2	0.162	0.251	0.226	0.224	0.108	0.210	1.664	0.987	1.177	0.594	8.949	1.840	0.843	0.406
Chebyshev	0.218	0.295	0.226	0.226	0.105	0.207	1.570	0.965	1.217	0.597	9.330	1.875	0.854	0.407
Laguerre	0.167	0.255	0.233	0.230	0.111	0.214	1.659	0.984	1.353	0.651	9.302	1.890	0.868	0.420
Legendre	0.161	0.250	0.243	0.235	0.109	0.213	1.685	0.995	1.177	0.603	<u>8.550</u>	<u>1.798</u>	0.841	0.404
Identity	0.163	0.252	0.227	0.225	0.106	0.209	<u>1.542</u>	<u>0.945</u>	1.153	0.587	8.856	1.819	0.848	0.408

consistently outperforms state-of-the-art Transformer-based and linear-based forecasters. Notably, these gains are achieved with high computational efficiency. (Figure 3 and Table 41).

We attribute this superior performance to the adopted OrthoTrans and NormLin modules. The effectiveness of the simpler NormLin module challenges the necessity of the widely adopted multi-head self-attention mechanism, which has been a dominant design in prior works. To further validate robustness, we evaluate OLinear under varying lookback lengths (Table 29), where it consistently outperforms existing state-of-the-art methods.

Figure 4 shows the prediction visualizations of OLinear. Moreover, Table 14 demonstrates that OLinear exhibits greater robustness to random seeds compared to state-of-the-art Transformer-based forecasters such as TimeMixer++ and iTransformer.

5.2 Model analysis

Transformation bases We replace OrthoTrans with several commonly used bases, including the Fourier basis, two wavelet bases (Haar and discrete Meyer wavelets), and three polynomial bases. As shown in Table 4, our method consistently outperforms all alternatives. Specifically, it achieves a 5.5% reduction in MSE over the Fourier basis, and a 10.0% improvement over the no-transformation baseline (i.e., the TF paradigm) on PEMS03.

OrthoTrans as a plug-in We further integrate OrthoTrans into three classic forecasters: iTrans-

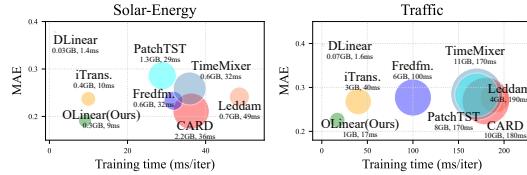


Figure 3: Model efficiency comparison. Bubble areas represent GPU memory usage during training, scaled independently in the two subfigures. Resource footprint data are from Table 41.

Table 5: Applying OrthoTrans (O.Trans) to iTransformer, PatchTST and RLinear. Average MSEs are reported. Full results are in Table 21.

Model	iTrans. 2024a		PatchTST 2023		RLinear 2023	
	Van.	+O.Trans	Van.	+O.Trans	Van.	+O.Trans
ETTm1	0.407	0.404	0.387	0.384	0.414	0.408
ECL	0.178	0.171	0.208	0.181	0.219	0.214
PEMS03	0.113	0.103	0.180	0.163	0.495	0.477
PEMS07	0.101	0.085	0.211	0.147	0.504	0.485
Solar	0.233	0.228	0.270	0.239	0.369	0.354
Weather	0.258	0.252	0.259	0.246	0.272	0.269
METR-LA	0.338	0.329	0.335	0.333	0.342	0.341

former, PatchTST, and RLinear [Li et al., 2023]. As shown in Table 5, OrthoTrans yields average MSE improvements of 5.1% and 10.1% for iTransformer and PatchTST, respectively, highlighting the benefit of incorporating dataset-specific statistical information into the model design. This improvement can be attributed to the increased attention matrix rank introduced by OrthoTrans (see Appendix E.3), which indicates enlarged representation space and enhanced model capacity.

Representation learning To validate the rationality of our CSL and ISL designs, we conduct ablation studies by replacing or removing their core components—NormLin and (standard) linear layers. As shown in Table 6, our design—applying NormLin along the variate dimension and standard linear layers along the temporal dimension—consistently achieves the best performance. Notably, applying NormLin along the *variate* dimension consistently outperforms its **self-attention** counterpart (last row), with reduced computational cost. Furthermore, removing the ISL module (third-last row) results in a 6.2% performance decline, highlighting the importance of updating temporal representations. Interestingly, on small-scale datasets with fewer variates (e.g., NASDAQ and ILI), the model with only temporal linear layers (third row) exhibits competitive performance, implying that NormLin is more beneficial when handling a larger number of variates.

Table 6: Ablations on the CSL and ISL design. ‘NormLin (Temporal)’ replaces the linear layer in Eq. 5 with the NormLin layer, and ‘Linear (Variate)’ replaces the NormLin layer in Eq. 4 with a standard linear layer. ‘Attn.’ denotes the self-attention mechanism. Results are averaged over four prediction lengths. Full results are shown in Table 22.

Variate	Temp.	ECL		Traffic		Solar		PEMS03		Weather		ETTm1		NASDAQ (S1)		ILI (S2)	
		MSE	MAE														
NormLin	Linear	0.159	0.248	0.451	0.247	0.215	0.217	0.095	0.199	0.237	0.260	0.374	0.377	0.055	0.125	1.764	0.802
Linear	Linear	0.178	0.272	0.606	0.320	0.246	0.238	0.121	0.226	0.238	0.261	0.377	0.380	0.057	0.132	1.938	0.837
	w/o	0.178	0.259	0.482	0.257	0.241	0.232	0.147	0.234	0.247	0.266	0.378	0.379	0.054	0.124	1.864	0.823
NormLin	NormLin	0.169	0.257	0.460	0.275	0.252	0.240	0.112	0.214	0.239	0.261	0.381	0.382	0.055	0.126	1.947	0.836
Linear	NormLin	0.183	0.276	0.578	0.339	0.262	0.254	0.143	0.246	0.240	0.261	0.383	0.384	0.057	0.133	2.037	0.867
	w/o	0.185	0.266	0.493	0.290	0.283	0.262	0.182	0.269	0.246	0.265	0.384	0.385	0.055	0.125	2.093	0.874
NormLin	w/o	0.169	0.257	0.460	0.275	0.253	0.241	0.114	0.215	0.239	0.261	0.380	0.382	0.055	0.126	1.940	0.837
Linear	w/o	0.183	0.276	0.591	0.341	0.262	0.254	0.142	0.246	0.240	0.262	0.384	0.384	0.057	0.132	2.073	0.874
Attn.	Linear	0.166	0.255	0.457	0.251	0.220	0.221	0.097	0.202	0.244	0.265	0.391	0.389	0.056	0.126	2.022	0.847

OLinear-C As shown in Figure 5, the learned weight matrix resemble Softmax (CorrMat_v), where CorrMat_v is the Pearson correlation matrix across variates. Motivated by this, we replace the learnable weights in NormLin with the pre-computed Softmax (CorrMat_v), resulting in a simplified variant: $\text{NormLin}_c(\mathbf{x}) = \text{Softmax}(\text{CorrMat}_v) \mathbf{x}$. **This indicates that the weight matrix in NormLin is optimized to align with the correlation matrix among channels CorrMat_v .** We refer to the model with this NormLin variant as **OLinear-C**. Table 7 shows that OLinear-C performs comparably to OLinear with reduced learnable parameters. Comprehensive evaluation of OLinear-C and the variants of NormLin_c are discussed in Appendix H.

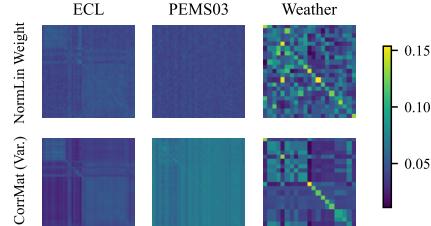


Figure 5: The learned NormLin weights exhibit similarity to the multivariate correlation matrix (after Softmax).

5.3 Generality and scalability of the NormLin module

In our CSL module, we employ the NormLin module—comprising the NormLin layer and its associated pre- and post-linear layers (Equation 4)—to capture multivariate correlations. Despite its simple architecture, the NormLin module demonstrates strong capability, generality, and scalability. Specifically, it consistently outperforms the classic multi-head self-attention mechanism and its variants, offering a compelling alternative for token dependency modeling in time series forecasting.

Comparison with attention mechanisms To assess its effectiveness, we compare NormLin with classic attention variants, such as Reformer [Kitaev et al., 2020], Flowformer [Wu et al., 2022], FLatten [Dao et al., 2022], and Mamba [Gu and Dao, 2023]. As shown in Table 8, NormLin

Table 7: Performance comparison of OLinear and OLinear-C. Full results are shown in Table 25.

Dataset	ECL		Traffic		ETT		Solar		PEMS		CarSales		ILI		COVID-19		Unemp	
Metric	MSE	MAE																
OLinear	0.159	0.248	0.451	0.247	0.359	0.376	0.215	0.217	0.094	0.187	0.330	0.305	1.429	0.690	5.187	1.211	0.729	0.461
OLinear-C	0.161	0.249	0.451	0.247	0.359	0.376	0.215	0.217	0.094	0.187	0.330	0.305	1.463	0.698	5.346	1.247	0.766	0.474

Table 8: Comparison of the NormLin module with state-of-the-art attention variants. Average MSEs are reported. Full results are in Table 23.

Dataset	NormLin	Trans.	Refm.	Flowfm.	Flatten	Mamba	E.Attn.	(Ours)	2017	2020	2022	2023	2023	2025
ECL	0.159	0.166	0.167	0.165	0.164	0.176	0.159							
Traffic	0.451	0.457	0.459	0.460	0.464	0.456	0.439							
PEMS03	0.095	0.097	0.096	0.099	0.101	0.104	0.097							
Weather	0.237	0.244	0.241	0.242	0.246	0.242	0.241							
Solar	0.215	0.223	0.216	0.222	0.231	0.228	0.217							
ILI	1.764	2.022	1.821	1.881	2.134	1.950	1.878							
NASDAQ	0.055	0.056	0.055	0.055	0.057	0.055	0.055							

Table 9: Applying NormLin (N.Lin) to state-of-the-art Transformer-based forecasters. Average MSEs are reported. Full results are in Table 24.

Dataset	iTrans.		PatchTST		Leddam		Fredformer	
	Van.	N.Lin	Van.	N.Lin	Van.	N.Lin	Van.	N.Lin
ETTm1	0.407	0.388	0.387	0.379	0.386	0.381	0.384	0.381
ECL	0.178	0.166	0.208	0.181	0.169	0.165	0.176	0.169
PEMS03	0.113	0.102	0.180	0.146	0.107	0.103	0.134	0.108
PEMS07	0.101	0.086	0.211	0.168	0.084	0.082	0.121	0.096
Solar	0.233	0.226	0.270	0.237	0.230	0.222	0.226	0.226
Weather	0.258	0.245	0.259	0.245	0.242	0.242	0.246	0.240
METR-LA	0.338	0.328	0.335	0.333	0.327	0.320	0.336	0.329

consistently outperforms these methods, indicating that simply normalizing the weight matrices can be more effective than self-attention for time series forecasting (see Paragraph 4.3).

NormLin as a plug-in To demonstrate the generality of the NormLin module, we replace self-attention with NormLin in Transformer-based forecasters. As shown in Table 9, this substitution leads to notable MSE improvements—6.7% for iTransformer and 10.3% for PatchTST—validating its plug-and-play effectiveness. These results also validate NormLin’s capability to model dependencies across multiple token types (e.g., variate, temporal patch, and frequency-domain patch), highlighting its potential as a universal token dependency learner for time series forecasting. Furthermore, the NormLin module consistently improves both training and inference efficiency across these models (see Table 42); for example, it boosts iTransformer’s inference efficiency by an average of 53%.

Scalability of the NormLin module Decoder-only Transformers are currently the *de facto* architecture of large time series models [Liu et al., 2024b, Ansari et al., 2024, Das et al., 2024]. We take Timer [Liu et al., 2024b] as a representative model and replace its self-attention with our NormLin module. To align with Timer’s decoder-only structure, we apply a causal mask by zeroing out the upper triangular part of NormLin’s weight matrix. The modified model is pre-trained on the UTSD dataset [Liu et al., 2024b], which spans seven domains and contains up to 1 billion time points. As shown in Table 10, NormLin improves performance in both zero-shot and fine-tuning scenarios. For example, the zero-shot MSE on ETTh2 and ETTm1 is reduced by 12.7% and 8.4%, respectively. These results demonstrate that the NormLin module adapts well to decoder-only architectures and scales effectively to large-scale pre-training scenarios.

Table 10: Applying NormLin to the large time series model Timer. For fine-tuning, 5% samples of trainset are used. MSEs are reported.

Dataset	Fine-tuning		Zero-shot	
	Van.	NormLin	Van.	NormLin
ETTh1	0.362	0.360	0.438	0.404
ETTh2	0.280	0.269	0.314	0.274
ETTm1	0.321	0.309	0.690	0.632
ETTm2	0.176	0.172	0.213	0.212
ECL	0.132	0.130	0.192	0.183
Traffic	0.361	0.353	0.458	0.462
Weather	0.151	0.149	0.181	0.174

6 Conclusion

In this work, we present OLinear, a simple yet effective linear-based forecaster that achieves state-of-the-art performance, built on two core components: (1) OrthoTrans, an orthogonal transformation that decorrelates temporal dependencies to facilitate better encoding and forecasting, and (2) the NormLin module, a powerful and general-purpose token dependency learner. Notably, both modules consistently improve existing forecasters when used as plug-ins. We hope OLinear will serve as a strong baseline and inspire future studies in time series forecasting.

Acknowledgments

This work was supported by Beijing Natural Science Foundation under Grant No. L247029, and the National Natural Science Foundation of China (NSFC) under Grant No. 62371009 and Grant No. 623B2002.

References

- Haixu Wu, Hang Zhou, Mingsheng Long, and Jianmin Wang. Interpretable weather forecasting for worldwide stations with a unified deep model. *Nature Machine Intelligence*, 5(6):602–611, 2023a.
- Changxi Ma, Guowen Dai, and Jibiao Zhou. Short-term traffic flow prediction for urban road sections based on time series analysis and lstm_bilstm method. *IEEE Transactions on Intelligent Transportation Systems*, 23(6):5615–5624, 2021.
- Haoyi Zhou, Shanghang Zhang, Jieqi Peng, Shuai Zhang, Jianxin Li, Hui Xiong, and Wancai Zhang. Informer: Beyond efficient transformer for long sequence time-series forecasting. In *AAAI*, volume 35, pages 11106–11115, 2021.
- Zonglei Chen, Minbo Ma, Tianrui Li, Hongjun Wang, and Chongshou Li. Long sequence time-series forecasting with deep learning: A survey. *Information Fusion*, 97:101819, 2023.
- Shiyu Wang, Jiawei Li, Xiaoming Shi, Zhou Ye, Baichuan Mo, Wenze Lin, Shengtong Ju, Zhixuan Chu, and Ming Jin. Timemixer++: A general time series pattern machine for universal predictive analysis. In *ICLR*, 2025a.
- Yong Liu, Tengge Hu, Haoran Zhang, Haixu Wu, Shiyu Wang, Lintao Ma, and Mingsheng Long. itransformer: Inverted transformers are effective for time series forecasting. In *ICLR*, 2024a.
- Yuqi Nie, Nam H. Nguyen, Phanwadee Sinthong, and Jayant Kalagnanam. A time series is worth 64 words: Long-term forecasting with transformers. *ICLR*, 2023.
- Shiyu Wang, Haixu Wu, Xiaoming Shi, Tengge Hu, Huakun Luo, Lintao Ma, James Y. Zhang, and Jun Zhou. Timemixer: Decomposable multiscale mixing for time series forecasting. In *ICLR*, 2024a.
- Kun Yi, Qi Zhang, Wei Fan, Shoujin Wang, Pengyang Wang, Hui He, Ning An, Defu Lian, Longbing Cao, and Zhendong Niu. Frequency-domain mlps are more effective learners in time series forecasting. In *NeurIPS*, 2023.
- Wenzhen Yue, Yong Liu, Xianghua Ying, Bowei Xing, Ruohao Guo, and Ji Shi. Freeformer: Frequency enhanced transformer for multivariate time series forecasting. *arXiv preprint arXiv:2501.13989*, 2025.
- Kun Yi, Jingru Fei, Qi Zhang, Hui He, Shufeng Hao, Defu Lian, and Wei Fan. Filternet: Harnessing frequency filters for time series forecasting. *Advances in Neural Information Processing Systems*, 37:55115–55140, 2024a.
- Luca Masserano, Abdul Fatir Ansari, Boran Han, Xiyuan Zhang, Christos Faloutsos, Michael W Mahoney, Andrew Gordon Wilson, Youngsuk Park, Syama Rangapuram, Danielle C Maddix, et al. Enhancing foundation models for time series forecasting via wavelet-based tokenization. *arXiv preprint arXiv:2412.05244*, 2024.
- Robert M Gray and Lee D Davisson. *An introduction to statistical signal processing*. Cambridge University Press, 2004.
- Guoqi Yu, Jing Zou, Xiaowei Hu, Angelica I Aviles-Rivero, Jing Qin, and Shujun Wang. Revitalizing multivariate time series forecasting: Learnable decomposition with inter-series dependencies and intra-series variations modeling. In *ICML*, 2024.
- Xue Wang, Tian Zhou, Qingsong Wen, Jinyang Gao, Bolin Ding, and Rong Jin. Card: Channel aligned robust blend transformer for time series forecasting. In *ICLR*, 2024b.

- Haixu Wu, Tengge Hu, Yong Liu, Hang Zhou, Jianmin Wang, and Mingsheng Long. Timesnet: Temporal 2d-variation modeling for general time series analysis. In *ICLR*, 2023b.
- Ailing Zeng, Muxi Chen, Lei Zhang, and Qiang Xu. Are transformers effective for time series forecasting? In *AAAI*, volume 37, pages 11121–11128, 2023.
- Hao Wang, Lichen Pan, Yuan Shen, Zhichao Chen, Degui Yang, Yifei Yang, Sen Zhang, Xinggao Liu, Haoxuan Li, and Dacheng Tao. Fredf: Learning to forecast in the frequency domain. In *ICLR*, 2025b.
- Zhijian Xu, Ailing Zeng, and Qiang Xu. Fits: Modeling time series with $10k$ parameters. In *ICLR*, 2024.
- Donghao Luo and Xue Wang. Moderntcn: A modern pure convolution structure for general time series analysis. In *ICLR*, 2024.
- Guokun Lai, Wei-Cheng Chang, Yiming Yang, and Hanxiao Liu. Modeling long-and short-term temporal patterns with deep neural networks. In *SIGIR*, pages 95–104, 2018.
- Syama Sundar Rangapuram, Matthias W Seeger, Jan Gasthaus, Lorenzo Stella, Yuyang Wang, and Tim Januschowski. Deep state space models for time series forecasting. *NeurIPS*, 31, 2018.
- Qihe Huang, Lei Shen, Ruixin Zhang, Shouhong Ding, Binwu Wang, Zhengyang Zhou, and Yang Wang. Crossgnn: Confronting noisy multivariate time series via cross interaction refinement. *NeurIPS*, 36:46885–46902, 2023a.
- Kun Yi, Qi Zhang, Wei Fan, Hui He, Liang Hu, Pengyang Wang, Ning An, Longbing Cao, and Zhendong Niu. Fouriergnn: Rethinking multivariate time series forecasting from a pure graph perspective. *Advances in Neural Information Processing Systems*, 36, 2024b.
- Taesung Kim, Jinhee Kim, Yunwon Tae, Cheonbok Park, Jang-Ho Choi, and Jaegul Choo. Reversible instance normalization for accurate time-series forecasting against distribution shift. In *ICLR*, 2021.
- Roger A Horn and Charles R Johnson. *Matrix analysis*. Cambridge university press, 2012.
- Christopher M Bishop and Nasser M Nasrabadi. *Pattern recognition and machine learning*, volume 4. Springer, 2006.
- Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Łukasz Kaiser, and Illia Polosukhin. Attention is all you need. *NIPS*, 30, 2017.
- Zhe Li, Shiyi Qi, Yiduo Li, and Zenglin Xu. Revisiting long-term time series forecasting: An investigation on linear mapping. *arXiv preprint arXiv:2305.10721*, 2023.
- Dongchen Han, Xuran Pan, Yizeng Han, Shiji Song, and Gao Huang. Flatten transformer: Vision transformer using focused linear attention. In *ICCV*, pages 5961–5971, 2023.
- Xihao Piao, Zheng Chen, Taichi Murayama, Yasuko Matsubara, and Yasushi Sakurai. Fredformer: Frequency debiased transformer for time series forecasting. In *SIGKDD*, 2024.
- Nikita Kitaev, Łukasz Kaiser, and Anselm Levskaya. Reformer: The efficient transformer. *ICLR*, 2020.
- Haixu Wu, Jialong Wu, Jiehui Xu, Jianmin Wang, and Mingsheng Long. Flowformer: Linearizing transformers with conservation flows. In *ICML*, 2022.
- Albert Gu and Tri Dao. Mamba: Linear-time sequence modeling with selective state spaces. *arXiv preprint arXiv:2312.00752*, 2023.
- Tri Dao, Dan Fu, Stefano Ermon, Atri Rudra, and Christopher Ré. Flashattention: Fast and memory-efficient exact attention with io-awareness. *NeurIPS*, 35:16344–16359, 2022.
- Yong Liu, Haoran Zhang, Chenyu Li, Xiangdong Huang, Jianmin Wang, and Mingsheng Long. Timer: Transformers for time series analysis at scale. In *ICML*, 2024b.

- Abdul Fatir Ansari, Lorenzo Stella, Caner Turkmen, Xiyuan Zhang, Pedro Mercado, Huibin Shen, Oleksandr Shchur, Syama Sundar Rangapuram, Sebastian Pineda Arango, Shubham Kapoor, et al. Chronos: Learning the language of time series. *arXiv preprint arXiv:2403.07815*, 2024.
- Abhimanyu Das, Weihao Kong, Rajat Sen, and Yichen Zhou. A decoder-only foundation model for time-series forecasting. In *Forty-first International Conference on Machine Learning*, 2024.
- Hao Wang, Zhiyu Wang, Yunlong Niu, Zhaoran Liu, Haozhe Li, Yilin Liao, Yuxin Huang, and Xinggao Liu. An accurate and interpretable framework for trustworthy process monitoring. *IEEE Transactions on Artificial Intelligence*, 5(5):2241–2252, 2023.
- Haotian Wang, Haoxuan Li, Hao Zou, Haoang Chi, Long Lan, Wanrong Huang, and Wenjing Yang. Effective and efficient time-varying counterfactual prediction with state-space models. In *The Thirteenth International Conference on Learning Representations*, 2025c.
- Lei Wang, Shanshan Huang, Chunyuan Zheng, Jun Liao, Xiaofei Zhu, Haoxuan Li, and Li Liu. Mitigating data imbalance in time series classification based on counterfactual minority samples augmentation. In *Proceedings of the 31st ACM SIGKDD Conference on Knowledge Discovery and Data Mining* V. 2, pages 2962–2973, 2025d.
- Hao Wang, Haoxuan Li, Xu Chen, Mingming Gong, Zhichao Chen, et al. Optimal transport for time series imputation. In *The Thirteenth International Conference on Learning Representations*, 2025e.
- Hao Wang, Licheng Pan, Zhichao Chen, Xu Chen, Qingyang Dai, Lei Wang, Haoxuan Li, and Zhouchen Lin. Time-o-l: Time-series forecasting needs transformed label alignment. *Advances in Neural Information Processing Systems*, 2025f.
- Xinyu Li, Yuchen Luo, Hao Wang, Haoxuan Li, Liuhua Peng, Feng Liu, Yandong Guo, Kun Zhang, and Mingming Gong. Towards accurate time series forecasting via implicit decoding. *Advances in Neural Information Processing Systems*, 2025.
- Xingjian Wu, Xiangfei Qiu, Zhengyu Li, Yihang Wang, Jilin Hu, Chenjuan Guo, Hui Xiong, and Bin Yang. CATCH: Channel-aware multivariate time series anomaly detection via frequency patching. In *ICLR*, 2025a.
- Xingjian Wu, Xiangfei Qiu, Hongfan Gao, Jilin Hu, Bin Yang, and Chenjuan Guo. K²VAE: A koopman-kalman enhanced variational autoencoder for probabilistic time series forecasting. In *ICML*, 2025b.
- Xingjian Wu, Xiangfei Qiu, Hanyin Cheng, Zhengyu Li, Jilin Hu, Chenjuan Guo, and Bin Yang. Enhancing time series forecasting through selective representation spaces: A patch perspective. In *NeurIPS*, 2025c.
- Xiangfei Qiu, Jilin Hu, Lekui Zhou, Xingjian Wu, Junyang Du, Buang Zhang, Chenjuan Guo, Aoying Zhou, Christian S. Jensen, Zhenli Sheng, and Bin Yang. TFB: Towards comprehensive and fair benchmarking of time series forecasting methods. In *Proc. VLDB Endow.*, pages 2363–2377, 2024.
- Xiangfei Qiu, Xingjian Wu, Hanyin Cheng, Xvyuan Liu, Chenjuan Guo, Jilin Hu, and Bin Yang. DBLoss: Decomposition-based loss function for time series forecasting. In *NeurIPS*, 2025a.
- Xiangfei Qiu, Xingjian Wu, Yan Lin, Chenjuan Guo, Jilin Hu, and Bin Yang. DUET: Dual clustering enhanced multivariate time series forecasting. In *SIGKDD*, pages 1185–1196, 2025b.
- Xiangfei Qiu, Zhe Li, Wanghui Qiu, Shiyan Hu, Lekui Zhou, Xingjian Wu, Zhengyu Li, Chenjuan Guo, Aoying Zhou, Zhenli Sheng, Jilin Hu, Christian S. Jensen, and Bin Yang. Tab: Unified benchmarking of time series anomaly detection methods. In *Proc. VLDB Endow.*, pages 2775–2789, 2025c.
- Jiaming Ma, Binwu Wang, Pengkun Wang, Zhengyang Zhou, Xu Wang, and Yang Wang. Bist: A lightweight and efficient bi-directional model for spatiotemporal prediction. *Proceedings of the VLDB Endowment*, 18(6):1663–1676, 2025a.

- Jiaming Ma, Zhiqing Cui, Binwu Wang, Pengkun Wang, Zhengyang Zhou, Zhe Zhao, and Yang Wang. Causal learning meet covariates: Empowering lightweight and effective nationwide air quality forecasting. In *International Joint Conference on Artificial Intelligence*, 2025b.
- Jiaming Ma, Binwu Wang, Guanjun Wang, Kuo Yang, Zhengyang Zhou, Pengkun Wang, Xu Wang, and Yang Wang. Less but more: Linear adaptive graph learning empowering spatiotemporal forecasting. In *Advances in Neural Information Processing Systems*, 2025c.
- Jiaming Ma, Binwu Wang, Pengkun Wang, Zhengyang Zhou, Yudong Zhang, Xu Wang, and Yang Wang. Mobimixer: A multi-scale spatiotemporal mixing model for mobile traffic prediction. *IEEE Transactions on Mobile Computing*, 2025d.
- Jiaming Ma, Binwu Wang, Qihe Huang, Guanjun Wang, Pengkun Wang, Zhengyang Zhou, and Yang Wang. Mofo: Empowering long-term time series forecasting with periodic pattern modeling. In *Advances in Neural Information Processing Systems*, 2025e.
- Jiaming Ma, Binwu Wang, Pengkun Wang, Zhengyang Zhou, Xu Wang, and Yang Wang. Robust spatio-temporal centralized interaction for ood learning. In *Forty-second International Conference on Machine Learning*, 2025f.
- Qihe Huang, Lei Shen, Ruixin Zhang, Shouhong Ding, Binwu Wang, Zhengyang Zhou, and Yang Wang. Crossggn: Confronting noisy multivariate time series via cross interaction refinement. *Advances in Neural Information Processing Systems*, 36:46885–46902, 2023b.
- Qihe Huang, Lei Shen, Ruixin Zhang, Jiahuan Cheng, Shouhong Ding, Zhengyang Zhou, and Yang Wang. Hdmixer: Hierarchical dependency with extendable patch for multivariate time series forecasting. In *Proceedings of the AAAI conference on artificial intelligence*, volume 38, pages 12608–12616, 2024.
- Qihe Huang, Zhengyang Zhou, , Yangze Li, Kuo Yang, Binwu Wang, and Yang Wang. Many minds, one goal: Time series forecasting via sub-task specialization and inter-agent cooperation. In *Advances in Neural Information Processing Systems*, 2025a.
- Qihe Huang, Zhengyang Zhou, Kuo Yang, Zhongchao Yi, Xu Wang, and Yang Wang. Timebase: The power of minimalism in efficient long-term time series forecasting. In *Forty-second International Conference on Machine Learning*, 2025b.
- Wenxuan Wang, Kai Wu, Yujian Betterest Li, Dan Wang, and Xiaoyu Zhang. Synthetic series-symbol data generation for time series foundation models. *Advances in Neural Information Processing Systems*, 2025g. URL <https://arxiv.org/abs/2510.08445>.
- Sai Surya Duvvuri and Inderjit S Dhillon. Laser: Attention with exponential transformation. *arXiv e-prints*, pages arXiv–2411, 2024.
- Haixu Wu, Jiehui Xu, Jianmin Wang, and Mingsheng Long. Autoformer: Decomposition transformers with auto-correlation for long-term series forecasting. *NeurIPS*, 34:22419–22430, 2021.
- Minhao Liu, Ailing Zeng, Muxi Chen, Zhijian Xu, Qiuxia Lai, Lingna Ma, and Qiang Xu. Scinet: Time series modeling and forecasting with sample convolution and interaction. *NeurIPS*, 35: 5816–5828, 2022.
- Yuzhou Chen, Ignacio Segovia-Dominguez, Baris Coskunuzer, and Yulia Gel. Tamp-s2genets: coupling time-aware multipersistence knowledge representation with spatio-supra graph convolutional networks for time-series forecasting. In *ICLR*, 2022.
- Diederik P. Kingma and Jimmy Ba. Adam: A method for stochastic optimization. In *ICLR*, 2015.
- Adam Paszke, Sam Gross, Francisco Massa, Adam Lerer, James Bradbury, Gregory Chanan, Trevor Killeen, Zeming Lin, Natalia Gimelshein, Luca Antiga, et al. Pytorch: An imperative style, high-performance deep learning library. *Advances in neural information processing systems*, 32, 2019.
- Ian T Jolliffe. *Principal component analysis for special types of data*. Springer, 2002.

- Qihe Huang, Zhengyang Zhou, Kuo Yang, Zhongchao Yi, Xu Wang, and Yang Wang. Timebase: The power of minimalism in efficient long-term time series forecasting. In *ICML*, 2025c.
- Wenzhen Yue, Xianghua Ying, Ruohao Guo, DongDong Chen, Ji Shi, Bowei Xing, Yuqing Zhu, and Taiyan Chen. Sub-adjacent transformer: Improving time series anomaly detection with reconstruction error from sub-adjacent neighborhoods. *arXiv preprint arXiv:2404.18948*, 2024.
- Hui Chen, Viet Luong, Lopamudra Mukherjee, and Vikas Singh. SimpleTM: A simple baseline for multivariate time series forecasting. In *ICLR*, 2025.
- Shengsheng Lin, Haojun Chen, Haijie Wu, Chunyun Qiu, and Weiwei Lin. Temporal query network for efficient multivariate time series forecasting. In *ICML*, 2025.
- Xiaowen Ma, Zhenliang Ni, Shuai Xiao, and Xinghao Chen. Timepro: Efficient multivariate long-term time series forecasting with variable- and time-aware hyper-state. In *ICML*, 2025g.
- Tom Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared D Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, et al. Language models are few-shot learners. *NIPS*, 33:1877–1901, 2020.
- Abhimanyu Das, Weihao Kong, Andrew Leach, Rajat Sen, and Rose Yu. Long-term forecasting with tide: Time-series dense encoder. *arXiv preprint arXiv:2304.08424*, 2023.
- Yunhao Zhang and Junchi Yan. Crossformer: Transformer utilizing cross-dimension dependency for multivariate time series forecasting. In *ICLR*, 2023.
- Tian Zhou, Ziqing Ma, Qingsong Wen, Xue Wang, Liang Sun, and Rong Jin. Fedformer: Frequency enhanced decomposed transformer for long-term series forecasting. In *ICML*, pages 27268–27286. PMLR, 2022.
- Nate Gruver, Marc Finzi, Shikai Qiu, and Andrew G Wilson. Large language models are zero-shot time series forecasters. *Advances in Neural Information Processing Systems*, 36:19622–19635, 2023.
- Xiaoming Shi, Shiyu Wang, Yuqi Nie, Dianqi Li, Zhou Ye, Qingsong Wen, and Ming Jin. Time-moe: Billion-scale time series foundation models with mixture of experts. In *ICLR*, 2025.
- Gerald Woo, Chenghao Liu, Akshat Kumar, Caiming Xiong, Silvio Savarese, and Doyen Sahoo. Unified training of universal time series forecasting transformers. In *ICML*, 2024.
- Yong Liu, Guo Qin, Zhiyuan Shi, Zhi Chen, Caiyin Yang, Xiangdong Huang, Jianmin Wang, and Mingsheng Long. Sundial: A family of highly capable time series foundation models. *arXiv preprint arXiv:2502.00816*, 2025.
- Mononito Goswami, Konrad Szafer, Arjun Choudhry, Yifu Cai, Shuo Li, and Artur Dubrawski. Moment: A family of open time-series foundation models. In *ICML*, 2024.
- Yunyang Xiong, Zhanpeng Zeng, Rudrasis Chakraborty, Mingxing Tan, Glenn Fung, Yin Li, and Vikas Singh. Nyströmformer: A nyström-based algorithm for approximating self-attention. In *Proceedings of the AAAI conference on artificial intelligence*, volume 35, pages 14138–14148, 2021.

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper's contributions and scope?

Answer: [Yes]

Justification: Our claims in the abstract and introduction clearly state our contributions and accurately match the theoretical and experimental results in Sections 3, 4 and 5, respectively.

Guidelines:

- The answer NA means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A No or NA answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [Yes]

Justification: The limitations are discussed in Appendix L.

Guidelines:

- The answer NA means that the paper has no limitation while the answer No means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate "Limitations" section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren't acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [Yes]

Justification: Theorem 1 is formally stated in Section 3, with its complete proof provided in Appendix A. Additionally, the derivations of the Jacobian matrices of the NormLin layer are detailed in Appendix B.

Guidelines:

- The answer NA means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: The implementation details are presented in Appendix D. The source code and datasets are available at <https://github.com/jackyue1994/OLinear>.

Guidelines:

- The answer NA means that the paper does not include experiments.
- If the paper includes experiments, a No answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
 - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [Yes]

Justification: Data and code are available at this anonymous repository: <https://github.com/jackyue1994/OLinear>.

Guidelines:

- The answer NA means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://nips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so “No” is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://nips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer, etc.) necessary to understand the results?

Answer: [Yes]

Justification: Experimental settings and implementation details are provided in Appendices C and D. The source code is publicly available.

Guidelines:

- The answer NA means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [Yes]

Justification: Error bars are reported in Appendices F and H.2. Specifically, Table 13 presents standard deviations over seven random seeds, and Table 14 shows 99% confidence intervals of averaged performance across four prediction lengths.

Guidelines:

- The answer NA means that the paper does not include experiments.
- The authors should answer "Yes" if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.

- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g. negative error rates).
- If error bars are reported in tables or plots, The authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: The compute resources are presented in Appendix K.

Guidelines:

- The answer NA means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: This work strictly conforms, in every respect, with the NeurIPS Code of Ethics.

Guidelines:

- The answer NA means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer No, they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [NA]

Justification: This is a technical work on time series forecasting without direct societal impact, hence we select [NA]. Broader technical implications are discussed in Appendix M.

Guidelines:

- The answer NA means that there is no societal impact of the work performed.

- If the authors answer NA or No, they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: [NA]

Justification: The datasets and models used in this work are standard in time series forecasting and pose no known safety or misuse risks.

Guidelines:

- The answer NA means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [Yes]

Justification: We use publicly available datasets and open-source code, all of which are properly cited in the main text. All terms of use have been respected.

Guidelines:

- The answer NA means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.

- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [Yes]

Justification: The released code repository includes all necessary source code to reproduce the experimental results and is well documented with usage instructions. All assets are anonymized for review.

Guidelines:

- The answer NA means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [NA]

Justification: This work does not involve crowdsourcing nor research with human subjects.

Guidelines:

- The answer NA means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. Institutional review board (IRB) approvals or equivalent for research with human subjects

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [NA]

Justification: This work does not involve crowdsourcing nor research with human subjects.

Guidelines:

- The answer NA means that the paper does not involve crowdsourcing nor research with human subjects.

- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

16. Declaration of LLM usage

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does not impact the core methodology, scientific rigorousness, or originality of the research, declaration is not required.

Answer: [NA]

Justification: LLMs are not used as an important, original, or non-standard component of our methods.

Guidelines:

- The answer NA means that the core method development in this research does not involve LLMs as any important, original, or non-standard components.
- Please refer to our LLM policy (<https://neurips.cc/Conferences/2025/LLM>) for what should or should not be described.

A Proof of Theorem 1

Proof. For clarity, we denote $\mathbf{z} \triangleq \begin{bmatrix} \mathbf{x} \\ y \end{bmatrix} \in \mathbb{R}^{t+1}$, $\mu_z \triangleq \begin{bmatrix} \mu_{\mathbf{x}} \\ \mu_y \end{bmatrix} \in \mathbb{R}^{t+1}$, and $\Sigma \triangleq \begin{bmatrix} \Sigma_{\mathbf{x}} & \Sigma_{\mathbf{x}y} \\ \Sigma_{\mathbf{x}y}^T & \sigma_y^2 \end{bmatrix}$. According to the definition, the probability density function of \mathbf{z} , i.e., the joint density of \mathbf{x} and y , is

$$p(\mathbf{z}) = p(\mathbf{x}, y) = \frac{1}{(2\pi)^{(t+1)/2} |\Sigma|^{1/2}} \exp\left(-\frac{1}{2}(\mathbf{z} - \mu_z)^T \Sigma^{-1} (\mathbf{z} - \mu_z)\right), \quad (6)$$

where $|\Sigma|$ denotes the determinant of Σ . In the following, we ignore the constant coefficient and focus on the exponential term.

Using the standard block matrix inverse formula [Horn and Johnson, 2012], we have

$$\Sigma^{-1} = \begin{bmatrix} \Sigma_{\mathbf{x}}^{-1} + d^{-1} \Sigma_{\mathbf{x}}^{-1} \Sigma_{\mathbf{x}y} \Sigma_{\mathbf{x}y}^T \Sigma_{\mathbf{x}}^{-1} & -d^{-1} \Sigma_{\mathbf{x}}^{-1} \Sigma_{\mathbf{x}y} \\ -d^{-1} \Sigma_{\mathbf{x}y}^T \Sigma_{\mathbf{x}}^{-1} & d^{-1} \end{bmatrix} \triangleq \begin{bmatrix} \mathbf{A} & -\mathbf{b} \\ -\mathbf{b}^T & d^{-1} \end{bmatrix}, \quad (7)$$

where the scalar $d = \sigma_y^2 - \Sigma_{\mathbf{x}y}^T \Sigma_{\mathbf{x}}^{-1} \Sigma_{\mathbf{x}y}$ is the Schur complement [Horn and Johnson, 2012]. Letting $\mathbf{b} \in \mathbb{R}^{t \times 1}$ denote $d^{-1} \Sigma_{\mathbf{x}}^{-1} \Sigma_{\mathbf{x}y}$, the matrix \mathbf{A} becomes $\Sigma_{\mathbf{x}}^{-1} + d \cdot \mathbf{b} \mathbf{b}^T$.

For notational simplicity, we define $\alpha \triangleq \mathbf{x} - \mu_{\mathbf{x}} \in \mathbb{R}^{t \times 1}$ and $\beta \triangleq y - \mu_y \in \mathbb{R}$. Then, the exponent in Equation 6 (ignoring the constant factor $-1/2$) becomes

$$\begin{aligned} (\mathbf{z} - \mu_z)^T \Sigma^{-1} (\mathbf{z} - \mu_z) &= [\alpha^T \quad \beta] \begin{bmatrix} \Sigma_{\mathbf{x}}^{-1} + d \cdot \mathbf{b} \mathbf{b}^T & -\mathbf{b} \\ -\mathbf{b}^T & d^{-1} \end{bmatrix} \begin{bmatrix} \alpha \\ \beta \end{bmatrix} \\ &= \alpha^T \Sigma_{\mathbf{x}}^{-1} \alpha + d \cdot \alpha^T \mathbf{b} \mathbf{b}^T - 2\beta \alpha^T \mathbf{b} + d^{-1} \beta^2 \\ &= \alpha^T \Sigma_{\mathbf{x}}^{-1} \alpha + d^{-1} (\beta^2 - 2d \cdot \alpha^T \mathbf{b} \cdot \beta + d^2 \cdot \alpha^T \mathbf{b} \mathbf{b}^T \alpha) \\ &= \alpha^T \Sigma_{\mathbf{x}}^{-1} \alpha + d^{-1} (\beta - d\alpha^T \mathbf{b})^2. \end{aligned} \quad (8)$$

Since $p(\mathbf{x}) \propto \exp(-\frac{1}{2}\alpha^T \Sigma_{\mathbf{x}}^{-1} \alpha)$, it follows that

$$\begin{aligned} p(y \mid \mathbf{x}) &= \frac{p(\mathbf{x}, y)}{p(\mathbf{x})} \\ &\propto \exp\left(-\frac{1}{2} (\alpha^T \Sigma_{\mathbf{x}}^{-1} \alpha + d^{-1} (\beta - d\alpha^T \mathbf{b})^2 - \alpha^T \Sigma_{\mathbf{x}}^{-1} \alpha)\right) \\ &= \exp\left(-\frac{(\beta - d\alpha^T \mathbf{b})^2}{2d}\right) \\ &= \exp\left(-\frac{(y - \mu_y - d\alpha^T \mathbf{b})^2}{2d}\right). \end{aligned} \quad (9)$$

Therefore, the conditional density of y given \mathbf{x} is also a Gaussian distribution, whose mean is

$$\begin{aligned} \mu_y + d\alpha^T \mathbf{b} &= \mu_y + d\mathbf{b}^T \alpha \\ &= \mu_y + d \cdot d^{-1} \Sigma_{\mathbf{x}y}^T \Sigma_{\mathbf{x}}^{-1} (\mathbf{x} - \mu_{\mathbf{x}}) \\ &= \mu_y + \Sigma_{\mathbf{x}y}^T \Sigma_{\mathbf{x}}^{-1} (\mathbf{x} - \mu_{\mathbf{x}}), \end{aligned} \quad (10)$$

and the variance is $d = \sigma_y^2 - \Sigma_{\mathbf{x}y}^T \Sigma_{\mathbf{x}}^{-1} \Sigma_{\mathbf{x}y}$. This completes the proof. \square

B Jacobian matrix comparison of self-attention and NormLin

In this section, we analyze the gradients of the non-linear transformations in the self-attention mechanism and the NormLin module, focusing on the attention/weight rows, denoted as $\mathbf{a} \in \mathbb{R}^N$.

B.1 Jacobian matrix of Softmax in self-attention

Let $\mathbf{c} \triangleq \text{Softmax}(\mathbf{a}) \in \mathbb{R}^N$. We first compute the partial derivatives element-wise, and then rewrite the results in matrix form. Based on the definition of Softmax, the i -th element of \mathbf{c} is $\mathbf{c}_i = \frac{e^{\mathbf{a}_i}}{\sum_{k=1}^N e^{\mathbf{a}_k}}$. The partial derivative of \mathbf{c}_i with respect to \mathbf{a}_i is:

$$\frac{\partial \mathbf{c}_i}{\partial \mathbf{a}_i} = \frac{e^{\mathbf{a}_i}}{\sum_{k=1}^N e^{\mathbf{a}_k}} - \frac{e^{\mathbf{a}_i} \cdot e^{\mathbf{a}_i}}{(\sum_{k=1}^N e^{\mathbf{a}_k})^2} = \mathbf{c}_i - \mathbf{c}_i^2. \quad (11)$$

For $j \neq i$, the partial derivative of \mathbf{c}_i with respect to \mathbf{a}_j is:

$$\frac{\partial \mathbf{c}_i}{\partial \mathbf{a}_j} = -\frac{e^{\mathbf{a}_i} \cdot e^{\mathbf{a}_j}}{(\sum_{k=1}^N e^{\mathbf{a}_k})^2} = -\mathbf{c}_i \mathbf{c}_j \quad (12)$$

Combining Equations (11) and (12), we obtain:

$$\frac{\partial \mathbf{c}_i}{\partial \mathbf{a}} = \mathbf{c}_i[\delta_{1i}, \delta_{2i}, \dots, \delta_{Ni}] - \mathbf{c}_i \mathbf{c}^\top, \quad (13)$$

where $\delta_{ji} = 1$ if $j = i$, otherwise 0. Therefore, the Jacobian matrix $\frac{\partial \mathbf{c}}{\partial \mathbf{a}}$ can be written as

$$\frac{\partial \mathbf{c}}{\partial \mathbf{a}} = \text{Diag}(\mathbf{c}) - \mathbf{c} \mathbf{c}^\top, \quad (14)$$

where $\text{Diag}(\mathbf{c})$ is the diagonal matrix with \mathbf{c} as its diagonal. We can observe that the Jacobian matrix is a function of \mathbf{c} . Since the Softmax function could cause sharp value concentration, the Jacobian matrix may exhibit sparsity, with most entries being close to zero [Surya Duvvuri and Dhillon, 2024].

B.2 Jacobian matrix of the NormLin layer

Let $\mathbf{c} \triangleq \text{Norm}_{L1}(\text{Softplus}(\mathbf{a})) \in \mathbb{R}^N$, where $\text{Norm}_{L1}(\cdot)$ denotes the L1 normalization. For notational simplicity, let $\mathbf{b} \triangleq \text{Softplus}(\mathbf{a})$. We first analyze $\frac{\partial \mathbf{c}}{\partial \mathbf{b}}$. Since all entries of \mathbf{b} are positive, we have:

$$\mathbf{c}_i = \frac{\mathbf{b}_i}{\sum_{k=1}^N \mathbf{b}_k}. \quad (15)$$

Therefore, it follows that

$$\frac{\partial \mathbf{c}_i}{\partial \mathbf{b}_i} = \frac{1}{\sum_{k=1}^N \mathbf{b}_k} - \frac{\mathbf{b}_i}{(\sum_{k=1}^N \mathbf{b}_k)^2} = \frac{1}{\|\mathbf{b}\|_1^2} (\|\mathbf{b}\|_1 - \mathbf{b}_i). \quad (16)$$

For $j \neq i$, the partial derivative $\frac{\partial \mathbf{c}_i}{\partial \mathbf{b}_j}$ is:

$$\frac{\partial \mathbf{c}_i}{\partial \mathbf{b}_j} = -\frac{\mathbf{b}_i}{(\sum_{k=1}^N \mathbf{b}_k)^2} = -\frac{1}{\|\mathbf{b}\|_1^2} \mathbf{b}_i. \quad (17)$$

Combining Equations (16) and (17), we obtain the Jacobian matrix:

$$\frac{\partial \mathbf{c}}{\partial \mathbf{b}} = \frac{1}{\|\mathbf{b}\|_1^2} (\|\mathbf{b}\|_1 \mathbf{I} - \mathbf{b} \mathbf{1}^\top), \quad (18)$$

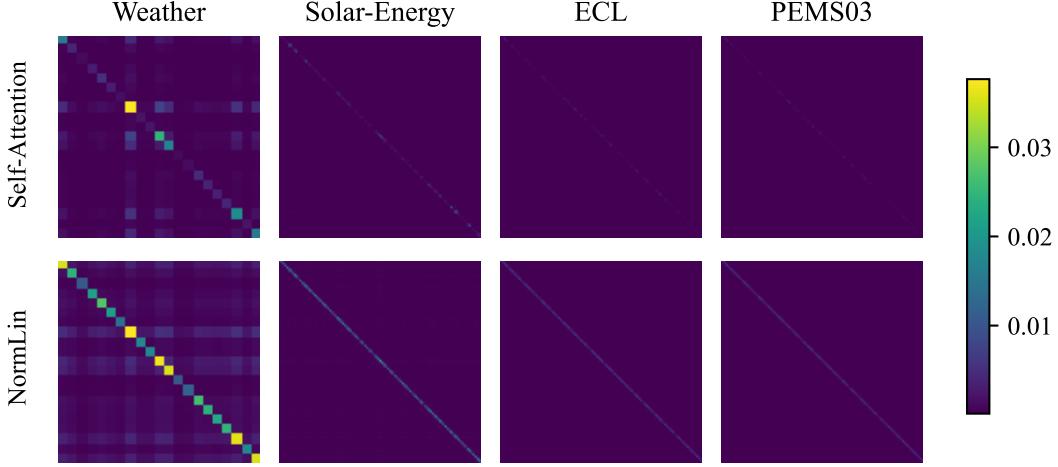


Figure 6: Comparison of Jacobian matrices of $\text{Softmax}(\cdot)$ in self-attention and $\text{Norm}_{\text{L}1}$ ($\text{Softplus}(\cdot)$) in the NormLin layer. The absolute values of the matrix entries are visualized. NormLin and self-attention exhibit similar patterns; however, NormLin generally yields diagonal entries of higher magnitude, suggesting more effective gradient propagation. The same input vector is used for both to ensure a fair comparison.

where $\mathbf{1} \in \mathbb{R}^N$ is the all-ones vector.

For the derivative of \mathbf{b}_i with respect to \mathbf{a}_j with $1 \leq i, j \leq N$, we have

$$\begin{aligned} \frac{\partial \mathbf{b}_i}{\partial \mathbf{a}_j} &= \delta_{ij} \cdot \frac{\partial \mathbf{b}_i}{\partial \mathbf{a}_i} = \delta_{ij} \cdot \frac{\partial}{\partial \mathbf{a}_i} \ln(1 + e^{\mathbf{a}_i}) \\ &= \delta_{ij} \cdot \frac{e^{\mathbf{a}_i}}{1 + e^{\mathbf{a}_i}} = \delta_{ij} \cdot \text{Sigmoid}(\mathbf{a}_i). \end{aligned} \quad (19)$$

Therefore, $\frac{\partial \mathbf{b}}{\partial \mathbf{a}}$ can be written as

$$\frac{\partial \mathbf{b}}{\partial \mathbf{a}} = \text{Diag}(\text{Sigmoid}(\mathbf{a})) \triangleq \text{Diag}(\tilde{\mathbf{b}}), \quad (20)$$

where $\text{Sigmoid}(\cdot)$ operates element-wise, and $\tilde{\mathbf{b}} \triangleq \text{Sigmoid}(\mathbf{a})$ is defined for clarity. Using the chain rule, the Jacobian matrix of \mathbf{c} with respect to \mathbf{a} can be derived as follows:

$$\begin{aligned} \frac{\partial \mathbf{c}}{\partial \mathbf{a}} &= \frac{\partial \mathbf{c}}{\partial \mathbf{b}} \cdot \frac{\partial \mathbf{b}}{\partial \mathbf{a}} = \frac{1}{\|\mathbf{b}\|_1^2} (\|\mathbf{b}\|_1 \mathbf{I} - \mathbf{b} \mathbf{1}^\top) \text{Diag}(\tilde{\mathbf{b}}) \\ &= \frac{1}{\|\mathbf{b}\|_1^2} \left(\|\mathbf{b}\|_1 \text{Diag}(\tilde{\mathbf{b}}) - \mathbf{b} \tilde{\mathbf{b}}^\top \right) \\ &= \frac{1}{\|\mathbf{b}\|_1} \left(\text{Diag}(\tilde{\mathbf{b}}) - \bar{\mathbf{b}} \tilde{\mathbf{b}}^\top \right), \end{aligned} \quad (21)$$

where we use the fact that $\mathbf{1}^\top \text{Diag}(\tilde{\mathbf{b}}) = \tilde{\mathbf{b}}^\top$, and $\bar{\mathbf{b}} = \frac{\mathbf{b}}{\|\mathbf{b}\|_1}$ is the normalized \mathbf{b} .

Equations (21) and (14) share similar structures, with the former introducing an additional learnable scaling factor $\frac{1}{\|\mathbf{b}\|_1}$, which offers greater flexibility. The entries of $\tilde{\mathbf{b}} = \text{Sigmoid}(\mathbf{a})$ in Equation (21) are generally larger than those of $\mathbf{c} = \text{Softmax}(\mathbf{a})$ in Equation (14), particularly when \mathbf{a} contains small values near zero. This property may contribute to stronger gradients in NormLin. Figure 6 illustrates the Jacobian matrices of the self-attention mechanism and the NormLin layer under the same input \mathbf{a} , highlighting that NormLin tends to produce stronger gradient values.

C Dataset description

In this work, the following real-world datasets are used for performance evaluation. The dataset details are presented in Table 11.

- **ETT** datasets [Zhou et al., 2021] record seven channels related to electricity transformers from July 2016 to July 2018. It contains four datasets: ETTh1 and ETTh2, with hourly recordings, and ETTm1 and ETTm2, with 15-minute recordings.
- **Weather** [Wu et al., 2021] contains 21 meteorological variables (e.g., air temperature, humidity) recorded every 10 minutes in 2020.
- **ECL** [Wu et al., 2021] records hourly electricity consumption of 321 consumers from July 2016 to July 2019.
- **Traffic** [Wu et al., 2021] includes hourly road occupancy rates from 862 sensors in the Bay Area from January 2015 to December 2016.
- **Exchange** [Wu et al., 2021] collects daily exchange rates for eight countries from January 1990 to October 2010.
- **Solar-Energy** [Lai et al., 2018] records the solar power output every 10 minutes from 137 photovoltaic plants in 2006.
- **PEMS** [Liu et al., 2022] provides public traffic sensor data from California, collected every 5 minutes. We use its four subsets (PEMS03, PEMS04, PEMS07, PEMS08) in this study.
- **ILI**² contains weekly records of influenza-like illness patient counts provided by the U.S. CDC from 2002 to 2021.
- **COVID-19** [Chen et al., 2022] includes daily records of COVID-19 hospitalizations in California in 2020, provided by Johns Hopkins University.
- **METR-LA**³ collects traffic network data in Los Angeles every 5 minutes from March to June 2012. A total of 207 channels are included.
- **NASDAQ**⁴ includes daily NASDAQ index and key economic indicators (e.g., interest rate and gold price) from 2010 to 2024.
- **Wiki**⁵ records daily page view counts for Wikipedia articles over two years (2018–2019). The first 99 channels are used in this study.
- **SP500** records daily SP500 index data (e.g., opening price, closing price, and trading volume) from January 1993 to February 2025.
- **DowJones** collects daily stock prices of 27 Dow Jones Industrial Average (DJIA) component companies from January 1999 to March 2025.
- **CarSales** collects daily sales of 10 vehicle brands (e.g., Toyota, Honda) in the U.S. from January 2005 to June 2023. The data are compiled from the *Vehicles Sales* dataset⁶ on Kaggle.
- **Power** contains daily wind and solar energy production (in MW) records for the French grid from April 2020 to June 2023. The data are compiled from the *Wind & Solar Daily Power Production* dataset⁷ on Kaggle.
- **Website**⁸ contains six years of daily visit data (e.g., first-time and returning visits) to an academic website, spanning from September 2014 to August 2020.
- **Unemp** contains monthly unemployment figures for 50 U.S. states and three other territories from January 1976 to March 2020, sourced from the official website of the U.S. Bureau of Labor Statistics⁹.

²<https://gis.cdc.gov/grasp/fluview/fluportaldashboard.html>

³<https://github.com/liyaguang/DCRNN>

⁴<https://www.kaggle.com/datasets/sai14karthik/nasdq-dataset>

⁵<https://www.kaggle.com/datasets/sandeshbhat/wikipedia-web-traffic-201819>

⁶<https://www.kaggle.com/datasets/crisbam/vehicles-sales/data>

⁷<https://www.kaggle.com/datasets/henriupton/wind-solar-electricity-production>

⁸<https://www.kaggle.com/datasets/bobnau/daily-website-visitors>

⁹<https://www.bls.gov/web/laus.supp.toc.htm>

Table 11: Detailed dataset descriptions and statistics. *Dim* denotes the number of variates for each dataset. *Frequency* refers to the time interval between consecutive steps. *Split* indicates the (Train, Validation, Test) ratio. *Prediction len.* represents the prediction lengths. For long-term forecasting, the input length is fixed at 96. For short-term forecasting, we adopt two settings: S1 (Input-12, Predict-{3, 6, 9, 12}) and S2 (Input-36, Predict-{24, 36, 48, 60}). In total, 140 prediction tasks across various datasets and prediction length settings are evaluated in this work.

Dataset	Dim	Frequency	Total len.	Split	Prediction len.	Information
ETTh1, ETTh2	7	Hourly	17420	6:2:2	{96,192,336,720}	Electricity
ETTm1, ETTm2	7	15 mins	69680	6:2:2	{96,192,336,720}	Electricity
Weather	21	10 mins	52696	7:1:2	{96,192,336,720}	Weather
ECL	321	Hourly	26304	7:1:2	{96,192,336,720}	Electricity
Traffic	862	Hourly	17544	7:1:2	{96,192,336,720}	Transportation
Exchange	8	Daily	7588	7:1:2	{96,192,336,720}	Economy
Solar-Energy	137	10 mins	52560	7:1:2	{96,192,336,720}	Energy
PEMS03	358	5 mins	26209	6:2:2	{12,24,48,96}	Transportation
PEMS04	307	5 mins	16992	6:2:2	{12,24,48,96}	Transportation
PEMS07	883	5 mins	28224	6:2:2	{12,24,48,96}	Transportation
PEMS08	170	5 mins	17856	6:2:2	{12,24,48,96}	Transportation
ILI	7	Weekly	966	7:1:2	{3,6,9,12} {24,36,48,60}	Health
COVID-19	55	Daily	335	7:1:2 6:2:2	{3,6,9,12} {24,36,48,60}	Health
METR-LA	207	5 mins	34272	7:1:2	{3,6,9,12} {24,36,48,60}	Transportation
NASDAQ	12	Daily	3914	7:1:2	{3,6,9,12} {24,36,48,60}	Finance
Wiki	99	Daily	730	7:1:2	{3,6,9,12} {24,36,48,60}	Web
SP500	5	Daily	8077	7:1:2	{3,6,9,12} {24,36,48,60}	Finance
DowJones	27	Daily	6577	7:1:2	{3,6,9,12} {24,36,48,60}	Finance
CarSales	10	Daily	6728	7:1:2	{3,6,9,12} {24,36,48,60}	Market
Power	2	Daily	1186	7:1:2	{3,6,9,12} {24,36,48,60}	Energy
Website	4	Daily	2167	7:1:2	{3,6,9,12} {24,36,48,60}	Web
Unemp	53	Monthly	531	6:2:2	{3,6,9,12} {24,36,48,60}	Society

D Implementation details

OLinear is optimized using the ADAM optimizer [Kingma and Ba, 2015], with the initial learning rate selected from $\{10^{-4}, 2 \times 10^{-4}, 5 \times 10^{-4}\}$. The model dimension D is chosen from $\{128, 256, 512\}$, while the embedding size d is set to 16. The batch size is selected from $\{4, 8, 16, 32\}$ depending on the dataset scale. The block number L is chosen from $\{1, 2, 3\}$. Training is performed for up to 50 epochs with early stopping, which halts training if the validation performance does not improve for 10 consecutive epochs. We adopt the weighted L1 loss function following CARD [Wang et al., 2024b]. The experiments are implemented in PyTorch [Paszke et al., 2019] and conducted on an NVIDIA GPU with 24 GB of memory. Hyperparameter sensitivity is discussed in Appendix I.8. For baseline models, we use the reported values from the original papers when available; otherwise, we produce the results using the official code. For the model FilterNet [Yi et al., 2024a], TexFilter is adopted in this work. The code and datasets are available at the following anonymous repository: <https://github.com/jackyue1994/OLinear>.

E Showcases

E.1 Time domain and transformed domain

Figure 7 illustrates the temporal domain and its corresponding transformed domain. In the transformed domain, the correlations along the sequence are effectively suppressed. Adjacent time series exhibit strong consistency in this new domain, which is desirable for forecasting tasks. The orthogonal transformation corresponds to projecting the time series onto the eigenvectors of the temporal Pearson correlation matrix CorrMat_t . The eigenvalues of this symmetric positive semi-definite matrix CorrMat_t typically decay rapidly [Jolliffe, 2002], with only a few being dominant. Consequently, the transformed series exhibits sparsity, with most of the energy concentrated in just a few dimensions. Moreover, since noise tends to be evenly distributed across the transformed dimensions, the signal-to-noise ratio (SNR) in the leading components is improved [Bishop and Nasrabadi, 2006].

E.2 Low-rank attention matrix and high-rank NormLin weight matrix

We replace the NormLin module in OLinear with a standard self-attention mechanism and observe the typical low-rank property in the resulting attention matrices, as shown in Figure 9. This phenomenon can be attributed to the sparsity induced by the transformed domain and the sharp focus introduced by the Softmax operation. The low-rank attention matrix could limit the expressive capacity of the model. In contrast, the weight matrices in NormLin exhibit higher rank, better preserving representation diversity.

Furthermore, the learned NormLin weights closely resemble the across-variate Pearson correlation matrix CorrMat_v , suggesting that the NormLin layer effectively captures multivariate correlations. By directly replacing the learnable weight matrix with $\text{Softmax}(\text{CorrMat}_v)$, we obtain **OLinear-C**, which also achieves competitive performance, as demonstrated in Appendix H.

Reason analysis A potential reason behind this phenomenon is that OrthoTrans can mitigate the intrinsic low-rank property of time series data. The empirical results are presented below. Following TimeBase [Huang et al., 2025c], we compute the median singular values of the correlation matrix among non-overlapping time series patches (patch length = 24, number of patches = 30). As shown in Table 12, OrthoTrans yields larger singular values than other transformation bases (with “Identity” denoting no transformation). In other words, OrthoTrans alleviates low-rank tendencies of time series data and enhances data diversity.

E.3 OrthoTrans enhances attention matrix rank

Figures 10–12 visualize the attention matrices of OLinear (with the NormLin module replaced by self-attention) under different transformation bases. Similar results for iTransformer and PatchTST are presented in Figures 13 and 14, respectively. As shown, OrthoTrans typically yields higher-rank attention matrices compared to DFT, wavelet transforms, or no transformation. As stated earlier,

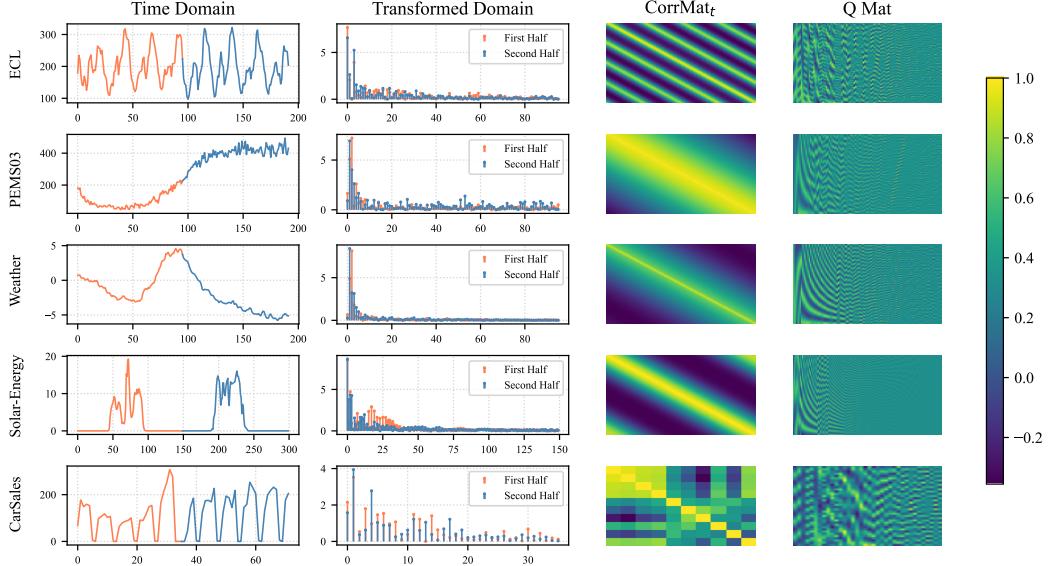


Figure 7: Visualization of temporal series, transformed series, the temporal Pearson correlation matrix, and its corresponding orthogonal matrix. The series are normalized prior to transformation, and the absolute values in the transformed domain are shown for clarity. The X-axis in the first and second columns represents the temporal and feature dimensions, respectively. Adjacent series exhibit consistency in the transformed domain, with sparsity clearly observable. Temporal Pearson correlation matrices on more datasets are presented in Figure 8.

higher-rank attention matrices better preserve the representation space and can potentially enhance the model’s expressive capacity [Han et al., 2023]. This may explain why integrating OrthoTrans as a plug-in consistently improves the performance of Transformer-based forecasters (as shown in Table 5).

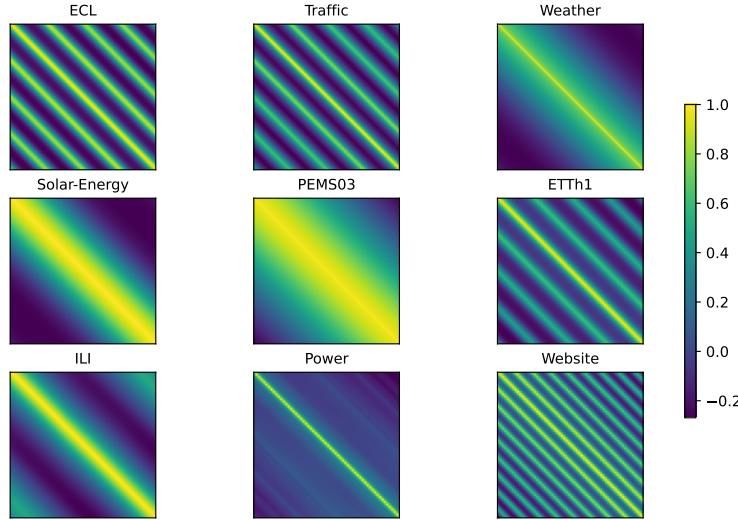


Figure 8: Temporal Pearson correlation matrices on various datasets. Periodic patterns are clearly reflected in the correlation matrices. We incorporate such *a priori* statistical information into the model design, leading to improved forecasting performance.

Table 12: OrthoTrans can effectively mitigate the low-rank characteristics of time series data. The median singular values of the correlation matrices among temporal patches are reported. The orthogonally transformed domain exhibits better data diversity.

Dataset	OrthoTrans	DFT	Wavelet	Identity
ECL	0.446	0.332	0.381	0.162
Traffic	0.568	0.177	0.049	0.002
Weather	0.572	0.352	0.292	0.012
Solar	0.481	0.032	0.023	0.007
PEMS03	0.450	0.356	0.395	0.241

Table 13: Robustness of OLinear performance. Standard deviations are calculated over seven random seeds. $S1$ and $S2$ correspond to Input – 12, Predict – $\{3, 6, 9, 12\}$ and Input – 36, Predict – $\{24, 36, 48, 60\}$, respectively.

Dataset		ECL		Traffic		ETTm1		Solar-Energy	
Metric	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE	
Horizon	96	0.131 \pm 4e-4	0.221 \pm 4e-4	0.398 \pm 3e-3	0.226 \pm 2e-4	0.302 \pm 7e-4	0.334 \pm 3e-4	0.179 \pm 6e-4	0.191 \pm 5e-4
	192	0.150 \pm 1e-3	0.238 \pm 1e-3	0.439 \pm 3e-3	0.241 \pm 4e-4	0.357 \pm 9e-4	0.363 \pm 5e-4	0.209 \pm 8e-4	0.213 \pm 2e-4
	336	0.165 \pm 1e-3	0.254 \pm 1e-3	0.464 \pm 4e-3	0.250 \pm 3e-4	0.387 \pm 2e-3	0.385 \pm 5e-4	0.231 \pm 8e-4	0.229 \pm 7e-5
	720	0.191 \pm 2e-3	0.279 \pm 2e-3	0.502 \pm 4e-3	0.270 \pm 4e-4	0.452 \pm 1e-3	0.426 \pm 6e-4	0.241 \pm 1e-3	0.236 \pm 4e-4
Dataset		Weather		PEMS03		NASDAQ (S2)		Wiki (S1)	
Metric	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE	
Horizon	H1	0.153 \pm 1e-3	0.190 \pm 1e-3	0.060 \pm 3e-4	0.159 \pm 3e-4	0.121 \pm 1e-3	0.216 \pm 9e-4	6.161 \pm 1e-2	0.368 \pm 7e-4
	H2	0.200 \pm 2e-3	0.235 \pm 2e-3	0.078 \pm 6e-4	0.179 \pm 5e-4	0.163 \pm 7e-4	0.261 \pm 9e-4	6.453 \pm 9e-3	0.385 \pm 1e-3
	H3	0.258 \pm 3e-3	0.280 \pm 2e-3	0.104 \pm 6e-4	0.210 \pm 6e-4	0.205 \pm 2e-3	0.296 \pm 2e-3	6.666 \pm 6e-3	0.398 \pm 1e-3
	H4	0.337 \pm 4e-3	0.333 \pm 2e-3	0.140 \pm 2e-3	0.247 \pm 1e-3	0.259 \pm 2e-3	0.336 \pm 2e-3	6.834 \pm 4e-3	0.406 \pm 4e-4
Dataset		DowJones (S2)		SP500 (S2)		CarSales (S1)		Power (S1)	
Metric	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE	
Horizon	H1	7.432 \pm 3e-2	0.664 \pm 9e-4	0.155 \pm 2e-3	0.271 \pm 2e-3	0.303 \pm 2e-3	0.277 \pm 1e-3	0.864 \pm 7e-3	0.688 \pm 4e-3
	H2	10.848 \pm 7e-2	0.799 \pm 1e-3	0.209 \pm 2e-3	0.317 \pm 2e-3	0.315 \pm 2e-3	0.285 \pm 2e-3	0.991 \pm 7e-3	0.742 \pm 3e-3
	H3	14.045 \pm 1e-1	0.914 \pm 1e-3	0.258 \pm 2e-3	0.358 \pm 1e-3	0.327 \pm 8e-4	0.293 \pm 8e-4	1.062 \pm 1e-2	0.770 \pm 4e-3
	H4	16.959 \pm 8e-2	1.017 \pm 3e-3	0.305 \pm 3e-3	0.387 \pm 2e-3	0.336 \pm 5e-4	0.301 \pm 4e-4	1.119 \pm 2e-2	0.789 \pm 6e-3

F Robustness under various random seeds

Table 13 presents the standard deviations for different datasets and prediction lengths using seven random seeds. OLinear demonstrates strong robustness across independent runs. Furthermore, as shown in Table 14, our model exhibits better robustness than state-of-the-art Transformer-based models, TimeMixer++ and iTransformer, as measured by the 99% confidence intervals.

Note that averaging over \mathcal{M} prediction lengths reduces the standard deviation by a factor of $\frac{1}{\sqrt{\mathcal{M}}}$, making the results more robust to the choice of random seeds. Therefore, we prefer to report the average results in this work to mitigate the influence of randomness.

We perform a significance test with 7 random seeds using Student’s t-test to compare OLinear and iTransformer across multiple datasets and prediction horizons. As shown in Table 15, OLinear consistently outperforms iTransformer, with the improvements being statistically significant at the 0.05 level ($p < 0.05$).

G Full results

Table 16 lists the simplified tables from the main text and their full versions in the appendix.

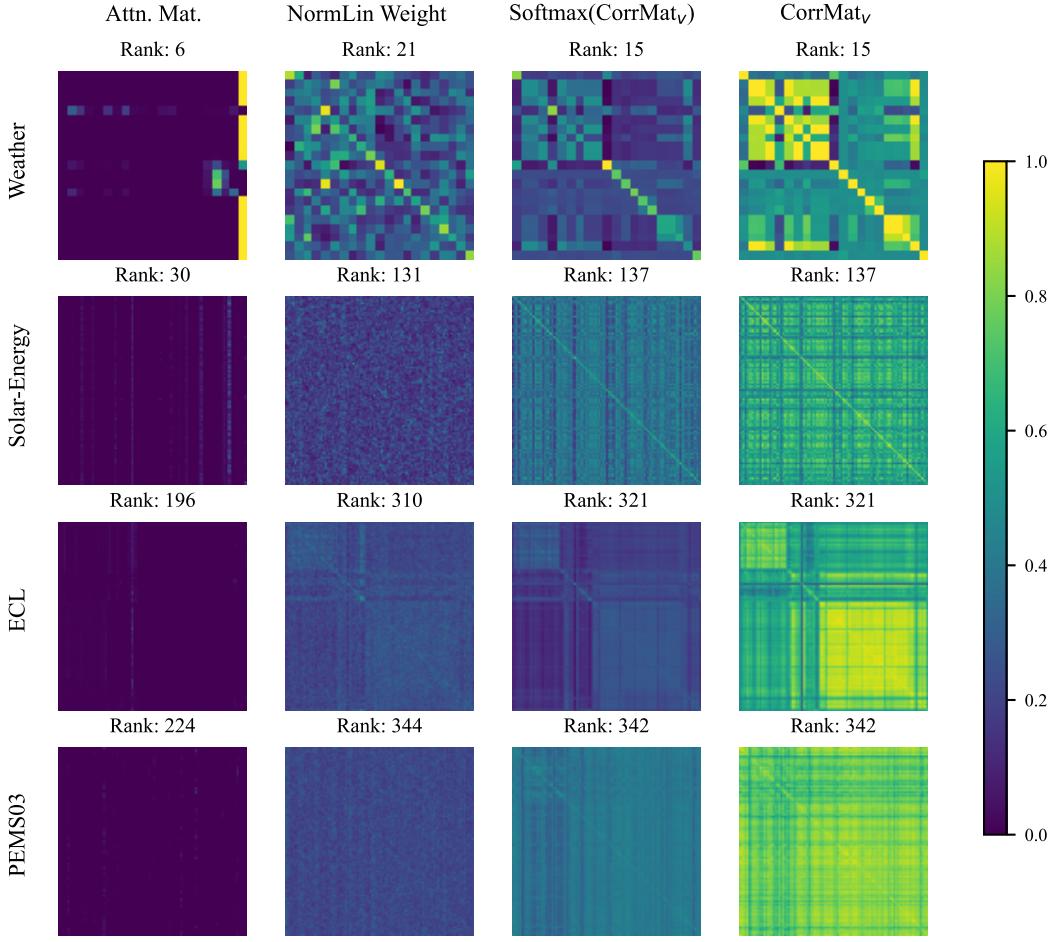


Figure 9: Illustration of attention/weight matrices from self-attention, NormLin, and the across-variate Pearson correlation matrices (with and without Softmax). When applied in OLinear, the self-attention mechanism results in low-rank attention matrices, whereas NormLin produces high-rank weight matrices that better preserve representation diversity. Moreover, the learned NormLin weights resemble the across-variate correlation matrices, suggesting effective modeling of multivariate relationships.

H OLinear-C

H.1 Forecasting performance

Table 25 compares OLinear and OLinear-C on both short- and long-term forecasting tasks. OLinear-C achieves performance comparable to OLinear while using fewer learnable parameters and offering improved efficiency (see Table 41). However, due to its fixed weight matrix, NormLin_c lacks the flexibility required to serve as a plug-and-play module for other forecasters. Therefore, we focus on OLinear in this work.

H.2 Robustness

Table 26 reports the standard deviation of OLinear-C across seven random seeds, demonstrating its robustness to independent runs. For a broader robustness comparison, Table 14 presents 99% confidence intervals for OLinear, OLinear-C, TimeMixer++, and iTransformer.

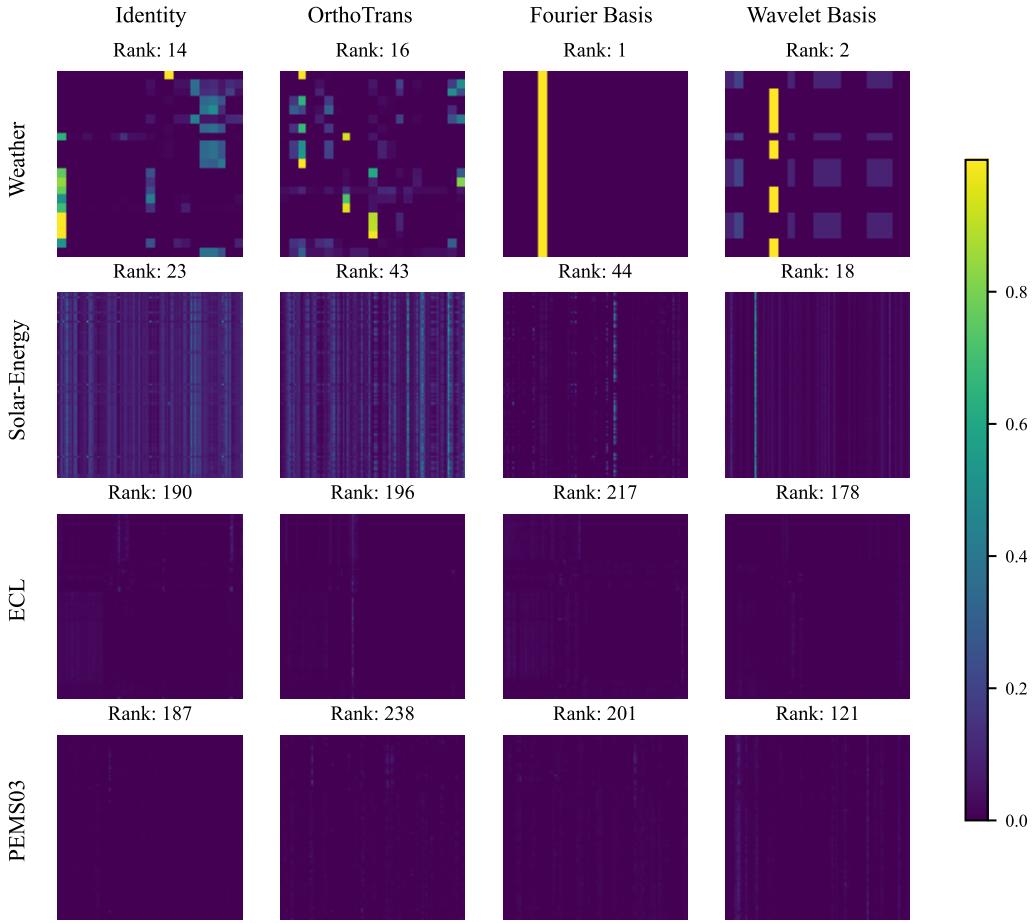


Figure 10: Comparison of attention matrices under various transformation bases for OLinear with the NormLin module replaced by self-attention (**Example 1**). For a fair comparison, the **same** input series is used for different bases on each dataset. ‘Identity’ denotes no transformation, and the Haar wavelet is used as the wavelet basis. In general, OrthoTrans yields higher-rank attention matrices and thus better preserves the representation space.

H.3 Ablation studies

We conduct ablation studies on different transformations of the Pearson correlation matrix CorrMat_v , including Softplus, Sigmoid, ReLU, and the identity function (i.e., no transformation). All variants are followed by row-wise L1 normalization. As shown in Table 27, $\text{Softmax}(\text{CorrMat}_v)$ consistently outperforms the other options and is adopted in the main experiments.

I More experiments

I.1 Quantitation of temporal decorrelation of OrthoTrans

We evaluate the decorrelation effectiveness of OrthoTrans, DFT, and Haar wavelet transforms using a window size of 96. The metric is the off-diagonal Frobenius norm of the temporal correlation matrix, normalized by the number of off-diagonal elements. As shown in Table 28, OrthoTrans reduces the correlation in the original series by 96%, and achieves relative improvements of 84% and 93% over DFT and Haar wavelet transforms, respectively.

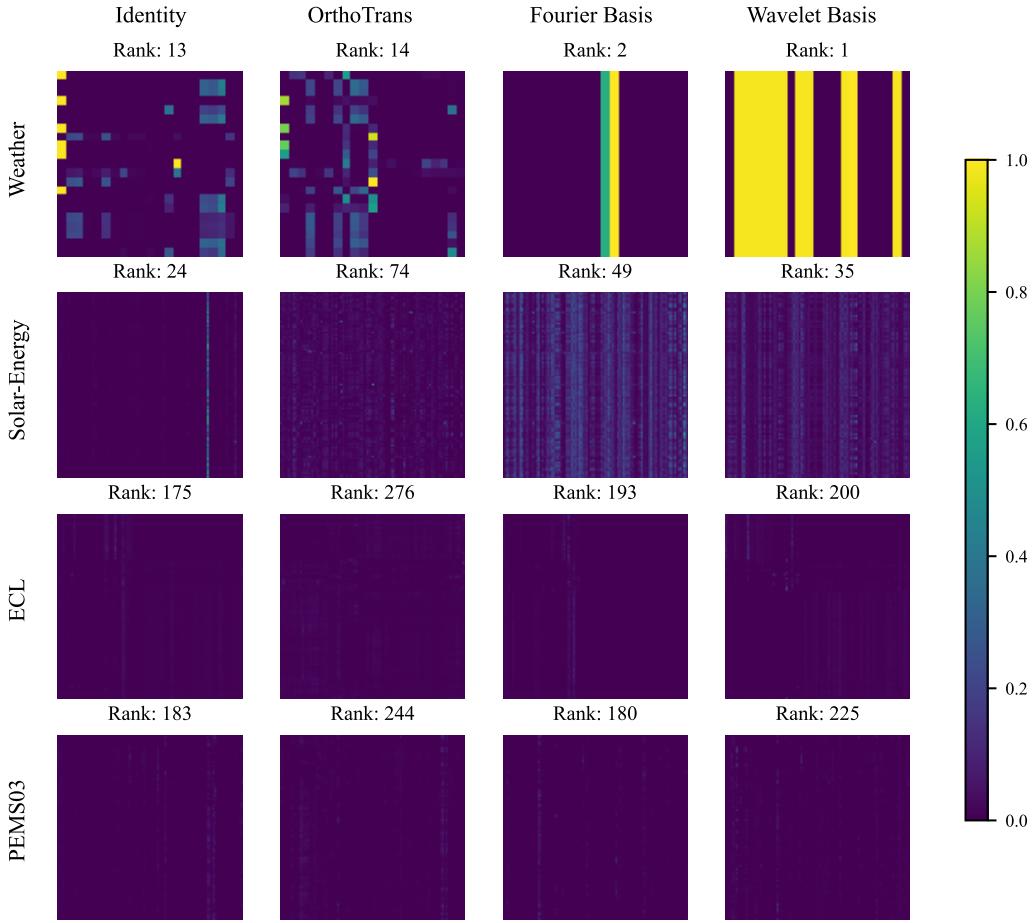


Figure 11: Comparison of attention matrices under various transformation bases for OLinear with the NormLin module replaced by self-attention (**Example 2**). For a fair comparison, the **same** input series is used across bases on each dataset. ‘Identity’ denotes no transformation, and the Haar wavelet is used as the wavelet basis. In general, OrthoTrans yields higher-rank attention matrices and thus better preserves the representation space.

I.2 Performance under varying lookback horizons

We evaluate the forecasting performance of OLinear and its variant, OLinear-C, under varying lookback horizons. Instead of using a fixed lookback length, we search for the optimal horizon within the range of 96 to 720. As shown in Table 29, both models consistently achieve state-of-the-art results across different forecasting lengths and datasets, demonstrating their robustness and strong performance.

I.3 More baselines

We further compare OLinear with several lightweight baselines, including **SimpleTM** [Chen et al., 2025], **TQNet** [Lin et al., 2025], **TimePro** [Ma et al., 2025g], and **TimeBase** [Huang et al., 2025c]. The lookback length is uniformly set to 96 for all methods. As shown in Table 30, OLinear achieves the lowest MAE across all datasets, outperforming SimpleTM, TQNet, TimePro, and TimeBase by 5%, 5%, 7%, and 16% on average, respectively. These results demonstrate that OLinear serves as a strong and efficient new baseline for lightweight time series forecasting.

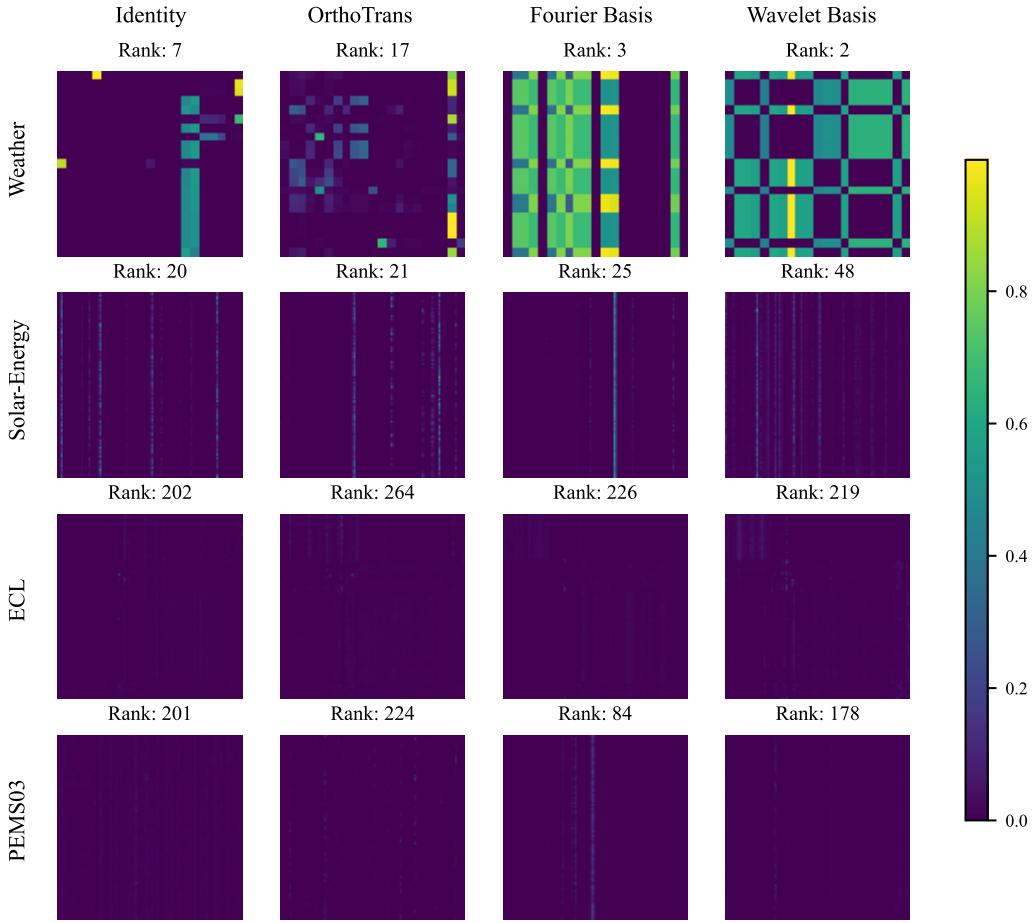


Figure 12: Comparison of attention matrices under various transformation bases for OLinear with the NormLin module replaced by self-attention (**Example 3**). For a fair comparison, the **same** input series is used across bases on each dataset. ‘Identity’ denotes no transformation, and the Haar wavelet is used as the wavelet basis. In general, OrthoTrans yields higher-rank attention matrices and thus better preserves the representation space.

I.4 Few-shot and zero-shot forecasting

To evaluate OLinear’s adaptability to sparse data and its transferability to unseen datasets, we conduct few-shot and zero-shot learning experiments, with the results shown in Tables 31 and 32, respectively. As a linear-based model, OLinear exhibits surprisingly strong generalization capabilities, a characteristic more commonly associated with Transformer-based models [Brown et al., 2020]. Notably, OLinear achieves this with high computational efficiency (see Table 41). Furthermore, OLinear delivers performance on par with or superior to TimeMixer++ [Wang et al., 2025a], despite the latter’s architectural complexity, which includes multi-scale design, series imaging, and dual-axis attention mechanisms.

This performance advantage may be attributed to the high-rank property of the NormLin weight matrix (see Appendix E.2) and its enhanced gradient flow (see Appendix B). In contrast, attention matrices derived from the classic query-key mechanism typically exhibit low-rank characteristics (see Figures 10–14), leading to a compressed representation space [Han et al., 2023].

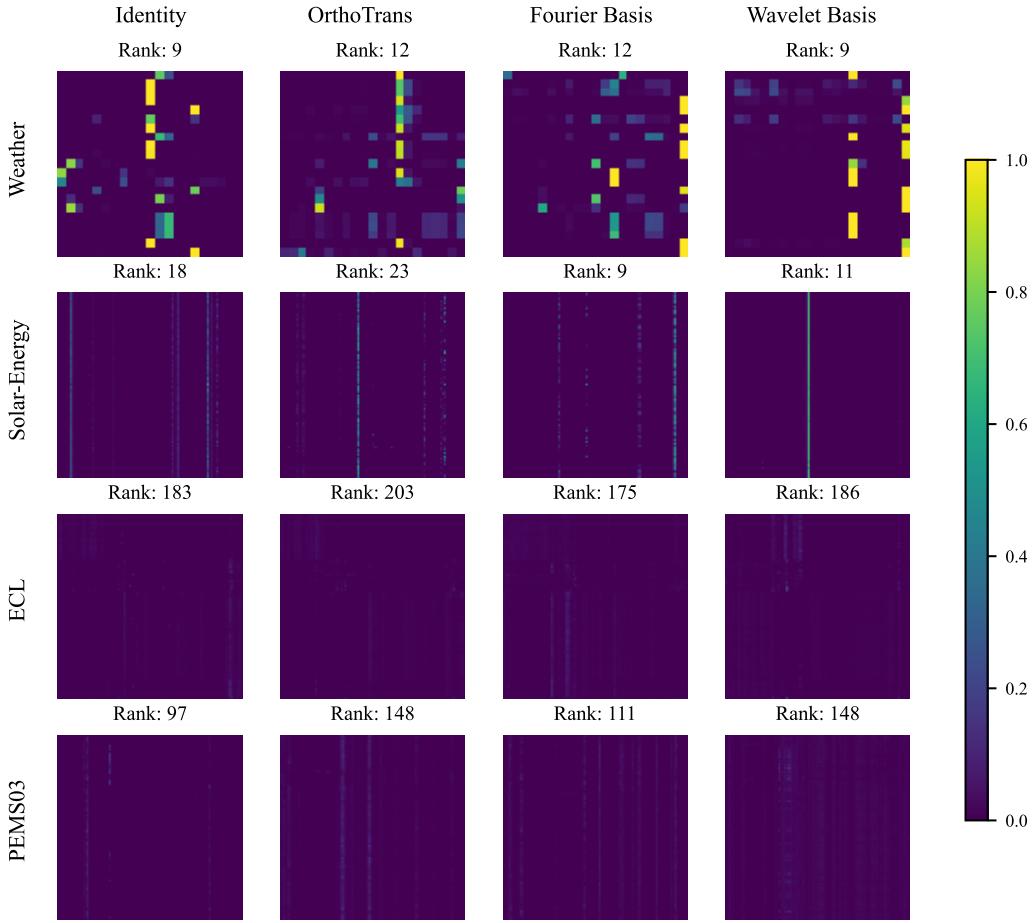


Figure 13: Comparison of attention matrices under various transformation bases for **iTransformer**. For a fair comparison, the **same** input series is used across bases on each dataset. ‘Identity’ denotes no transformation, and the Haar wavelet is used as the wavelet basis. In general, OrthoTrans yields higher-rank attention matrices and thus better preserves the representation space.

I.5 Increasing lookback lengths

Performance under increasing lookback lengths reflects a model’s ability to effectively utilize historical information [Liu et al., 2024a]. Linear-based forecasters have strong potential to benefit from longer lookback windows [Zeng et al., 2023]. As shown in Figure 15, OLinear demonstrates consistent improvements as the lookback horizon increases from 48 to 720, and consistently outperforms state-of-the-art forecasters.

I.6 Training with less data

Training with less data reflects a model’s adaptability and learning efficacy. As shown in Figure 16, OLinear adapts well to decreasing training set sizes, from 100% down to 5%. Our model consistently outperforms others across all training ratios, highlighting its robustness to data-sparse conditions.

I.7 OLinear versus large time series models

We further compare OLinear with the large pre-trained time series models in the full-shot, few-shot and zero-shot scenarios. As shown in Table 33, OLinear surprisingly achieves comparable

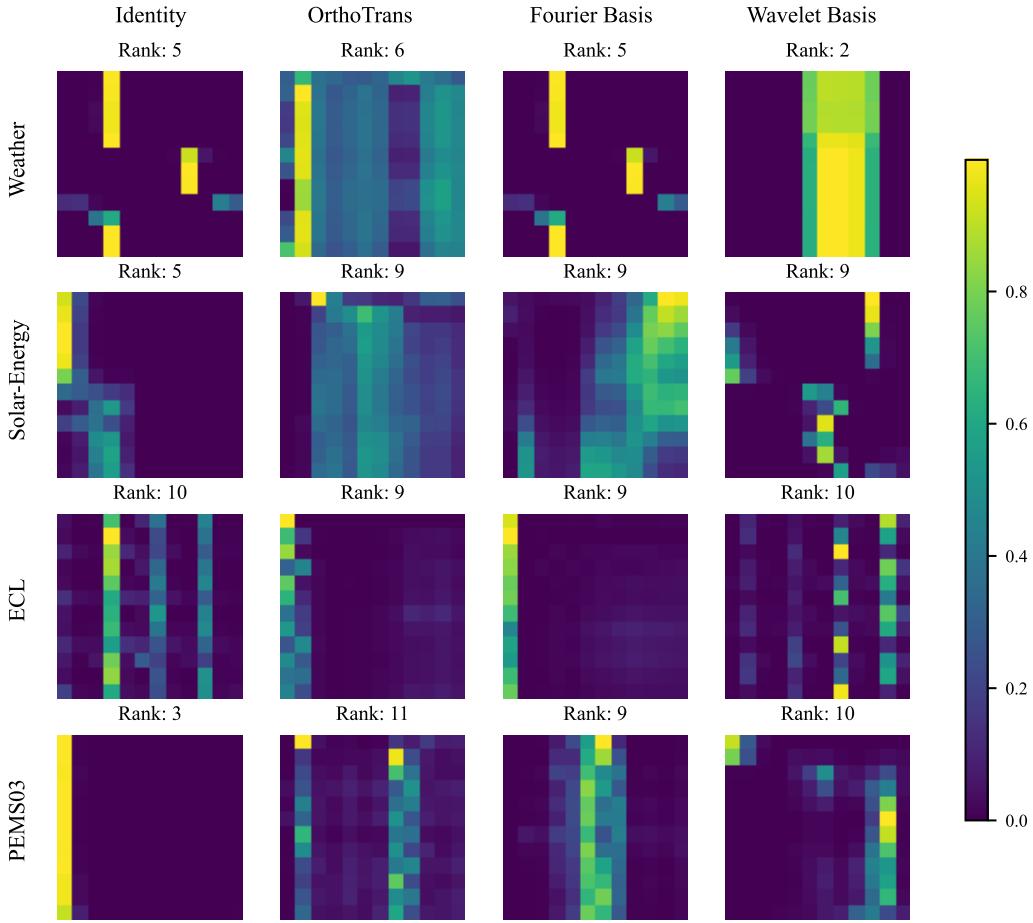


Figure 14: Comparison of attention matrices under various transformation bases for **PatchTST**. For a fair comparison, the **same** input series is used across all bases for each dataset. ‘Identity’ denotes no transformation, and the Haar wavelet is used as the wavelet basis. In general, OrthoTrans yields higher-rank attention matrices and thus better preserves the representation space.

performance to Timer [Liu et al., 2024b], despite not being pre-trained on large-scale datasets. When the full training set is used, OLinear outperforms the fine-tuned Timer on the ECL, Traffic, Weather, and PEMS datasets. Even when only 5% of training data are available, OLinear still performs better on the ETTm1 and ETTm2 datasets. In zero-shot forecasting, our model—trained on ETTh1 and ETTm1 respectively—outperforms state-of-the-art large time series models when evaluated on ETTh2 and ETTm2. These results not only demonstrate OLinear’s strong generalization capability, but also imply that large time series models are still in the early stages of development.

I.8 Hyperparameter sensitivity

We evaluate the hyperparameter sensitivity of OLinear with respect to four key factors: the learning rate (lr), the number of blocks (L), the model dimension (D), and the embedding size (d). As shown in Figure 17, OLinear generally maintains stable performance across different choices of lr , L , and D , although each dataset exhibits its own preferences. Regarding the embedding size d , Table 34 indicates that $d = 16$ represents a *sweet spot*, which we adopt as the default setting in our experiments.

Table 14: 99% confidence intervals of the average performance across four prediction lengths, computed using three times the standard deviation from seven random seeds. The smallest standard deviations are highlighted in **bold**. *OLinear-C* denotes the variant of OLinear where the weight matrix in NormLin is set to Softmax (CorrMat_v), with CorrMat_v being the multivariate correlation matrix.

Model	OLinear (Ours)		OLinear-C (Ours)		TimeMixer++ 2025a		iTransformer 2024a	
	Metric	MSE	MAE	MSE	MAE	MSE	MAE	MSE
Weather	0.237±0.006	0.260± 0.003	0.238± 0.005	0.259±0.004	0.226±0.008	0.262±0.007	0.258±0.009	0.278±0.006
Solar	0.215± 0.001	0.217±0.001	0.215± 0.001	0.217± 3e-4	0.203± 0.001	0.238±0.010	0.233±0.009	0.262±0.007
ECL	0.159± 0.001	0.248±0.002	0.161±0.006	0.249±0.005	0.165±0.011	0.253± 0.001	0.178±0.002	0.270±0.005
Traffic	0.451± 0.003	0.247± 0.001	0.451±0.006	0.247± 0.001	0.416±0.015	0.264±0.013	0.428±0.008	0.282±0.002
ETTh1	0.424±0.003	0.424±0.002	0.424± 0.002	0.424± 0.001	0.419±0.011	0.432±0.015	0.454±0.004	0.447±0.007
ETTh2	0.367± 0.002	0.388±0.002	0.368± 0.002	0.389± 0.001	0.339±0.009	0.380±0.002	0.383±0.004	0.407±0.007
ETTm1	0.374± 0.001	0.377± 0.001	0.375± 0.001	0.378± 0.001	0.369±0.005	0.378±0.007	0.407±0.004	0.410±0.009
ETTm2	0.270± 3e-4	0.313± 2e-4	0.270±4e-4	0.313± 2e-4	0.269±0.002	0.320±0.012	0.288±0.010	0.332±0.003

Table 15: Significance test with Student’s t-test and performance comparison between OLinear and iTransformer across various datasets and prediction horizons.

Dataset & Horizon	OLinear (MAE)	iTrans.(MAE)	p-value between OLinear and iTrans.
ECL: 96	0.221±4e-4	0.240±4e-4	5.12E-11
ECL: 192	0.238±1e-3	0.253±2e-3	3.06E-06
ECL: 336	0.254±1e-3	0.269±1e-3	9.96E-07
ECL: 720	0.279±2e-3	0.317±7e-3	5.59E-05
Traffic: 96	0.226±2e-4	0.268±1e-3	1.61E-12
Traffic: 192	0.241±4e-4	0.276±1e-3	9.48E-10
Traffic: 336	0.250±3e-4	0.283±4e-4	1.87E-13
Traffic: 720	0.270±4e-4	0.302±4e-4	6.15E-13
Weather: 96	0.190±1e-3	0.214±3e-4	1.29E-08
Weather: 192	0.235±2e-3	0.254±1e-3	1.04E-06
Weather: 336	0.280±2e-3	0.296±1e-3	1.93E-06
Weather: 720	0.333±2e-3	0.349±4e-4	5.68E-06

I.9 More metrics

We further conduct a comprehensive evaluation of forecasting performance using three scale-free metrics: the Coefficient of Determination (R^2), the Pearson Correlation Coefficient (r), and the Mean Absolute Scaled Error (MASE), which are defined as follows:

$$\begin{aligned}
 R^2 &= \frac{1}{N} \sum_{n=1}^N \left(1 - \frac{\|\hat{\mathbf{Y}}_{n:} - \mathbf{Y}_{n:}\|_2^2}{\|\mathbf{Y}_{n:} - \bar{\mathbf{Y}}_{n:}\|_2^2} \right), \\
 r &= \frac{1}{N} \sum_{n=1}^N \frac{\sum_{t=1}^{\tau} (\hat{\mathbf{Y}}_{n,t} - \bar{\hat{\mathbf{Y}}}_{n:}) (\mathbf{Y}_{n,t} - \bar{\mathbf{Y}}_{n:})}{\sqrt{\sum_{t=1}^{\tau} (\hat{\mathbf{Y}}_{n,t} - \bar{\hat{\mathbf{Y}}}_{n:})^2} \sqrt{\sum_{t=1}^{\tau} (\mathbf{Y}_{n,t} - \bar{\mathbf{Y}}_{n:})^2}}, \\
 \text{MASE} &= \frac{1}{N} \sum_{n=1}^N \frac{\frac{1}{\tau} \|\hat{\mathbf{Y}}_{n:} - \mathbf{Y}_{n:}\|_1}{\frac{1}{\tau-1} \sum_{t=1}^{\tau-1} |\mathbf{Y}_{n,t} - \mathbf{Y}_{t-1,n}|}.
 \end{aligned} \tag{22}$$

Table 16: Overview of tables in the main paper and their full versions in the appendix.

Tables in paper	Tables in Appendix	Content
Table 2	Table 17	Long-term forecasting
Table 3	Tables 18 and 19	Short-term forecasting
Table 4	Table 20	Ablation studies on various bases
Table 5	Table 21	OrthoTrans as a plug-in
Table 7	Table 25	Performance of OLinear-C
Table 6	Table 22	Ablation studies of OLinear
Table 8	Table 23	NormLin vs self-attention and its variants
Table 9	Table 24	NormLin as a plug-in

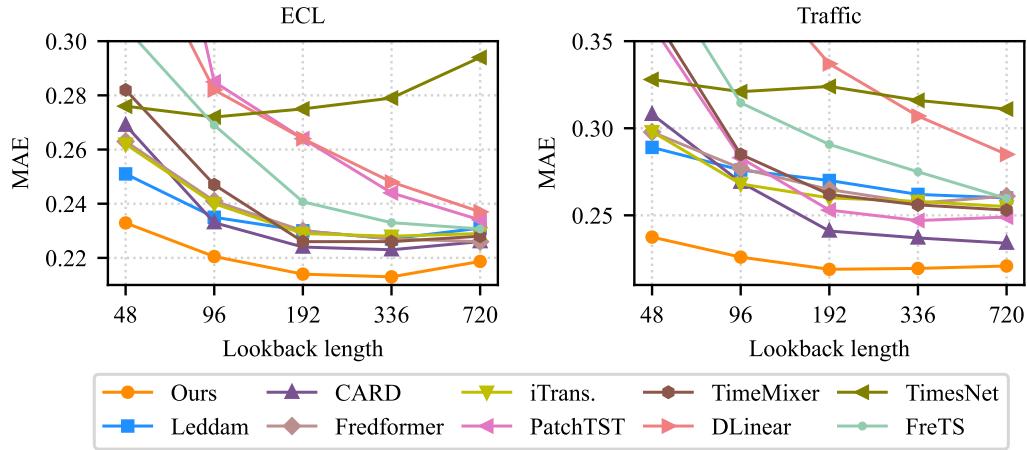


Figure 15: Performance change with increasing lookback lengths. The prediction length is $\tau = 96$.

Here $\mathbf{Y}_{n:}$ and $\hat{\mathbf{Y}}_{n:}$ denote the n -th variate of the ground truth \mathbf{Y} and the prediction $\hat{\mathbf{Y}}$, respectively. $\mathbf{Y}_{n,t}$ refers to the t -th timestep of the n -th variate, and $\overline{(\cdot)}$ represents the mean.

As shown in Table 35, OLinear consistently achieves state-of-the-art performance on these metrics. For example, it an average Pearson correlation coefficient of $r = 0.903$ on the ECL dataset, indicating strong alignment between predictions and ground truth. On the Traffic dataset, OLinear achieves a MASE of 0.764, outperforming the second-best model, CARD [Wang et al., 2024b], by a notable margin of 11.4%.

I.10 Less training data to compute Q matrices

To evaluate the amount of training data required for computing the orthogonal transformation matrices \mathbf{Q}_i and \mathbf{Q}_o , we conduct experiments using only partial data of the training set. As shown in Table 36, OLinear exhibits strong robustness, with minimal performance variation across different data proportions. This property not only enables more efficient preprocessing before training, but also ensures robust performance in few-shot learning scenarios (see Appendix I.4).

J More ablation studies

J.1 Ablations on NormLin design

In the NormLin layer, the Softplus function followed by L1 row-wise normalization is applied to the learnable weight matrix \mathbf{W} to produce attention-like matrices, i.e., $\text{RowNorm}_{\text{L1}}(\text{Softplus}(\mathbf{W}))$. To assess the effects of different transformations and normalization strategies, we consider transformation

Table 17: Full results for the long-term forecasting task. The lookback length T is set to 96 for all baselines. Avg means the average results across four prediction lengths. This table presents the full version of Table 2.

Category	Linear-Based												Transformer-Based												TCN-Based			
	Model	OLinear (Ours)	TimeMixed 2024a	FilterNet 2024a	FITS 2024	DLinear 2023	TimeMixer++ 2025a	Leddam 2024	CARD 2024b	Fredformer 2024	iTrans. 2024a	PatchTST 2023	TimesNet 2023b															
Metric	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE						
ETTh1	96	0.302	0.334	0.320	0.357	0.321	0.361	0.353	0.375	0.345	0.372	0.310	0.334	0.319	0.359	0.316	0.347	0.326	0.361	0.334	0.368	0.329	0.367	0.338	0.375			
	192	0.357	0.363	0.361	0.381	0.367	0.387	0.486	0.445	0.380	0.389	0.48	0.362	0.369	0.383	0.363	0.370	0.363	0.388	0.377	0.391	0.367	0.385	0.374	0.387			
	336	0.387	0.385	0.390	0.404	0.401	0.409	0.531	0.475	0.413	0.376	0.391	0.394	0.402	0.392	0.390	0.395	0.403	0.426	0.399	0.410	0.410	0.411	0.411				
	720	0.452	0.426	0.454	0.441	0.477	0.448	0.600	0.513	0.474	0.453	0.440	0.423	0.460	0.442	0.458	0.425	0.453	0.438	0.491	0.459	0.454	0.439	0.478	0.450			
Avg		0.374	0.377	0.381	0.395	0.392	0.401	0.493	0.452	0.403	0.407	0.369	0.378	0.386	0.397	0.383	0.384	0.384	0.395	0.407	0.410	0.387	0.400	0.400	0.406			
ETTh2	96	0.169	0.249	0.175	0.258	0.175	0.258	0.182	0.266	0.193	0.292	0.170	0.245	0.176	0.257	0.169	0.248	0.177	0.259	0.180	0.264	0.175	0.259	0.187	0.267			
	192	0.232	0.230	0.237	0.299	0.240	0.301	0.253	0.312	0.284	0.362	0.229	0.291	0.243	0.303	0.234	0.292	0.243	0.301	0.250	0.309	0.241	0.302	0.249	0.309			
	336	0.291	0.328	0.298	0.340	0.311	0.347	0.313	0.349	0.369	0.427	0.303	0.343	0.303	0.341	0.294	0.339	0.302	0.340	0.311	0.348	0.305	0.343	0.321	0.351			
	720	0.389	0.387	0.391	0.396	0.414	0.405	0.416	0.406	0.554	0.522	0.373	0.399	0.400	0.398	0.390	0.388	0.397	0.396	0.412	0.407	0.402	0.400	0.408	0.403			
Avg		0.270	0.313	0.275	0.323	0.285	0.328	0.291	0.333	0.350	0.401	0.269	0.320	0.281	0.325	0.272	0.317	0.279	0.324	0.288	0.332	0.281	0.326	0.291	0.333			
ETTh3	96	0.360	0.382	0.375	0.400	0.382	0.402	0.385	0.394	0.386	0.400	0.361	0.403	0.377	0.394	0.383	0.391	0.373	0.392	0.386	0.405	0.414	0.419	0.384	0.402			
	192	0.416	0.414	0.429	0.421	0.430	0.429	0.434	0.422	0.437	0.432	0.416	0.441	0.424	0.422	0.435	0.420	0.433	0.420	0.441	0.436	0.460	0.445	0.436	0.429			
	336	0.457	0.438	0.484	0.458	0.472	0.451	0.476	0.444	0.481	0.459	0.430	0.434	0.459	0.472	0.447	0.470	0.437	0.487	0.458	0.501	0.466	0.491	0.469	0.469			
	720	0.463	0.462	0.498	0.482	0.481	0.473	0.465	0.462	0.519	0.516	0.467	0.451	0.463	0.459	0.471	0.461	0.467	0.456	0.503	0.491	0.500	0.488	0.521	0.500			
Avg		0.424	0.424	0.447	0.440	0.441	0.440	0.431	0.456	0.452	0.419	0.432	0.431	0.429	0.442	0.429	0.435	0.426	0.454	0.447	0.469	0.454	0.458	0.450				
ETL	96	0.284	0.329	0.289	0.341	0.293	0.343	0.292	0.340	0.333	0.387	0.276	0.328	0.292	0.343	0.293	0.342	0.292	0.342	0.330	0.349	0.329	0.342	0.340	0.374			
	192	0.360	0.379	0.372	0.392	0.374	0.396	0.377	0.391	0.477	0.476	0.342	0.379	0.367	0.389	0.363	0.381	0.371	0.389	0.380	0.400	0.387	0.400	0.402	0.414			
	336	0.409	0.415	0.386	0.414	0.417	0.430	0.416	0.425	0.594	0.541	0.346	0.398	0.412	0.424	0.411	0.418	0.382	0.409	0.428	0.432	0.426	0.433	0.452	0.452			
	720	0.415	0.431	0.412	0.434	0.449	0.460	0.418	0.437	0.831	0.657	0.392	0.415	0.419	0.438	0.416	0.431	0.415	0.434	0.427	0.445	0.431	0.446	0.462	0.468			
Avg		0.367	0.388	0.365	0.395	0.383	0.407	0.376	0.398	0.559	0.515	0.339	0.380	0.373	0.399	0.368	0.390	0.365	0.393	0.393	0.383	0.407	0.384	0.405	0.414	0.427		
ETL2	96	0.131	0.221	0.153	0.247	0.147	0.245	0.198	0.274	0.197	0.282	0.135	0.222	0.141	0.235	0.141	0.233	0.147	0.234	0.141	0.235	0.146	0.250	0.168	0.272			
	192	0.150	0.238	0.166	0.256	0.160	0.250	0.163	0.263	0.422	0.196	0.285	0.147	0.235	0.159	0.250	0.160	0.258	0.162	0.253	0.199	0.289	0.184	0.289	0.289			
	336	0.163	0.254	0.185	0.277	0.173	0.267	0.444	0.490	0.209	0.301	0.164	0.245	0.173	0.268	0.173	0.263	0.177	0.273	0.178	0.269	0.215	0.305	0.198	0.300			
	720	0.191	0.279	0.225	0.310	0.210	0.309	0.532	0.551	0.245	0.333	0.212	0.310	0.201	0.295	0.197	0.284	0.213	0.304	0.225	0.317	0.256	0.337	0.220	0.320			
Avg		0.367	0.388	0.365	0.395	0.383	0.407	0.376	0.398	0.559	0.515	0.339	0.380	0.373	0.399	0.368	0.390	0.365	0.393	0.393	0.383	0.407	0.384	0.405	0.414	0.427		
ETL3	96	0.082	0.200	0.086	0.205	0.091	0.211	0.087	0.208	0.088	0.218	0.085	0.214	0.086	0.207	0.084	0.202	0.084	0.202	0.086	0.206	0.088	0.205	0.107	0.234			
	192	0.171	0.293	0.193	0.312	0.186	0.305	0.185	0.306	0.176	0.315	0.175	0.313	0.171	0.298	0.178	0.302	0.177	0.299	0.176	0.299	0.226	0.344					
	336	0.331	0.414	0.356	0.433	0.380	0.449	0.342	0.425	0.313	0.427	0.316	0.420	0.325	0.435	0.313	0.418	0.382	0.409	0.428	0.432	0.426	0.433	0.452	0.452			
	720	0.837	0.688	0.912	0.712	0.896	0.712	0.846	0.694	0.839	0.695	0.851	0.689	0.831	0.686	0.851	0.691	0.749	0.651	0.847	0.691	0.901	0.714	0.964	0.746			
Avg		0.355	0.399	0.387	0.416	0.388	0.419	0.365	0.408	0.354	0.414	0.357	0.409	0.354	0.402	0.362	0.402	0.333	0.391	0.360	0.403	0.367	0.404	0.416	0.443			
ETL4	96	0.398	0.226	0.462	0.283	0.430	0.294	0.601	0.361	0.650	0.396	0.392	0.253	0.426	0.276	0.419	0.269	0.404	0.277	0.395	0.268	0.446	0.283	0.593	0.321			
	192	0.439	0.241	0.473	0.296	0.452	0.307	0.603	0.365	0.598	0.370	0.402	0.258	0.458	0.289	0.443	0.276	0.426	0.290	0.417	0.276	0.540	0.354	0.617	0.336			
	336	0.464	0.250	0.498	0.296	0.470	0.316	0.609	0.366	0.605	0.373	0.428	0.263	0.486	0.297	0.460	0.283	0.437	0.292	0.433	0.283	0.551	0.358	0.629	0.336			
	720	0.502	0.270	0.506	0.313	0.498	0.323	0.648	0.387	0.645	0.394	0.441	0.282	0.498	0.313	0.490	0.299	0.462	0.305	0.467	0.302	0.586	0.375	0.640	0.350			
Avg		0.451	0.485	0.298	0.463	0.310	0.615	0.370	0.625	0.383	0.416	0.264	0.467	0.294	0.453	0.282	0.433	0.291	0.428	0.282	0.531	0.343	0.620	0.336				
ETL5	96	0.153	0.190	0.163	0.209	0.162	0.207	0.196	0.236	0.196	0.255	0.155	0.205	0.156	0.202	0.150	0.188	0.163	0.207	0.174	0.214	0.177	0.218	0.220				

Table 18: Full results for the short-term forecasting task (Part 1). For prediction lengths $\tau \in \{3, 6, 9, 12\}$, the lookback horizon is $T = 12$; for $\tau \in \{24, 36, 48, 60\}$, it is $T = 36$. The best results are highlighted in **bold**, and the second-best results are underlined. This table presents the detailed version of Table 3.

Model	OrthoLienar (Ours)	TimeMix. 2024a	FilterNet 2024a	FITS 2024	DLinear 2023	TimeMix.++ 2025a	Leddam 2024	CARD 2024b	Fredformer 2024	iTrans. 2024a	PatchTST 2023	TimesNet 2023b												
Metric	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE						
ILI																								
3	0.468	0.349	0.659	0.435	0.660	0.437	1.461	0.743	1.280	0.747	0.658	0.430	0.551	0.388	0.597	0.392	0.528	0.403	0.555	0.395	0.646	0.417	0.627	0.420
6	0.923	0.516	1.306	0.643	1.140	0.606	2.337	0.974	2.054	0.967	1.273	0.645	<u>1.021</u>	0.569	1.246	0.610	1.128	0.604	1.124	0.586	1.269	0.629	1.147	0.610
9	1.289	0.655	2.070	0.842	1.815	0.798	3.397	1.197	2.771	1.138	2.009	0.840	1.881	0.795	2.041	0.829	1.804	0.786	1.794	0.772	2.021	0.830	1.796	0.785
12	1.698	0.791	2.792	1.018	2.435	0.953	4.244	1.361	3.497	1.284	2.683	1.004	2.421	0.964	2.746	0.996	2.610	0.989	<u>2.272</u>	0.884	2.788	1.016	2.349	0.920
Avg	1.094	0.578	1.707	0.734	1.512	0.698	2.860	1.069	2.400	1.034	1.656	0.730	1.468	0.679	1.658	0.707	1.518	0.696	1.437	0.659	1.681	0.723	1.480	0.684
COVID-19																								
3	1.100	0.487	1.237	0.547	1.195	0.555	2.039	0.790	2.386	0.909	1.298	0.584	1.216	0.570	<u>1.103</u>	0.521	1.165	0.548	1.193	0.561	1.220	0.573	1.201	0.704
6	1.750	0.619	2.003	0.739	1.839	0.711	2.683	0.919	3.220	1.053	1.833	0.682	1.782	0.689	1.919	0.735	1.465	0.685	1.933	0.755	1.982	0.762	2.405	0.808
9	2.239	0.734	2.594	0.860	2.537	0.897	3.147	1.050	3.803	1.160	2.472	0.822	2.407	0.866	2.358	0.841	2.145	0.845	2.441	0.879	2.633	0.916	2.858	0.969
12	2.538	0.831	3.103	0.981	<u>2.782</u>	0.956	3.630	1.156	4.524	1.288	3.273	1.084	2.851	0.991	2.857	0.971	2.833	0.984	2.819	0.984	3.050	1.030	2.993	0.964
Avg	1.764	0.802	2.020	0.878	2.073	0.885	4.130	1.465	3.083	1.217	1.955	0.857	1.982	0.875	2.260	0.938	<u>1.947</u>	0.899	1.993	0.887	2.128	0.885	2.139	0.931
METRA																								
3	0.207	0.171	0.205	0.192	0.204	0.189	0.211	0.198	0.218	0.231	0.204	0.189	0.204	0.191	0.210	0.180	0.205	0.188	0.205	0.188	0.204	0.190	0.221	0.204
6	0.301	0.207	0.297	0.230	0.296	0.228	0.306	0.237	0.307	0.278	0.298	0.230	0.293	0.227	0.311	0.219	0.298	0.227	0.300	0.229	0.298	0.227	0.308	0.238
9	0.382	0.238	0.381	0.264	0.377	0.261	0.392	0.271	0.386	0.316	0.384	0.266	0.369	0.264	0.401	0.253	0.385	0.263	0.386	0.265	0.382	0.263	0.387	0.273
12	0.452	0.263	0.455	0.295	0.449	0.290	0.467	0.300	0.452	0.353	0.456	0.294	0.442	0.292	0.474	0.281	0.457	0.292	0.460	0.295	0.456	0.292	0.462	0.298
Avg	0.335	0.220	0.334	0.245	0.331	0.242	0.344	0.251	0.341	0.294	0.335	0.245	0.327	0.243	0.349	0.233	0.336	0.242	0.338	0.244	0.335	0.243	0.344	0.253
NASDAQ																								
3	0.036	0.092	0.035	0.093	0.038	0.100	0.042	0.111	0.044	0.123	0.036	0.098	0.040	0.103	0.037	0.096	0.039	0.102	0.040	0.105	0.038	0.099	0.049	0.123
6	0.049	0.117	0.049	0.118	0.052	0.126	0.056	0.134	0.062	0.155	0.050	0.121	0.054	0.128	0.052	0.123	0.053	0.126	0.054	0.129	0.053	0.124	0.061	0.142
9	0.062	0.137	0.062	0.139	0.065	0.145	0.069	0.154	0.082	0.189	0.063	0.144	0.066	0.147	0.063	0.141	0.067	0.147	0.068	0.150	0.065	0.145	0.073	0.161
12	0.073	0.154	0.073	0.156	0.076	0.162	0.081	0.172	0.100	0.215	0.075	0.161	0.078	0.164	0.075	0.158	0.079	0.165	0.078	0.165	0.077	0.161	0.088	0.179
Avg	0.0838	0.402	0.881	0.499	0.875	0.489	0.927	0.510	0.819	0.550	0.799	0.480	0.878	0.490	0.929	0.466	0.898	0.495	0.916	0.501	0.893	0.502	0.890	0.488
Wiki																								
3	6.161	0.368	6.209	0.392	6.234	0.402	7.470	0.496	6.254	0.438	6.149	0.389	6.148	0.383	6.183	0.378	6.190	0.387	6.237	0.393	6.112	0.380	7.597	0.510
6	6.453	0.385	6.475	0.402	6.460	0.401	8.326	0.544	6.579	0.467	6.436	0.401	6.455	0.397	6.465	0.393	6.696	0.404	6.484	0.400	6.425	0.395	7.962	0.515
9	6.666	0.398	6.702	0.418	6.697	0.416	8.869	0.564	6.776	0.508	6.714	0.420	6.687	0.412	6.714	0.415	6.768	0.411	6.689	0.411	6.743	0.426	8.150	0.524
12	6.834	0.406	6.902	0.426	6.899	0.426	9.394	0.608	6.927	0.513	6.852	0.421	6.899	0.424	6.852	0.415	7.168	0.424	6.868	0.419	6.814	0.414	8.117	0.533
Avg	6.262	0.440	6.315	0.468	6.342	0.467	7.813	0.681	<u>6.205</u>	0.539	6.323	0.479	6.286	0.463	6.285	0.453	5.931	0.453	6.275	0.458	6.212	0.444	7.310	0.623
1 st Count	19	47	8	1	2	1	0	0	0	0	7	0	6	0	0	0	10	1	0	0	5	1	0	0

Table 19: Full results for the short-term forecasting task (Part 2). For prediction lengths $\tau \in \{3, 6, 9, 12\}$, the lookback horizon is $T = 12$; for $\tau \in \{24, 36, 48, 60\}$, it is $T = 36$. The best and second-best results are highlighted in **bold**, and underlined, respectively. This table presents the detailed version of Table 3.

	Model	OrthoLienar (Ours)	TimeMix. 2024a	FilterNet 2024a	FITS 2024	DLinear 2023	TimeMix.++ 2025a	Leddam 2024	CARD 2024b	Fredformer 2024	iTrans. 2024a	PatchTST 2023	TimesNet 2023b						
Metric	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE	
SP500	3	0.035	<u>0.126</u>	0.038	0.137	0.046	0.154	0.046	0.157	0.047	0.152	0.040	0.141	0.047	0.155	<u>0.036</u>	<u>0.130</u>	0.045	0.153
	6	0.053	<u>0.158</u>	0.056	0.167	0.071	0.193	0.067	0.190	0.070	0.197	0.057	0.168	0.064	<u>0.181</u>	<u>0.054</u>	<u>0.162</u>	0.066	0.187
	9	0.070	<u>0.181</u>	0.074	0.192	0.080	0.203	0.087	0.218	0.091	0.220	0.076	0.196	0.081	0.206	0.072	0.184	0.080	0.205
	12	0.088	<u>0.204</u>	0.092	0.213	0.097	0.224	0.106	0.242	0.113	0.251	0.093	0.213	0.104	0.236	0.089	<u>0.206</u>	0.103	0.234
	Avg	0.061	<u>0.167</u>	0.065	0.177	0.073	0.193	0.076	0.202	0.080	0.205	0.066	0.179	0.074	0.194	0.063	<u>0.170</u>	0.073	0.195
	24	0.155	<u>0.271</u>	0.159	0.288	0.181	0.317	0.193	0.334	0.189	0.330	<u>0.172</u>	<u>0.305</u>	0.175	0.308	0.156	<u>0.276</u>	0.181	0.315
	36	0.209	<u>0.317</u>	0.218	0.343	0.224	<u>0.341</u>	0.259	0.389	0.250	0.363	0.227	0.344	0.232	0.358	0.206	<u>0.319</u>	0.365	0.225
	48	0.258	<u>0.358</u>	0.264	0.367	0.280	0.384	0.324	0.439	0.291	0.398	0.272	0.383	<u>0.276</u>	<u>0.388</u>	0.258	<u>0.354</u>	0.283	0.394
	60	0.305	<u>0.387</u>	0.322	0.416	0.332	0.416	0.391	0.486	0.377	0.475	<u>0.319</u>	<u>0.413</u>	0.325	0.423	<u>0.303</u>	<u>0.385</u>	0.341	0.438
	Avg	0.231	<u>0.333</u>	0.241	0.353	0.254	0.365	0.291	0.412	0.277	0.391	0.247	<u>0.361</u>	0.252	0.369	0.231	<u>0.333</u>	<u>0.261</u>	0.378
	3	1.550	<u>0.276</u>	1.568	<u>0.276</u>	1.587	0.278	1.611	0.283	1.574	0.280	1.555	0.278	1.538	<u>0.277</u>	<u>1.563</u>	0.276	1.563	<u>0.277</u>
	6	2.569	<u>0.361</u>	2.593	<u>0.362</u>	2.634	0.364	2.661	0.370	2.613	0.367	2.591	<u>0.364</u>	2.588	0.364	<u>2.579</u>	<u>0.362</u>	2.619	0.364
	9	3.565	<u>0.430</u>	3.564	<u>0.431</u>	3.651	0.434	3.704	0.443	3.621	0.436	3.572	<u>0.433</u>	3.644	0.440	3.569	0.430	<u>3.575</u>	0.432
	12	4.517	<u>0.490</u>	4.472	<u>0.490</u>	<u>4.634</u>	0.496	4.666	0.503	4.586	0.497	4.630	0.496	4.622	0.498	4.528	0.490	4.605	0.493
	Avg	3.050	<u>0.389</u>	3.049	<u>0.390</u>	3.126	0.393	3.160	0.400	3.099	0.395	3.087	0.393	3.098	0.395	3.060	0.389	<u>0.390</u>	0.391
	24	7.432	<u>0.664</u>	8.327	0.683	8.006	0.683	7.974	0.690	7.590	0.670	8.283	0.689	8.029	7.416	<u>0.665</u>	7.758	0.672	
	36	10.848	<u>0.799</u>	11.192	0.813	12.011	0.823	11.907	0.837	10.986	0.803	14.754	0.856	11.962	10.799	<u>0.798</u>	11.456	0.808	
	48	14.045	<u>0.914</u>	15.278	0.945	14.814	0.933	15.821	0.969	14.157	0.922	16.893	0.970	<u>15.266</u>	0.937	13.881	<u>0.912</u>	14.696	<u>0.921</u>
	60	16.959	<u>1.017</u>	20.997	1.067	18.932	<u>1.054</u>	19.320	1.077	18.018	1.035	18.881	1.059	<u>18.407</u>	<u>1.045</u>	17.257	<u>1.021</u>	18.058	1.032
	Avg	12.321	<u>0.848</u>	13.948	0.877	13.439	0.873	13.755	0.893	12.688	0.857	14.703	0.893	13.416	<u>0.872</u>	12.338	<u>0.849</u>	12.992	0.858
	3	0.303	<u>0.277</u>	<u>0.307</u>	0.296	0.304	0.291	0.411	0.375	0.396	0.383	<u>0.309</u>	0.289	0.316	0.300	<u>0.295</u>	<u>0.311</u>	<u>0.281</u>	<u>0.276</u>
	6	0.315	<u>0.285</u>	0.327	<u>0.317</u>	0.300	0.401	0.367	0.390	0.376	0.324	<u>0.300</u>	<u>0.320</u>	0.303	0.335	<u>0.307</u>	0.326	0.309	<u>0.294</u>
	9	0.327	<u>0.293</u>	0.336	<u>0.315</u>	<u>0.329</u>	0.308	0.418	0.380	0.406	0.390	0.331	<u>0.306</u>	<u>0.332</u>	0.311	<u>0.319</u>	0.339	<u>0.318</u>	<u>0.305</u>
	12	0.336	<u>0.301</u>	0.343	0.319	<u>0.335</u>	0.312	0.424	0.385	0.410	0.390	0.343	<u>0.314</u>	<u>0.336</u>	0.313	<u>0.357</u>	<u>0.324</u>	0.348	0.324
	Avg	0.320	<u>0.289</u>	0.328	0.309	<u>0.321</u>	0.303	0.413	0.377	0.400	0.385	0.327	0.302	0.326	0.306	<u>0.310</u>	<u>0.311</u>	<u>0.298</u>	<u>0.288</u>
	24	0.320	<u>0.302</u>	0.320	0.318	0.318	0.319	0.359	0.347	<u>0.354</u>	0.350	0.323	0.320	0.325	0.322	0.337	<u>0.321</u>	0.319	0.326
	36	0.334	<u>0.315</u>	0.332	0.331	<u>0.331</u>	0.330	0.373	0.360	<u>0.368</u>	0.365	0.351	0.348	0.337	0.333	0.348	<u>0.333</u>	0.333	0.335
	48	0.347	<u>0.327</u>	0.345	<u>0.342</u>	0.341	0.385	0.370	<u>0.382</u>	<u>0.379</u>	<u>0.351</u>	0.342	0.351	0.346	<u>0.345</u>	0.349	<u>0.331</u>	<u>0.332</u>	0.347
	60	0.358	<u>0.335</u>	0.355	<u>0.351</u>	<u>0.352</u>	0.349	0.399	0.385	<u>0.388</u>	0.380	0.363	0.352	0.361	0.353	<u>0.353</u>	0.359	0.349	<u>0.344</u>
	Avg	0.340	<u>0.320</u>	0.338	0.336	<u>0.336</u>	0.335	0.379	0.365	<u>0.373</u>	0.368	0.347	0.340	0.343	0.338	<u>0.338</u>	<u>0.324</u>	<u>0.327</u>	0.338
	3	0.864	<u>0.688</u>	0.850	<u>0.683</u>	<u>0.842</u>	<u>0.685</u>	0.899	0.712	0.876	0.708	0.843	<u>0.685</u>	0.859	0.696	0.896	0.691	0.865	0.697
	6	0.991	<u>0.742</u>	<u>0.971</u>	<u>0.743</u>	0.988	0.747	1.079	0.785	0.991	0.761	0.997	0.748	1.023	0.761	<u>1.013</u>	0.750	0.994	0.751
	9	1.062	<u>0.770</u>	1.024	<u>0.763</u>	1.063	0.778	1.150	0.804	1.051	0.790	1.050	<u>0.767</u>	1.071	0.783	<u>1.085</u>	0.781	1.073	0.784
	12	1.119	<u>0.789</u>	1.087	<u>0.788</u>	1.125	0.803	1.266	0.851	1.110	0.814	1.107	0.792	1.130	0.804	<u>1.124</u>	0.792	1.135	0.807
	Avg	1.009	<u>0.747</u>	0.983	<u>0.744</u>	1.004	0.753	1.098	0.788	1.007	0.768	0.999	0.748	1.021	0.761	<u>1.029</u>	0.753	1.017	0.760
	24	1.343	<u>0.870</u>	1.341	0.881	1.410	0.916	1.491	0.944	1.390	0.916	1.340	<u>0.877</u>	1.397	0.909	1.406	<u>0.886</u>	1.410	0.913
	36	1.445	<u>0.903</u>	1.420	<u>0.914</u>	1.590	0.968	1.621	0.994	1.518	0.957	1.446	0.920	1.509	0.951	<u>1.506</u>	<u>0.921</u>	1.538	0.952
	48	1.559	<u>0.946</u>	1.567	<u>0.963</u>	1.680	1.009	1.775	1.052	1.610	0.995	1.467	<u>0.933</u>	1.646	0.999	1.583	0.957	1.652	1.008
	60	1.602	<u>0.971</u>	1.609	<u>0.988</u>	1.776	1.053	1.958	1.122	1.679	1.020	<u>1.626</u>	1.006	1.727	1.043	1.693	1.003	1.752	1.049
	Avg	1.487	<u>0.922</u>	1.484	<u>0.937</u>	1.614	0.986	1.711	1.028	1.549	0.972	1.470	<u>0.934</u>	1.570	0.975	1.547	0.942	1.588	0.981
	3	0.077	<u>0.199</u>	0.086	0.215	0.084	0.213	0.191	0.320	0.159	0.288	0.083	0.208	<u>0.082</u>	0.210	0.099	<u>0.225</u>	0.080	0.207
	6	0.103	<u>0.227</u>	0.124	0.248	0.116	0.242	0.235	0.356	0.182	0.302	<u>0.115</u>	0.237	0.114	0.239	0.135	<u>0.255</u>	0.116	0.241
	9	0.135	<u>0.251</u>	0.159	0.275	0.151	0.269	0.276	0.372	0.220	0.330	0.156	0.273	<u>0.147</u>	0.246	0.176	<u>0.282</u>	0.150	<u>0.268</u>
	12	0.176	<u>0.281</u>	0.204	0.306	0.194	0.297	0.409	0.484	0.255	0.355	0.193	0.296	0.213	0.307	0.196	0.301	<u>0.144</u>	<u>0.269</u>
	Avg	0.123	<u>0.240</u>	0.143	0.261	0.136	0.255	0.278	0.383	0.204	0.319	0.137	0.253	0.135	0.253	0.154	<u>0.267</u>	0.135	0.254
	24	0.186	<u>0.306</u>	0.229	0.335	0.273	0.357	0.431	0.469	0.315	0.393	0.231	0.349	0.240	0.345	0.325	0.370	0.216	0.335
	36	0.272	<u>0.356</u>	0.361	0.420	0.401	0.441	0.554	0.552	<u>0.385</u>	0.447	0.328	0.397	0.327	0.405	0.428			

Table 20: Comparison of different transformation bases. For $S1$ and $S2$, the prediction lengths are $\tau \in \{3, 6, 9, 12\}$ and $\tau \in \{24, 36, 48, 60\}$, respectively. *Wavelet1* and *Wavelet2* denote the classic Haar and discrete Meyer wavelets, respectively. *Identity* means that no transformation is applied. This is a complete version of Table 4.

Base	ECL		Solar-Energy		PEMS03		Power (S2)		ILI (S1)		COVID-19 (S2)		METR-LA (S2)		
Metric	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE	
Ours	H1	0.131	0.221	0.179	0.191	0.060	0.159	1.343	0.870	0.468	0.349	4.474	1.180	0.650	0.337
	H2	0.150	0.238	0.209	0.213	0.078	0.179	1.445	0.903	0.923	0.516	7.241	1.670	0.800	0.388
	H3	0.165	0.254	0.231	0.229	0.104	0.210	1.559	0.946	1.289	0.655	10.076	1.985	0.905	0.427
	H4	0.191	0.279	0.241	0.236	0.140	0.247	1.602	0.971	1.698	0.791	12.079	2.182	0.999	0.457
	Avg	0.159	0.248	0.215	0.217	0.095	0.199	1.487	0.922	1.094	0.578	8.467	1.754	0.838	0.402
Fourier	H1	0.132	0.222	0.181	0.193	0.061	0.161	1.469	0.917	0.486	0.350	5.192	1.378	0.654	0.337
	H2	0.153	0.241	0.214	0.215	0.078	0.182	1.584	0.962	0.885	0.500	7.553	1.699	0.799	0.390
	H3	0.165	0.253	0.236	0.232	0.110	0.216	1.625	0.970	1.438	0.664	11.096	2.070	0.910	0.427
	H4	0.195	0.283	0.247	0.238	0.155	0.258	1.778	1.020	2.264	0.822	12.819	2.209	1.007	0.459
	Avg	0.161	0.250	0.219	0.219	0.101	0.204	1.614	0.967	1.268	0.584	9.165	1.839	0.843	0.403
Wavelet1	H1	0.134	0.223	0.182	0.194	0.061	0.161	1.476	0.917	0.483	0.351	4.826	1.317	0.647	0.336
	H2	0.153	0.241	0.217	0.216	0.079	0.183	1.586	0.959	0.943	0.520	7.117	1.652	0.801	0.390
	H3	0.166	0.255	0.238	0.234	0.114	0.219	1.711	1.011	1.251	0.653	10.133	1.995	0.903	0.428
	H4	0.186	0.276	0.250	0.240	0.176	0.278	1.880	1.060	1.786	0.795	12.588	2.233	1.011	0.464
	Avg	0.160	0.249	0.221	0.221	0.107	0.210	1.663	0.987	1.116	0.580	8.666	1.799	0.840	0.404
Wavelet2	H1	0.133	0.222	0.186	0.198	0.062	0.162	1.511	0.929	0.461	0.349	5.398	1.413	0.653	0.340
	H2	0.154	0.243	0.222	0.220	0.080	0.184	1.570	0.953	0.935	0.526	7.199	1.663	0.803	0.392
	H3	0.165	0.255	0.242	0.236	0.115	0.221	1.701	1.005	1.393	0.683	10.616	2.051	0.909	0.432
	H4	0.196	0.283	0.254	0.243	0.175	0.272	1.876	1.060	1.921	0.818	12.583	2.232	1.007	0.462
	Avg	0.162	0.251	0.226	0.224	0.108	0.210	1.664	0.987	1.177	0.594	8.949	1.840	0.843	0.406
Chebyshev	H1	0.134	0.223	0.184	0.198	0.061	0.160	1.338	0.864	0.456	0.341	5.381	1.391	0.661	0.339
	H2	0.156	0.244	0.220	0.222	0.079	0.182	1.564	0.969	0.945	0.524	8.001	1.767	0.803	0.391
	H3	0.182	0.270	0.244	0.240	0.118	0.220	1.604	0.950	1.527	0.700	11.244	2.108	0.923	0.433
	H4	0.399	0.444	0.256	0.246	0.163	0.266	1.774	1.076	1.940	0.824	12.695	2.233	1.028	0.466
	Avg	0.218	0.295	0.226	0.226	0.105	0.207	1.570	0.965	1.217	0.597	9.330	1.875	0.854	0.407
Laguerre	H1	0.139	0.227	0.192	0.204	0.065	0.167	1.425	0.899	0.595	0.395	5.632	1.456	0.659	0.346
	H2	0.156	0.243	0.227	0.227	0.085	0.191	1.658	0.985	1.083	0.577	8.729	1.859	0.829	0.407
	H3	0.173	0.262	0.251	0.244	0.122	0.226	1.774	1.013	1.581	0.737	10.381	2.024	0.933	0.444
	H4	0.199	0.286	0.261	0.248	0.171	0.271	1.780	1.039	2.155	0.896	12.467	2.221	1.053	0.484
	Avg	0.167	0.255	0.233	0.230	0.111	0.214	1.659	0.984	1.353	0.651	9.302	1.890	0.868	0.420
Legendre	H1	0.134	0.224	0.200	0.208	0.074	0.181	1.502	0.926	0.501	0.366	4.965	1.350	0.652	0.337
	H2	0.152	0.239	0.238	0.231	0.080	0.183	1.603	0.962	0.964	0.529	7.396	1.703	0.799	0.389
	H3	0.165	0.256	0.262	0.247	0.113	0.219	1.790	1.031	1.446	0.695	9.481	1.935	0.900	0.427
	H4	0.193	0.282	0.274	0.253	0.168	0.269	1.846	1.060	1.799	0.822	12.358	2.203	1.015	0.463
	Avg	0.161	0.250	0.243	0.235	0.109	0.213	1.685	0.995	1.177	0.603	8.550	1.798	0.841	0.404
Identity	H1	0.134	0.224	0.186	0.199	0.063	0.163	1.419	0.900	0.477	0.352	4.841	1.284	0.659	0.342
	H2	0.156	0.245	0.221	0.220	0.080	0.184	1.453	0.903	0.956	0.523	7.828	1.757	0.807	0.395
	H3	0.167	0.257	0.244	0.237	0.117	0.222	1.582	0.956	1.307	0.659	10.153	2.003	0.921	0.433
	H4	0.194	0.282	0.256	0.244	0.165	0.266	1.715	1.023	1.874	0.812	12.603	2.232	1.006	0.461
	Avg	0.163	0.252	0.227	0.225	0.106	0.209	1.542	0.945	1.153	0.587	8.856	1.819	0.848	0.408

functions from the set {Softplus, Identity, Softmax, Sigmoid, ReLU} and normalization methods from {L1, L2} (all applied row-wise). Here, Identity denotes no transformation.

As shown in Table 37, the combination ‘Softplus + L1’ yields the best performance among the 10 candidates. On average, L1 normalization leads to a 10% performance improvement over L2, while Softplus slightly outperforms Softmax and other transformation functions.

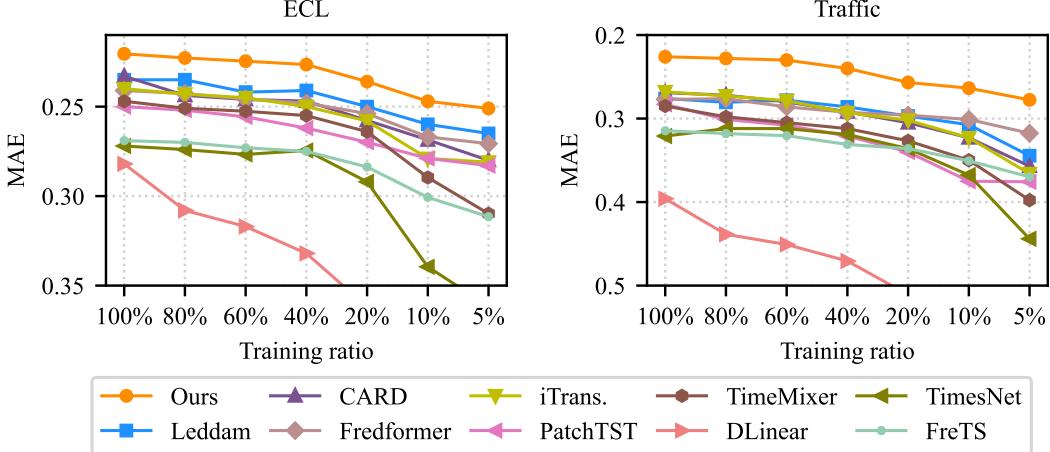


Figure 16: Performance change with less training data. The lookback and prediction lengths are set to 96. The Y-axis is inverted for clarity.

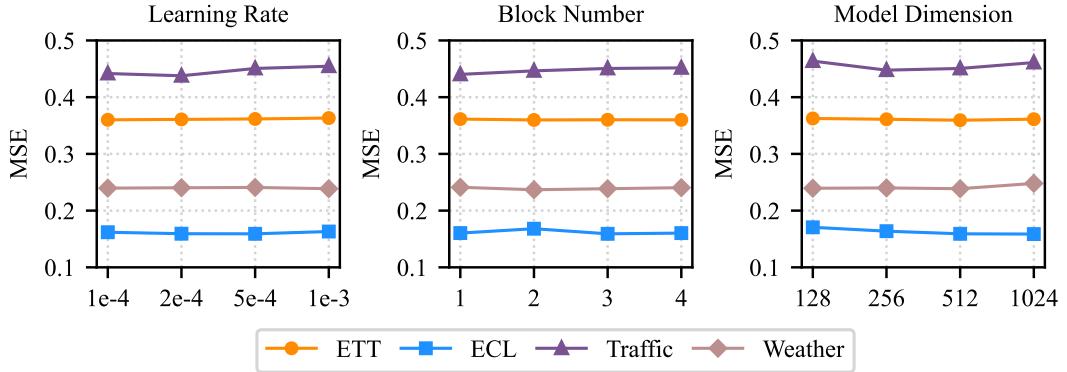


Figure 17: Hyperparameter sensitivity with respect to the learning rate lr , the number of blocks L , and the model dimension D . Average MSEs are reported across four prediction lengths $\tau \in \{96, 192, 336, 720\}$. ETT denotes the average performance over four subsets: ETT $\{h1, h2, m1, m2\}$. The lookback length T is uniformly set as 96.

J.2 Pre- and post-linear layers in NormLin module

In CSL, the pre- and post-linear layers, named based on their positions in the execution sequence, are applied around the NormLin layer (see Equation 4): Post_Lin (NormLin (Pre_Lin(\cdot))).

We conduct ablation studies to evaluate the contributions of these two layers. As shown in Table 38, incorporating both pre- and post-linear layers yields an average performance gain of 6% compared to the variant without them, highlighting their effectiveness in refining inputs for multivariate correlation modeling and downstream series representation learning.

J.3 Ablations on Q matrices

The effectiveness of OrthoTrans has been demonstrated in Tables 4 and 5. To further analyze the individual contributions of \mathbf{Q}_i and \mathbf{Q}_o , we conduct ablation studies on these transformation matrices.

As shown in Table 39, using both \mathbf{Q}_i and \mathbf{Q}_o achieves the best performance—reducing the average MSE by 5.3% compared to the baseline without transformations, by 4.4% compared to using \mathbf{Q}_o only, and by 0.8% compared to using \mathbf{Q}_i only. Notably, the variant using only \mathbf{Q}_i outperforms the one using only \mathbf{Q}_o , highlighting the critical role of temporal decorrelation in effective encoding and representation learning.

J.4 Swapping Q matrices

Table 40 reports the forecasting performance when both training and testing are conducted using Q matrices (\mathbf{Q}_i and \mathbf{Q}_o) derived from different datasets. As shown, using the Q matrices from the same dataset yields the best performance, highlighting the importance of incorporating dataset-specific temporal correlation information in time series forecasting.

K Model efficiency

As a linear model, OLinear achieves remarkable training and inference efficiency while delivering state-of-the-art forecasting performance. Figure 3 shows that OLinear achieves better efficiency than Transformer-based forecasters, benefiting from its relatively simple architecture. Detailed resource consumption for both training and inference is summarized in Table 41. Moreover, Table 42 presents the efficiency differences before and after applying the NormLin module to Transformer-based forecasters.

L Limitations

While this work includes extensive experiments, it primarily focuses on time series forecasting, and our claims are confined to this setting. Future work includes larger-scale pre-training and a wider range of time series analysis tasks (e.g., imputation and anomaly detection), in order to further validate the robustness of OLinear and its core components, OrthoTrans and NormLin. In addition, although NormLin is linear-based, it still incurs $\mathcal{O}(N^2)$ computational and memory complexity with respect to the number of variates N . Developing simpler yet effective alternatives for modeling multivariate correlations—and more generally, for learning general-purpose token dependencies—remains a promising direction.

M Broader impacts

Multivariate time series forecasting is a fundamental field with a wide range of real-world applications across domains such as energy, finance, healthcare, and transportation. While this work does not target a specific application domain, OLinear could have general real-world applicability (e.g., weather forecasting and traffic planning). We contribute two key insights: (1) the OrthoTrans scheme, which reformulates temporally correlated forecasting as inter-independent feature prediction; and (2) the linear-based NormLin module, which effectively models multivariate correlations, enhances representation learning among temporal and frequency tokens, and remains compatible with decoder architectures and large-scale pre-training. As a step toward simpler and more efficient token dependency learners, the NormLin module shows potential for broader use in deep learning and may inspire future research in this direction.

Table 21: Full results of applying OrthoTrans to iTransformer, PatchTST and RLinear. ‘Imp.’ denotes the improvement over the corresponding vanilla model. Note that RLinear consists of only a single linear layer that projects the lookback length T to the prediction horizon τ . Nevertheless, OrthoTrans consistently improves RLinear’s performance, demonstrating its ability to enhance model capacity. This table is the complete version of Table 5.

Model	iTransformer 2024a				PatchTST 2023				RLinear 2023				
	Vanilla		+OrthoTrans		Vanilla		+OrthoTrans		Vanilla		+OrthoTrans		
	Metric	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE
ETTm1	96	0.334	0.368	0.334	0.367	0.329	0.367	0.322	0.359	0.355	0.376	0.348	0.366
	192	0.377	0.391	0.377	0.389	0.367	0.385	0.363	0.383	0.391	0.392	0.386	0.385
	336	0.426	0.420	0.414	0.413	0.399	0.410	0.394	0.405	0.424	0.415	0.419	0.406
	720	0.491	0.459	0.490	0.452	0.454	0.439	0.457	0.445	0.487	0.450	0.478	0.439
	Avg	0.407	0.410	0.404	0.405	0.387	0.400	0.384	0.398	0.414	0.407	0.408	0.399
	Imp.	–	–	8.00%	1.16%	–	–	0.82%	0.53%	–	–	1.51%	1.97%
ECL	96	0.148	0.240	0.136	0.233	0.161	0.250	0.153	0.243	0.201	0.281	0.197	0.273
	192	0.162	0.253	0.155	0.250	0.199	0.289	0.165	0.254	0.201	0.283	0.196	0.276
	336	0.178	0.269	0.169	0.265	0.215	0.305	0.182	0.272	0.215	0.298	0.211	0.291
	720	0.225	0.317	0.225	0.311	0.256	0.337	0.222	0.306	0.257	0.331	0.253	0.324
	Avg	0.178	0.270	0.171	0.265	0.208	0.295	0.181	0.269	0.219	0.298	0.214	0.291
	Imp.	–	–	3.97%	1.87%	–	–	13.12%	8.98%	–	–	2.17%	2.35%
PEMS03	12	0.071	0.174	0.062	0.164	0.099	0.216	0.073	0.176	0.126	0.236	0.117	0.226
	24	0.093	0.201	0.080	0.185	0.142	0.259	0.104	0.210	0.246	0.334	0.233	0.322
	48	0.125	0.236	0.112	0.222	0.211	0.319	0.181	0.268	0.551	0.529	0.532	0.514
	96	0.164	0.275	0.160	0.269	0.269	0.370	0.293	0.338	1.057	0.787	1.024	0.760
	Avg	0.113	0.221	0.103	0.210	0.180	0.291	0.163	0.248	0.495	0.472	0.477	0.456
	Imp.	–	–	8.45%	5.02%	–	–	9.69%	14.81%	–	–	3.74%	3.50%
PEMS07	12	0.067	0.165	0.055	0.147	0.095	0.207	0.065	0.160	0.118	0.235	0.109	0.221
	24	0.088	0.190	0.071	0.165	0.150	0.262	0.097	0.192	0.242	0.341	0.228	0.324
	48	0.110	0.215	0.093	0.188	0.253	0.340	0.164	0.247	0.562	0.541	0.539	0.522
	96	0.139	0.245	0.122	0.218	0.346	0.404	0.264	0.311	1.096	0.795	1.065	0.773
	Avg	0.101	0.204	0.085	0.179	0.211	0.303	0.147	0.227	0.504	0.478	0.485	0.460
	Imp.	–	–	15.59%	11.94%	–	–	30.17%	24.94%	–	–	3.76%	3.77%
Solar-Energy	96	0.203	0.237	0.194	0.237	0.234	0.286	0.208	0.248	0.322	0.339	0.302	0.327
	192	0.233	0.261	0.231	0.265	0.267	0.310	0.237	0.268	0.359	0.356	0.342	0.346
	336	0.248	0.273	0.242	0.272	0.290	0.315	0.255	0.281	0.397	0.369	0.384	0.362
	720	0.249	0.275	0.245	0.275	0.289	0.317	0.255	0.281	0.397	0.356	0.387	0.356
	Avg	0.233	0.262	0.228	0.262	0.270	0.307	0.239	0.269	0.369	0.356	0.354	0.348
	Imp.	–	–	2.16%	0.00%	–	–	11.57%	12.22%	–	–	4.14%	2.37%
Weather	96	0.174	0.214	0.164	0.208	0.177	0.218	0.164	0.209	0.192	0.232	0.192	0.232
	192	0.221	0.254	0.221	0.259	0.225	0.259	0.211	0.251	0.240	0.271	0.236	0.268
	336	0.278	0.296	0.271	0.295	0.278	0.297	0.266	0.292	0.292	0.307	0.288	0.304
	720	0.358	0.349	0.351	0.347	0.354	0.348	0.345	0.343	0.364	0.353	0.359	0.350
	Avg	0.258	0.279	0.252	0.277	0.259	0.281	0.246	0.274	0.272	0.291	0.269	0.289
	Imp.	–	–	2.42%	0.63%	–	–	4.86%	2.59%	–	–	1.22%	0.86%
METR-LA	3	0.205	0.188	0.205	0.189	0.204	0.190	0.203	0.187	0.210	0.195	0.208	0.194
	6	0.300	0.229	0.292	0.227	0.298	0.227	0.296	0.228	0.304	0.234	0.303	0.235
	9	0.386	0.265	0.375	0.260	0.382	0.263	0.380	0.263	0.389	0.268	0.387	0.266
	12	0.460	0.295	0.444	0.291	0.456	0.292	0.452	0.294	0.465	0.298	0.464	0.297
	Avg	0.338	0.244	0.329	0.242	0.335	0.243	0.333	0.243	0.342	0.249	0.341	0.248
	Imp.	–	–	2.55%	1.02%	–	–	0.67%	0.00%	–	–	0.44%	0.28%

Table 22: Ablation results for replacing or removing components along the variate and temporal dimensions. ‘Hor.’ denotes the prediction horizon. The simplified results are presented in Table 6.

Var.	Temp.	Hor.	ECL		Traffic		Solar		PEMS03		Weather		ETTm1		NASDAQ (S1)		ILI (S2)	
			MSE	MAE														
NormLin	Linear	H1	0.131	0.221	0.398	0.226	0.179	0.191	0.060	0.159	0.153	0.190	0.302	0.334	0.036	0.092	1.737	0.800
		H2	0.150	0.238	0.439	0.241	0.209	0.213	0.078	0.179	0.200	0.235	0.357	0.363	0.049	0.117	1.714	0.795
		H3	0.165	0.254	0.464	0.250	0.231	0.229	0.104	0.210	0.258	0.280	0.387	0.385	0.062	0.137	1.821	0.804
		H4	0.191	0.279	0.502	0.270	0.241	0.236	0.140	0.247	0.337	0.333	0.452	0.426	0.073	0.154	1.785	0.810
	Linear	Avg	0.159	0.248	0.451	0.247	0.215	0.217	0.095	0.199	0.237	0.260	0.374	0.377	0.055	0.125	1.764	0.802
		H1	0.147	0.244	0.564	0.308	0.207	0.213	0.069	0.173	0.149	0.187	0.307	0.339	0.038	0.101	2.015	0.828
		H2	0.168	0.262	0.585	0.311	0.246	0.238	0.090	0.197	0.200	0.236	0.358	0.365	0.051	0.123	1.938	0.865
		H3	0.184	0.279	0.607	0.319	0.265	0.251	0.135	0.241	0.264	0.286	0.394	0.390	0.064	0.143	1.856	0.810
		H4	0.213	0.304	0.668	0.342	0.268	0.250	0.189	0.292	0.338	0.334	0.451	0.425	0.075	0.160	1.942	0.847
		Avg	0.178	0.272	0.606	0.320	0.246	0.238	0.121	0.226	0.238	0.261	0.377	0.380	0.057	0.132	1.938	0.837
w/o	Linear	H1	0.147	0.230	0.441	0.241	0.201	0.208	0.069	0.170	0.164	0.199	0.304	0.336	0.035	0.091	1.767	0.805
		H2	0.165	0.246	0.471	0.251	0.236	0.229	0.099	0.200	0.210	0.243	0.362	0.366	0.049	0.116	1.905	0.826
		H3	0.181	0.263	0.486	0.259	0.260	0.244	0.168	0.254	0.264	0.283	0.392	0.388	0.061	0.136	1.909	0.826
		H4	0.219	0.297	0.529	0.278	0.267	0.248	0.254	0.314	0.349	0.339	0.456	0.428	0.073	0.153	1.876	0.834
		Avg	0.178	0.259	0.482	0.257	0.241	0.232	0.147	0.234	0.247	0.266	0.378	0.379	0.054	0.124	1.864	0.823
	NormLin	H1	0.138	0.229	0.427	0.260	0.205	0.213	0.060	0.160	0.152	0.190	0.312	0.342	0.035	0.093	2.093	0.850
		H2	0.155	0.243	0.447	0.269	0.245	0.235	0.079	0.184	0.202	0.238	0.361	0.368	0.049	0.118	1.818	0.822
		H3	0.173	0.261	0.466	0.276	0.277	0.255	0.120	0.227	0.259	0.280	0.393	0.388	0.062	0.138	1.974	0.842
		H4	0.212	0.294	0.500	0.294	0.283	0.258	0.189	0.286	0.342	0.335	0.457	0.430	0.073	0.154	1.905	0.832
		Avg	0.169	0.257	0.460	0.275	0.252	0.240	0.112	0.214	0.239	0.261	0.381	0.382	0.055	0.126	1.947	0.836
Linear	NormLin	H1	0.156	0.252	0.557	0.329	0.214	0.226	0.072	0.178	0.155	0.191	0.313	0.345	0.038	0.101	2.135	0.885
		H2	0.171	0.264	0.555	0.333	0.255	0.250	0.100	0.209	0.202	0.238	0.364	0.369	0.051	0.122	1.986	0.865
		H3	0.187	0.282	0.567	0.336	0.287	0.270	0.154	0.264	0.263	0.283	0.396	0.392	0.064	0.144	1.968	0.848
		H4	0.220	0.306	0.632	0.359	0.294	0.272	0.245	0.334	0.340	0.334	0.461	0.431	0.077	0.164	2.059	0.870
		Avg	0.183	0.276	0.578	0.339	0.262	0.254	0.143	0.246	0.240	0.261	0.383	0.384	0.057	0.133	2.037	0.867
	w/o	H1	0.158	0.242	0.474	0.282	0.237	0.237	0.073	0.178	0.163	0.199	0.314	0.346	0.035	0.092	2.154	0.886
		H2	0.168	0.251	0.479	0.284	0.273	0.257	0.111	0.218	0.210	0.242	0.365	0.370	0.049	0.117	2.074	0.870
		H3	0.185	0.268	0.493	0.289	0.310	0.277	0.202	0.293	0.264	0.282	0.397	0.392	0.062	0.138	2.105	0.874
		H4	0.228	0.304	0.526	0.307	0.313	0.276	0.344	0.387	0.346	0.336	0.462	0.433	0.073	0.154	2.039	0.866
		Avg	0.185	0.266	0.493	0.290	0.283	0.262	0.182	0.269	0.246	0.265	0.384	0.385	0.055	0.125	2.093	0.874
NormLin	w/o	H1	0.138	0.228	0.427	0.260	0.207	0.214	0.060	0.160	0.151	0.189	0.311	0.342	0.036	0.093	2.117	0.852
		H2	0.155	0.243	0.446	0.269	0.246	0.236	0.079	0.184	0.202	0.238	0.360	0.367	0.049	0.118	1.767	0.812
		H3	0.171	0.260	0.467	0.277	0.280	0.257	0.121	0.227	0.260	0.282	0.392	0.388	0.063	0.139	1.960	0.846
		H4	0.214	0.296	0.500	0.295	0.281	0.258	0.195	0.291	0.343	0.336	0.459	0.431	0.073	0.155	1.916	0.837
		Avg	0.169	0.257	0.460	0.275	0.253	0.241	0.114	0.215	0.239	0.261	0.380	0.382	0.055	0.126	1.940	0.837
	Linear	H1	0.155	0.252	0.575	0.331	0.216	0.228	0.073	0.179	0.152	0.190	0.315	0.346	0.037	0.099	2.190	0.891
		H2	0.171	0.266	0.575	0.333	0.254	0.249	0.100	0.210	0.204	0.238	0.362	0.369	0.052	0.123	1.978	0.861
		H3	0.189	0.282	0.572	0.334	0.286	0.269	0.154	0.264	0.262	0.283	0.396	0.392	0.064	0.145	2.074	0.874
		H4	0.217	0.306	0.642	0.368	0.292	0.270	0.242	0.332	0.342	0.337	0.462	0.431	0.076	0.162	2.051	0.872
		Avg	0.183	0.276	0.591	0.341	0.262	0.254	0.142	0.246	0.240	0.262	0.384	0.384	0.057	0.132	2.073	0.874
Attn.	Linear	H1	0.138	0.226	0.407	0.228	0.185	0.194	0.060	0.160	0.155	0.192	0.317	0.346	0.036	0.094	2.404	0.888
		H2	0.157	0.245	0.445	0.243	0.218	0.218	0.075	0.179	0.206	0.240	0.369	0.375	0.050	0.118	1.979	0.857
		H3	0.180	0.267	0.461	0.252	0.235	0.234	0.105	0.212	0.267	0.288	0.416	0.403	0.063	0.138	1.895	0.827
		H4	0.190	0.281	0.516	0.282	0.240	0.237	0.149	0.256	0.350	0.339	0.465	0.433	0.074	0.155	1.811	0.818
		Avg	0.166	0.255	0.457	0.251	0.220	0.221	0.097	0.202	0.244	0.265	0.391	0.389	0.056	0.126	2.022	0.847

Table 23: Comparison of NormLin with the attention variants using the OLinear architecture. This is a complete version of Table 8.

Model	OLinear (Ours)	Trans. 2017	Reformer 2020	Flowformer 2022	Flashformer 2022	Flatten 2023	Mamba 2023	Lin.Attn. 2024	Enh.Attn. 2025
Metric									
ECL	96	0.131 0.221	0.138 0.226	0.134 0.223	0.135 0.225	0.138 0.227	0.135 0.225	0.142 0.231	0.134 0.222 0.131 0.221
	192	0.150 0.238	0.157 0.245	0.153 0.240	0.157 0.245	0.155 0.243	0.156 0.245	0.164 0.250	0.155 0.243
	336	0.165 0.254	0.180 0.267	0.169 0.257	0.171 0.261	0.177 0.266	0.172 0.262	0.180 0.267	0.174 0.261
	720	0.191 0.279 0.190	0.281 0.211	0.293 0.196	0.283 0.197	0.285 0.194	0.283 0.194	0.219 0.302	0.190 0.279 0.187 0.277
Avg									
Traffic	96	0.398 0.226	0.407 0.228	0.413 0.226	0.411 0.229	0.407 0.229	0.417 0.229	0.420 0.250	0.409 0.228 0.403
	192	0.439 0.241	0.445 0.243	0.440 0.241	0.446 0.246	0.441 0.243	0.441 0.242	0.445 0.263	0.441 0.424 0.243
	336	0.464 0.250 0.461 0.252	0.470 0.250	0.468 0.255	0.464 0.250	0.474 0.252	0.463 0.271	0.469 0.252 0.447 0.252	
	720	0.502 0.270	0.516 0.282	0.514 0.270	0.516 0.275	0.528 0.283	0.526 0.273	0.496 0.290	0.513 0.274
Avg									
PEMS03	12	0.060 0.159 0.060 0.160	0.060 0.159 0.060 0.160	0.061 0.160 0.060 0.160	0.061 0.160 0.060 0.160	0.062 0.161	0.062 0.162	0.061 0.160	
	24	0.078 0.179 0.075 0.179	0.077 0.178 0.076 0.179	0.075 0.178 0.075 0.178	0.077 0.181	0.078 0.181	0.078 0.182	0.078 0.181	
	48	0.104 0.210	0.105 0.212	0.103 0.208	0.109 0.214	0.104 0.210	0.112 0.216	0.115 0.220	0.111 0.215
	96	0.140 0.247	0.149 0.256	0.144 0.248	0.149 0.255	0.147 0.255	0.155 0.259	0.163 0.266	0.156 0.261
Avg									
Weather	96	0.153 0.190 0.155	0.192 0.153 0.190	0.156 0.191	0.157 0.194	0.156 0.193	0.155 0.194	0.193 0.156	
	192	0.200 0.235	0.206 0.240	0.204 0.237	0.209 0.242	0.205 0.240	0.211 0.245	0.205 0.241	
	336	0.258 0.280	0.267 0.288	0.267 0.285	0.264 0.284	0.264 0.285	0.268 0.288	0.263 0.285	
	720	0.337 0.333	0.350 0.339	0.340 0.332	0.341 0.336	0.341 0.337	0.349 0.340	0.345 0.339	0.361 0.347
Avg									
PEMS08	12	0.068 0.159	0.070 0.164	0.069 0.157 0.068	0.161 0.070	0.164 0.071	0.164 0.072	0.167 0.070	0.164 0.068 0.159
	24	0.089 0.178	0.095 0.187	0.087 0.175	0.090 0.181	0.096 0.188	0.094 0.185	0.095 0.190	0.094 0.184 0.084 0.178
	48	0.123 0.204	0.130 0.215	0.118 0.201	0.130 0.212	0.128 0.216	0.141 0.220	0.139 0.228	0.133 0.217
	96	0.173 0.236	0.207 0.265	0.174 0.234	0.195 0.252	0.206 0.263	0.213 0.260	0.226 0.276	0.192 0.255
Avg									
Solar-Energy	96	0.179 0.191	0.192 0.198	0.179 0.192	0.180 0.190	0.193 0.198	0.200 0.198	0.188 0.202	0.193 0.195
	192	0.209 0.213	0.220 0.219	0.211 0.214	0.218 0.220	0.215 0.215	0.230 0.222	0.221 0.224	0.230 0.223
	336	0.231 0.229	0.237 0.236	0.232 0.230	0.237 0.234	0.240 0.237	0.246 0.238	0.247 0.241	0.244 0.237
	720	0.241 0.236	0.243 0.240	0.243 0.236	0.252 0.242	0.246 0.241	0.251 0.242	0.257 0.245	0.249 0.242
Avg									
ILI	24	1.737 0.800	2.404 0.888	1.776 0.805	2.015 0.842	2.565 0.895	2.723 0.909	1.867 0.795	2.580 0.903
	36	1.714 0.795	1.979 0.857	1.769 0.806	1.820 0.832	1.973 0.856	1.944 0.864	2.002 0.846	1.832 0.835
	48	1.821 0.804	1.895 0.827	1.911 0.820	1.867 0.824	1.891 0.831	1.884 0.839	1.996 0.837	1.854 0.829
	60	1.785 0.810	1.811 0.818	1.827 0.814	1.823 0.815	1.837 0.816	1.987 0.872	1.935 0.835	1.972 0.868
Avg									
NASDAQ	3	0.036 0.092	0.036 0.094	0.035 0.092	0.036 0.093	0.036 0.094	0.094	0.037 0.096	0.036 0.093 0.036 0.093
	6	0.049 0.117	0.050 0.118	0.049 0.118	0.049 0.117	0.050 0.119	0.052 0.122	0.050 0.118	0.051 0.119
	9	0.062 0.137	0.063 0.138	0.062 0.137	0.062 0.137	0.063 0.138	0.064 0.140	0.062 0.138	0.062 0.139
	12	0.073 0.154	0.074 0.155	0.074 0.155	0.073 0.154	0.074 0.155	0.075 0.157	0.074 0.155	0.074 0.155
Avg									
1 st Count		27	26	3	0	10	21	6	5
		1	2	1	2	1	0	2	1
		1	0	2	1	1	1	0	15
		1	0	2	1	1	0	15	5

Table 24: Applying the NormLin module to Transformer-based forecasters: iTransformer, PatchTST, Leddam, and Fredformer. For a fair comparison, model hyperparameters, loss functions, and training strategies are kept unchanged. In the case of Leddam, only the *cross-channel attention* module is updated. This table presents the complete results corresponding to Table 9 in the main text.

Model	iTransformer 2024a				PatchTST 2023				Leddam 2024				Fredformer 2024				
	Vanilla	+ NormLin	Vanilla	+ NormLin	Vanilla	+ NormLin	Vanilla	+ NormLin	Vanilla	+ NormLin	Vanilla	+ NormLin	MSE	MAE	MSE	MAE	
ETTm1	96	0.334	0.368	0.320	0.359	0.329	0.367	0.323	0.361	0.319	0.359	0.319	0.357	0.326	0.361	0.318	0.356
	192	0.377	0.391	0.365	0.382	0.367	0.385	0.357	0.384	0.369	0.383	0.360	0.380	0.363	0.380	0.362	0.383
	336	0.426	0.420	0.400	0.405	0.399	0.410	0.390	0.408	0.394	0.402	0.391	0.401	0.395	0.403	0.391	0.404
	720	0.491	0.459	0.467	0.443	0.454	0.439	0.448	0.443	0.460	0.442	0.455	0.439	0.453	0.438	0.454	0.440
	Avg	0.407	0.410	0.388	0.397	0.387	0.400	0.379	0.399	0.386	0.397	0.381	0.394	0.384	0.396	0.381	0.396
ECL	96	0.148	0.240	0.138	0.232	0.161	0.250	0.155	0.246	0.141	0.235	0.139	0.234	0.147	0.241	0.138	0.235
	192	0.162	0.253	0.154	0.248	0.199	0.289	0.167	0.256	0.159	0.252	0.155	0.248	0.165	0.258	0.155	0.251
	336	0.178	0.269	0.171	0.265	0.215	0.305	0.183	0.273	0.173	0.268	0.171	0.265	0.177	0.273	0.172	0.271
	720	0.225	0.317	0.201	0.297	0.256	0.337	0.221	0.307	0.201	0.295	0.195	0.293	0.213	0.304	0.210	0.304
	Avg	0.178	0.270	0.166	0.260	0.208	0.295	0.181	0.270	0.169	0.263	0.165	0.260	0.176	0.269	0.169	0.265
PEMS03	12	0.071	0.174	0.062	0.164	0.099	0.216	0.071	0.176	0.063	0.164	0.063	0.164	0.068	0.174	0.064	0.168
	24	0.093	0.201	0.078	0.184	0.142	0.259	0.103	0.211	0.080	0.185	0.079	0.182	0.093	0.202	0.081	0.188
	48	0.125	0.236	0.110	0.219	0.211	0.319	0.160	0.248	0.124	0.226	0.110	0.217	0.146	0.258	0.119	0.228
	96	0.164	0.275	0.159	0.269	0.269	0.370	0.251	0.331	0.160	0.266	0.159	0.266	0.228	0.330	0.170	0.284
	Avg	0.113	0.221	0.102	0.209	0.180	0.291	0.146	0.241	0.107	0.210	0.103	0.207	0.134	0.241	0.108	0.217
PEMS07	12	0.067	0.165	0.055	0.145	0.095	0.207	0.084	0.199	0.055	0.145	0.054	0.145	0.063	0.158	0.057	0.150
	24	0.088	0.190	0.070	0.162	0.150	0.262	0.098	0.198	0.070	0.164	0.067	0.160	0.089	0.190	0.073	0.168
	48	0.110	0.215	0.091	0.187	0.253	0.340	0.243	0.341	0.094	0.192	0.087	0.184	0.135	0.239	0.099	0.201
	96	0.139	0.245	0.128	0.220	0.346	0.404	0.249	0.310	0.117	0.217	0.119	0.222	0.196	0.294	0.157	0.267
	Avg	0.101	0.204	0.086	0.178	0.211	0.303	0.168	0.262	0.084	0.180	0.082	0.178	0.121	0.220	0.096	0.196
Solar-Energy	96	0.203	0.237	0.194	0.233	0.234	0.286	0.208	0.249	0.197	0.241	0.186	0.227	0.185	0.233	0.190	0.231
	192	0.233	0.261	0.225	0.258	0.267	0.310	0.234	0.266	0.231	0.264	0.221	0.258	0.227	0.253	0.222	0.254
	336	0.248	0.273	0.241	0.273	0.290	0.315	0.252	0.278	0.241	0.268	0.242	0.270	0.246	0.284	0.245	0.275
	720	0.249	0.275	0.244	0.277	0.289	0.317	0.253	0.277	0.250	0.281	0.239	0.271	0.247	0.276	0.249	0.281
	Avg	0.233	0.262	0.226	0.260	0.270	0.307	0.237	0.267	0.230	0.264	0.222	0.257	0.226	0.262	0.226	0.260
Weather	96	0.174	0.214	0.160	0.205	0.177	0.218	0.161	0.206	0.156	0.202	0.156	0.201	0.163	0.207	0.156	0.202
	192	0.221	0.254	0.209	0.250	0.225	0.259	0.207	0.249	0.207	0.250	0.206	0.249	0.211	0.251	0.205	0.248
	336	0.278	0.296	0.268	0.294	0.278	0.297	0.266	0.291	0.262	0.291	0.262	0.291	0.267	0.292	0.261	0.289
	720	0.358	0.349	0.344	0.344	0.354	0.348	0.345	0.344	0.343	0.343	0.345	0.344	0.343	0.341	0.340	0.341
	Avg	0.258	0.279	0.245	0.273	0.259	0.281	0.245	0.272	0.242	0.272	0.242	0.271	0.246	0.273	0.240	0.270
METR-LA	3	0.205	0.188	0.202	0.189	0.204	0.190	0.205	0.188	0.204	0.191	0.202	0.189	0.205	0.188	0.200	0.186
	6	0.300	0.229	0.293	0.226	0.298	0.227	0.297	0.229	0.293	0.227	0.289	0.226	0.298	0.227	0.296	0.227
	9	0.386	0.265	0.375	0.259	0.382	0.263	0.379	0.262	0.369	0.264	0.366	0.259	0.385	0.263	0.369	0.263
	12	0.460	0.295	0.442	0.289	0.456	0.292	0.453	0.293	0.442	0.292	0.423	0.287	0.457	0.292	0.451	0.291
	Avg	0.338	0.244	0.328	0.241	0.335	0.243	0.333	0.243	0.327	0.243	0.320	0.240	0.336	0.242	0.329	0.242

Table 25: Full evaluation of OLinear-C. S_1 and S_2 correspond to ‘Input-12, Predict-{3, 6, 9, 12}’ and ‘Input-36, Predict-{24, 36, 48, 60}’, respectively. This table presents the full results of Table 7.

Dataset		ETTm1		ETTm2		ETTh1		ETTh2		ECL		Exchange		Traffic	
Metric		MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE
OLinear	96	0.302	0.334	0.169	0.249	0.360	0.382	0.284	0.329	0.131	0.221	0.082	0.200	0.398	0.226
	192	0.357	0.363	0.232	0.290	0.416	0.414	0.360	0.379	0.150	0.238	0.171	0.293	0.439	0.241
	336	0.387	0.385	0.291	0.328	0.457	0.438	0.409	0.415	0.165	0.254	0.331	0.414	0.464	0.250
	720	0.452	0.426	0.389	0.387	0.463	0.462	0.415	0.431	0.191	0.279	0.837	0.688	0.502	0.270
	Avg	0.374	0.377	0.270	0.313	0.424	0.424	0.367	0.388	0.159	0.248	0.355	0.399	0.451	0.247
OLinear-C	96	0.303	0.335	0.169	0.249	0.362	0.383	0.285	0.329	0.130	0.220	0.081	0.198	0.404	0.227
	192	0.357	0.364	0.232	0.290	0.415	0.414	0.358	0.377	0.156	0.243	0.171	0.293	0.435	0.241
	336	0.389	0.386	0.291	0.328	0.460	0.440	0.407	0.415	0.167	0.256	0.325	0.410	0.465	0.250
	720	0.452	0.426	0.389	0.387	0.461	0.460	0.423	0.436	0.192	0.279	0.845	0.692	0.502	0.271
	Avg	0.375	0.378	0.270	0.313	0.424	0.424	0.368	0.389	0.161	0.249	0.355	0.398	0.451	0.247
Dataset		Weather		Solar-Energy		PEMS03		PEMS04		PEMS07		PEMS08		ILI (S1)	
Metric		MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE
OLinear	H1	0.153	0.190	0.179	0.191	0.060	0.159	0.068	0.163	0.052	0.138	0.068	0.159	0.468	0.349
	H2	0.200	0.235	0.209	0.213	0.078	0.179	0.079	0.176	0.065	0.151	0.089	0.178	0.923	0.516
	H3	0.258	0.280	0.231	0.229	0.104	0.210	0.095	0.197	0.084	0.171	0.123	0.204	1.289	0.655
	H4	0.337	0.333	0.241	0.236	0.140	0.247	0.122	0.226	0.108	0.196	0.173	0.236	1.698	0.791
	Avg	0.237	0.260	0.215	0.217	0.095	0.199	0.091	0.190	0.077	0.164	0.113	0.194	1.094	0.578
OLinear-C	H1	0.152	0.189	0.178	0.191	0.061	0.160	0.068	0.163	0.052	0.138	0.068	0.159	0.485	0.352
	H2	0.202	0.237	0.209	0.213	0.076	0.178	0.078	0.176	0.064	0.151	0.089	0.177	0.951	0.523
	H3	0.258	0.279	0.231	0.229	0.105	0.211	0.097	0.197	0.083	0.171	0.120	0.203	1.336	0.667
	H4	0.340	0.333	0.241	0.236	0.142	0.248	0.121	0.227	0.108	0.195	0.177	0.237	1.693	0.790
	Avg	0.238	0.259	0.215	0.217	0.096	0.199	0.091	0.191	0.077	0.164	0.113	0.194	1.116	0.583
Dataset		ILI (S2)		COVID-19 (S1)		COVID-19 (S2)		METR-LA (S1)		METR-LA (S2)		NASDAQ (S1)		NASDAQ (S2)	
Metric		MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE
OLinear	H1	1.737	0.800	1.100	0.487	4.474	1.180	0.207	0.171	0.650	0.337	0.036	0.092	0.121	0.216
	H2	1.714	0.795	1.750	0.619	7.241	1.670	0.301	0.207	0.800	0.388	0.049	0.117	0.163	0.261
	H3	1.821	0.804	2.239	0.734	10.076	1.985	0.382	0.238	0.905	0.427	0.062	0.137	0.205	0.296
	H4	1.785	0.810	2.538	0.831	12.079	2.182	0.452	0.263	0.999	0.457	0.073	0.154	0.259	0.336
	Avg	1.764	0.802	1.907	0.668	8.467	1.754	0.335	0.220	0.838	0.402	0.055	0.125	0.187	0.277
OLinear-C	H1	1.850	0.835	1.107	0.493	4.980	1.336	0.207	0.172	0.654	0.337	0.035	0.092	0.119	0.214
	H2	1.749	0.801	1.693	0.630	7.605	1.719	0.301	0.207	0.804	0.388	0.049	0.117	0.162	0.259
	H3	1.847	0.808	2.251	0.740	10.336	2.015	0.382	0.238	0.916	0.428	0.061	0.137	0.207	0.298
	H4	1.792	0.808	2.545	0.843	12.251	2.204	0.450	0.263	1.016	0.462	0.073	0.153	0.255	0.334
	Avg	1.809	0.813	1.899	0.676	8.793	1.819	0.335	0.220	0.847	0.404	0.054	0.125	0.186	0.276
Dataset		Wiki (S1)		Wiki (S2)		SP500 (S1)		SP500 (S2)		DowJones (S1)		DowJones (S2)		CarSales (S1)	
Metric		MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE
OLinear	H1	6.161	0.368	6.894	0.423	0.035	0.126	0.155	0.271	1.550	0.276	7.432	0.664	0.303	0.277
	H2	6.453	0.385	6.446	0.439	0.053	0.158	0.209	0.317	2.569	0.361	10.848	0.799	0.315	0.285
	H3	6.666	0.398	6.004	0.446	0.070	0.181	0.258	0.358	3.565	0.430	14.045	0.914	0.327	0.293
	H4	6.834	0.406	5.705	0.454	0.088	0.204	0.305	0.387	4.517	0.490	16.959	1.017	0.336	0.301
	Avg	6.528	0.389	6.262	0.440	0.061	0.167	0.231	0.333	3.050	0.389	12.321	0.848	0.320	0.289
OLinear-C	H1	6.187	0.368	6.888	0.422	0.035	0.126	0.156	0.271	1.553	0.276	7.496	0.665	0.303	0.276
	H2	6.457	0.384	6.437	0.438	0.054	0.159	0.210	0.319	2.570	0.361	10.965	0.802	0.317	0.287
	H3	6.652	0.396	6.003	0.446	0.071	0.182	0.262	0.358	3.573	0.431	14.161	0.915	0.328	0.294
	H4	6.835	0.406	5.717	0.456	0.088	0.205	0.296	0.378	4.538	0.490	17.084	1.018	0.337	0.300
	Avg	6.533	0.388	6.261	0.440	0.062	0.168	0.231	0.331	3.058	0.389	12.427	0.850	0.321	0.289
Dataset		CarSales (S2)		Power (S1)		Power (S2)		Website (S1)		Website (S2)		Unemp (S1)		Unemp (S2)	
Metric		MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE
OLinear	H1	0.320	0.302	0.864	0.688	1.343	0.870	0.077	0.199	0.186	0.306	0.012	0.047	0.458	0.447
	H2	0.334	0.315	0.991	0.742	1.445	0.903	0.103	0.227	0.272	0.356	0.041	0.108	0.870	0.619
	H3	0.347	0.327	1.062	0.770	1.559	0.946	0.135	0.251	0.365	0.391	0.084	0.170	1.643	0.900
	H4	0.358	0.337	1.119	0.789	1.602	0.971	0.176	0.281	0.486	0.481	0.131	0.220	2.593	1.182
	Avg	0.340	0.320	1.009	0.747	1.487	0.922	0.123	0.240	0.327	0.383	0.067	0.136	1.391	0.787
OLinear-C	H1	0.320	0.301	0.873	0.696	1.460	0.922	0.092	0.222	0.176	0.304	0.012	0.047	0.437	0.434
	H2	0.334	0.315	0.997	0.744	1.527	0.939	0.110	0.232	0.283	0.364	0.039	0.105	0.899	0.633
	H3	0.347	0.327	1.060	0.768	1.721	1.016	0.144	0.257	0.402	0.421	0.079	0.163	1.963	1.026
	H4	0.358	0.337	1.128	0.796	1.869	1.077	0.188	0.286	0.517	0.490	0.132	0.222	2.571	1.162
	Avg	0.340	0.320	1.014	0.751	1.644	0.988	0.133	0.249	0.344	0.395	0.065	0.134	1.467	0.814

Table 26: Robustness of OLinear-C performance. Standard deviations are calculated over seven random seeds.

Dataset		ECL		Traffic		ETTm1		Solar-Energy	
Metric		MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE
Horizon	96	0.130±3e-4	0.220±4e-4	0.404±4e-3	0.227±7e-4	0.303±7e-4	0.335±7e-4	0.178±1e-3	0.191±7e-4
	192	0.156±1e-3	0.243±9e-4	0.435±2e-3	0.241±4e-4	0.357±2e-4	0.364±2e-4	0.209±8e-4	0.213±3e-4
	336	0.167±1e-3	0.256±8e-4	0.465±5e-3	0.250±2e-4	0.389±9e-4	0.386±5e-4	0.231±5e-4	0.229±4e-5
	720	0.192±9e-3	0.279±6e-3	0.502±5e-3	0.271±6e-4	0.452±6e-4	0.426±4e-4	0.241±4e-4	0.236±4e-4
Dataset		Weather		PEMS03		NASDAQ (S1)		Wiki (S1)	
Metric		MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE
Horizon	H1	0.154±1e-3	0.191±8e-4	0.061±4e-4	0.160±4e-4	0.035±2e-4	0.092±5e-4	6.187±9e-3	0.368±7e-4
	H2	0.204±3e-3	0.238±2e-3	0.076±8e-4	0.178±8e-4	0.049±8e-5	0.117±4e-4	6.457±7e-3	0.384±6e-4
	H3	0.259±4e-3	0.280±4e-3	0.105±9e-4	0.211±6e-4	0.061±0.000	0.137±3e-4	6.652±4e-3	0.396±6e-4
	H4	0.343±6e-3	0.334±3e-3	0.142±2e-3	0.248±1e-3	0.073±5e-5	0.153±2e-4	6.835±3e-3	0.406±2e-4
Dataset		DowJones (S2)		SP500 (S2)		CarSales (S1)		Power (S2)	
Metric		MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE
Horizon	H1	7.496±4e-2	0.665±9e-4	0.156±2e-3	0.271±2e-3	0.303±1e-3	0.276±1e-3	1.460±4e-2	0.922±1e-2
	H2	10.965±7e-2	0.802±1e-3	0.210±1e-3	0.319±1e-3	0.317±7e-4	0.287±6e-4	1.527±3e-2	0.939±2e-2
	H3	14.161±5e-2	0.915±7e-4	0.262±2e-3	0.358±2e-3	0.328±7e-4	0.294±6e-4	1.721±7e-2	1.016±3e-2
	H4	17.084±7e-2	1.018±1e-3	0.296±3e-3	0.378±3e-3	0.337±4e-4	0.300±4e-4	1.869±1e-2	1.077±5e-3

Table 27: Ablation study on different transformations of CorrMat_v

Dataset		ECL		Traffic		Solar		PEMS03		Weather		NASDAQ S1		ILI S2		Web_S1		Cars_S2	
Metric		MSE	MAE																
Softmax	H1	0.130	0.220	0.404	0.227	0.178	0.191	0.061	0.160	0.152	0.189	0.035	0.092	1.850	0.835	0.078	0.200	0.320	0.302
	H2	0.156	0.243	0.435	0.241	0.209	0.213	0.076	0.178	0.202	0.237	0.049	0.117	1.749	0.801	0.108	0.230	0.334	0.315
	H3	0.167	0.256	0.465	0.250	0.231	0.229	0.105	0.211	0.258	0.279	0.061	0.137	1.847	0.808	0.143	0.256	0.347	0.327
	H4	0.192	0.279	0.502	0.271	0.241	0.236	0.142	0.248	0.340	0.333	0.073	0.153	1.792	0.808	0.185	0.285	0.358	0.337
Avg		0.161	0.249	0.451	0.247	0.215	0.217	0.096	0.199	0.238	0.259	0.054	0.125	1.809	0.813	0.128	0.243	0.340	0.320
Softplus	H1	0.131	0.221	0.401	0.227	0.178	0.191	0.061	0.160	0.154	0.191	0.035	0.092	1.854	0.838	0.078	0.200	0.320	0.302
	H2	0.153	0.241	0.436	0.241	0.209	0.213	0.078	0.180	0.204	0.240	0.049	0.117	1.741	0.801	0.108	0.230	0.334	0.315
	H3	0.166	0.255	0.461	0.249	0.231	0.229	0.106	0.211	0.262	0.281	0.061	0.137	1.843	0.807	0.143	0.256	0.347	0.327
	H4	0.191	0.280	0.511	0.271	0.242	0.236	0.140	0.247	0.342	0.336	0.073	0.154	1.793	0.808	0.185	0.285	0.358	0.337
Avg		0.160	0.249	0.452	0.247	0.215	0.217	0.096	0.199	0.240	0.262	0.054	0.125	1.808	0.813	0.128	0.243	0.340	0.320
Sigmoid	H1	0.131	0.221	0.398	0.227	0.178	0.192	0.060	0.159	0.152	0.189	0.036	0.093	1.727	0.796	0.078	0.200	0.320	0.302
	H2	0.152	0.240	0.441	0.241	0.213	0.214	0.077	0.179	0.206	0.241	0.049	0.117	1.908	0.832	0.108	0.230	0.334	0.315
	H3	0.166	0.256	0.460	0.250	0.230	0.229	0.104	0.210	0.266	0.285	0.061	0.137	1.845	0.807	0.143	0.256	0.347	0.327
	H4	0.191	0.280	0.509	0.270	0.241	0.236	0.140	0.247	0.345	0.337	0.073	0.154	1.795	0.809	0.185	0.285	0.358	0.337
Avg		0.160	0.249	0.452	0.247	0.216	0.218	0.095	0.199	0.242	0.263	0.055	0.125	1.819	0.811	0.128	0.243	0.340	0.320
ReLU	H1	0.132	0.222	0.403	0.227	0.179	0.192	0.060	0.159	0.154	0.190	0.035	0.092	1.844	0.832	0.078	0.201	0.320	0.302
	H2	0.153	0.241	0.446	0.241	0.213	0.214	0.076	0.178	0.205	0.241	0.049	0.116	1.757	0.802	0.108	0.229	0.334	0.315
	H3	0.171	0.259	0.464	0.250	0.230	0.229	0.106	0.211	0.258	0.280	0.061	0.136	1.841	0.807	0.143	0.256	0.347	0.327
	H4	0.193	0.282	0.519	0.271	0.242	0.236	0.141	0.248	0.343	0.335	0.073	0.153	1.790	0.808	0.185	0.285	0.358	0.337
Avg		0.162	0.251	0.458	0.247	0.216	0.218	0.096	0.199	0.240	0.261	0.054	0.124	1.808	0.812	0.128	0.243	0.340	0.320
Identity	H1	0.132	0.222	0.402	0.226	0.179	0.192	0.060	0.159	0.150	0.188	0.035	0.092	1.843	0.832	0.078	0.201	0.320	0.301
	H2	0.152	0.240	0.445	0.241	0.213	0.214	0.076	0.178	0.204	0.240	0.049	0.116	1.757	0.802	0.108	0.229	0.334	0.315
	H3	0.170	0.259	0.464	0.250	0.230	0.229	0.106	0.211	0.262	0.282	0.061	0.136	1.842	0.807	0.144	0.257	0.347	0.327
	H4	0.194	0.284	0.520	0.271	0.242	0.236	0.141	0.248	0.341	0.334	0.073	0.153	1.791	0.808	0.188	0.285	0.358	0.337
Avg		0.162	0.251	0.458	0.247	0.216	0.218	0.096	0.199	0.239	0.261	0.054	0.124	1.808	0.812	0.129	0.243	0.340	0.320

Table 28: Temporal decorrelation comparison of OrthoTrans, DFT and Wavelet transforms. We report the off-diagonal Frobenius norm of the temporal correlation matrix, divided by the number of off-diagonal elements.

Dataset	Training Set				Test Set			
	Original	OrthoTrans	DFT	Wavelet	Original	OrthoTrans	DFT	Wavelet
ETTh1	6.7E-03	1.6E-04	5.8E-04	3.7E-03	5.2E-03	2.1E-04	6.0E-04	3.1E-03
ETTm2	9.1E-03	1.4E-04	1.1E-03	4.6E-03	7.2E-03	2.2E-04	1.3E-03	3.7E-03
ECL	5.0E-03	1.9E-04	4.7E-04	3.6E-03	5.0E-03	2.2E-04	4.9E-04	3.6E-03
Traffic	3.8E-03	1.9E-04	5.5E-04	2.6E-03	3.8E-03	2.3E-04	5.7E-04	2.6E-03
Weather	6.7E-03	3.1E-04	4.1E-03	3.6E-03	7.2E-03	4.0E-04	4.6E-03	3.7E-03
Avg	6.3E-03	2.0E-04	1.4E-03	3.6E-03	5.7E-03	2.6E-04	1.5E-03	3.3E-03

Table 29: Performance with varying lookback horizons. The best and second-best results are highlighted in **bold** and underlined, respectively. *OLinear-C* refers to the OLinear variant where the weight matrix in NormLin is replaced by Softmax (CorrMat_v), with CorrMat_v being the cross-variate correlation matrix.

Model	OLinear (Ours)	OLinear-C (Ours)	Leddam 2024	CARD 2024b	Fredformer 2024	iTrans. 2024a	TimeMixer 2024a	PatchTST 2023	TimesNet 2023b	DLinear 2023									
Metric MSE MAE																			
ECL	96 0.123 <u>0.212</u> 0.123 <u>0.212</u> 0.134 0.227 <u>0.129</u> 0.223 <u>0.129</u> 0.226 0.132 0.227 <u>0.129</u> 0.224 <u>0.129</u> 0.222 0.168 0.272 0.140 0.237	192 <u>0.143</u> <u>0.232</u> 0.144 0.233 0.156 0.248 0.154 0.245 0.148 0.244 0.154 0.251 0.140 0.220 0.147 0.240 0.184 0.322 0.153 0.249	336 <u>0.157</u> 0.247 <u>0.156</u> 0.247 0.166 0.264 0.161 0.257 0.165 0.262 0.170 0.268 0.161 <u>0.255</u> 0.163 0.259 0.198 0.300 0.169 0.267	720 <u>0.181</u> 0.270 <u>0.180</u> 0.270 0.195 0.291 0.185 <u>0.278</u> 0.193 0.286 0.193 0.288 0.194 0.287 0.197 0.290 0.220 0.320 0.203 0.301	Avg 0.151 0.240 <u>0.151</u> 0.240 0.163 0.257 0.157 0.251 0.159 0.254 0.162 0.258 <u>0.156</u> 0.247 0.159 0.253 0.193 0.304 0.166 0.264														
Traffic	96 0.338 <u>0.221</u> <u>0.340</u> 0.221 0.366 0.260 0.341 <u>0.229</u> 0.358 0.257 0.359 0.262 0.360 0.249 0.360 0.249 0.593 0.321 0.410 0.282	192 0.361 <u>0.233</u> 0.369 <u>0.235</u> 0.394 0.270 <u>0.367</u> 0.243 0.381 0.272 0.376 0.270 0.375 0.250 0.379 0.256 0.617 0.336 0.423 0.287	336 <u>0.386</u> 0.241 0.390 <u>0.241</u> 0.400 0.283 0.388 <u>0.254</u> 0.396 0.277 0.393 0.279 0.385 0.270 0.392 0.264 0.629 0.336 0.436 0.296	720 0.444 0.264 0.451 <u>0.266</u> 0.442 0.297 <u>0.427</u> 0.276 0.424 0.296 0.434 0.293 0.430 0.281 0.432 0.286 0.640 0.350 0.466 0.315	Avg <u>0.382</u> 0.240 0.387 <u>0.241</u> 0.400 0.278 0.381 0.251 0.390 0.275 0.390 0.276 0.388 0.263 0.391 0.264 0.620 0.336 0.434 0.295														
Weather	96 0.144 <u>0.184</u> 0.144 <u>0.183</u> 0.149 0.199 <u>0.145</u> 0.186 0.150 0.203 0.165 0.214 0.147 0.197 0.149 0.198 0.172 0.220 0.176 0.237	192 0.190 0.230 <u>0.189</u> <u>0.228</u> 0.196 0.243 0.187 <u>0.227</u> 0.194 0.246 0.208 0.253 <u>0.189</u> 0.239 0.194 0.241 0.219 0.261 0.220 0.282	336 0.235 <u>0.268</u> <u>0.236</u> <u>0.268</u> 0.243 0.280 0.238 <u>0.258</u> 0.243 0.284 0.257 0.292 0.241 0.280 0.306 0.282 0.246 0.337 0.265 0.319	720 0.316 <u>0.323</u> 0.316 <u>0.323</u> 0.321 0.334 0.308 <u>0.321</u> 0.308 0.333 0.331 0.343 <u>0.310</u> 0.330 0.314 0.334 0.365 0.359 0.323 0.362	Avg <u>0.221</u> 0.251 <u>0.221</u> <u>0.250</u> 0.227 0.264 0.220 <u>0.248</u> 0.224 0.266 0.240 0.275 0.222 0.262 0.241 0.264 0.251 0.294 0.246 0.300														
Solar-Energy	96 <u>0.169</u> 0.194 0.195 <u>0.207</u> 0.186 0.242 0.170 <u>0.207</u> 0.187 0.236 0.190 0.241 0.167 0.220 0.224 0.278 0.219 0.314 0.289 0.377	192 0.190 0.212 <u>0.189</u> 0.212 0.208 0.262 0.192 <u>0.219</u> 0.196 0.251 0.233 0.261 0.187 0.249 0.253 0.298 0.231 0.322 0.319 0.397	336 0.192 <u>0.215</u> <u>0.194</u> <u>0.219</u> 0.218 0.265 0.226 0.233 0.208 0.265 0.226 0.275 0.200 0.258 0.273 0.306 0.246 0.337 0.352 0.415	720 <u>0.204</u> 0.225 <u>0.203</u> <u>0.226</u> 0.208 0.273 0.217 0.243 0.209 0.272 0.220 0.282 0.215 0.250 0.272 0.308 0.280 0.363 0.356 0.412	Avg 0.189 <u>0.211</u> 0.195 <u>0.216</u> 0.205 0.261 0.201 0.225 0.200 0.256 0.217 0.265 <u>0.192</u> 0.244 0.256 0.298 0.244 0.334 0.329 0.400														
ETTm1	96 0.275 <u>0.326</u> 0.275 <u>0.326</u> 0.294 0.347 0.288 <u>0.332</u> <u>0.284</u> 0.338 0.309 0.357 0.291 0.340 0.293 0.346 0.338 0.375 0.299 0.343	192 <u>0.318</u> <u>0.352</u> 0.317 <u>0.351</u> 0.336 0.369 0.332 0.357 0.323 0.364 0.346 0.383 0.327 0.365 0.333 0.370 0.374 0.387 0.335 0.365	336 0.352 <u>0.373</u> <u>0.352</u> <u>0.373</u> 0.364 0.389 0.364 <u>0.376</u> <u>0.358</u> 0.387 0.385 0.410 0.360 0.381 0.369 0.392 0.410 0.411 0.369 0.386	720 0.406 <u>0.408</u> 0.406 <u>0.408</u> 0.421 0.419 <u>0.414</u> 0.407 0.420 0.417 0.440 0.442 0.415 0.417 0.416 0.420 0.478 0.450 0.425 0.421	Avg <u>0.338</u> <u>0.365</u> 0.338 <u>0.364</u> 0.354 0.381 0.350 0.368 <u>0.346</u> 0.376 0.370 0.398 0.348 0.376 0.353 0.382 0.400 0.406 0.357 0.379														
1 st Count	12	16	11	12	0	0	4	4	2	0	0	0	4	1	0	0	0	0	0

Table 30: Comparison of OLinear with more baselines. MAEs are reported.

Model	OLinear (Ours)	SimpleTM 2025	TQNet 2025	TimePro 2025g	TimeBase 2025c
ETTm1	96 0.334	0.361	0.353	0.364	0.388
	192 0.363	0.380	0.378	0.383	0.409
	336 0.385	0.404	0.401	0.409	0.421
	720 0.426	0.438	0.440	0.446	0.461
	Avg 0.377	0.396	0.393	0.400	0.420
ETTm2	96 0.249	0.257	0.256	0.260	0.271
	192 0.290	0.299	0.298	0.303	0.309
	336 0.328	0.338	0.340	0.342	0.346
	720 0.387	0.395	0.396	0.399	0.401
	Avg 0.313	0.322	0.323	0.326	0.332
ETTh1	96 0.382	0.392	0.393	0.398	0.392
	192 0.414	0.421	0.426	0.429	0.423
	336 0.438	0.438	0.446	0.450	0.443
	720 0.462	0.462	0.470	0.474	0.458
	Avg 0.424	0.428	0.434	0.438	0.429
ETTh2	96 0.329	0.338	0.343	0.345	0.376
	192 0.379	0.387	0.393	0.394	0.405
	336 0.415	0.401	0.427	0.431	0.440
	720 0.431	0.436	0.446	0.445	0.477
	Avg 0.388	0.391	0.402	0.403	0.424
ECL	96 0.221	0.235	0.229	0.234	0.279
	192 0.238	0.247	0.247	0.249	0.281
	336 0.254	0.267	0.264	0.267	0.295
	720 0.279	0.293	0.294	0.299	0.327
	Avg 0.248	0.260	0.259	0.262	0.295
Traffic	96 0.226	0.274	0.261	0.269	0.384
	192 0.241	0.280	0.271	0.276	0.362
	336 0.250	0.290	0.277	0.287	0.365
	720 0.270	0.309	0.295	0.312	0.386
	Avg 0.247	0.289	0.276	0.286	0.374
Weather	96 0.190	0.207	0.200	0.207	0.215
	192 0.235	0.248	0.245	0.254	0.256
	336 0.280	0.290	0.287	0.296	0.297
	720 0.333	0.341	0.342	0.346	0.348
	Avg 0.260	0.271	0.269	0.276	0.279
Solar-Energy	96 0.191	0.232	0.233	0.237	0.363
	192 0.213	0.247	0.257	0.263	0.404
	336 0.229	0.257	0.263	0.281	0.398
	720 0.236	0.252	0.270	0.285	0.388
	Avg 0.217	0.247	0.256	0.266	0.388

Table 31: Few-shot forecasting performance with 10% training data. Results are averaged over four prediction lengths {96, 192, 336, 720}.

Model	OLinear (Ours)	TimeMixer++ 2025a	TimeMixer 2024a	iTrans. 2024a	PatchTST 2023	TimesNet 2023b	DLinear 2023	TiDE 2023	Crossfm. 2023	FEDformer 2022				
Metric	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE
ETTh1	0.478	0.456	0.517	0.512	0.613	0.520	0.510	0.597	0.633	0.542	0.869	0.628	0.691	0.600
ETTh2	0.390	0.407	0.379	0.391	0.402	0.433	0.455	0.461	0.415	0.431	0.479	0.465	0.605	0.538
ETTm1	0.405	0.400	0.398	0.431	0.487	0.461	0.491	0.516	0.501	0.466	0.677	0.537	0.411	0.429
ETTm2	0.287	0.330	0.291	0.351	0.311	0.367	0.375	0.412	0.296	0.343	0.320	0.353	0.316	0.368
Weather	0.245	0.265	0.241	0.271	0.281	0.291	0.331	0.242	0.279	0.279	0.301	0.241	0.283	0.249
	0.291	0.271	0.242	0.279	0.301	0.241	0.283	0.249	0.291	0.267	0.306	0.284	0.324	

Table 32: Zero-shot forecasting performance. Results are averaged over four prediction lengths {96, 192, 336, 720}. $\mathcal{A} \rightarrow \mathcal{B}$ denotes that the model is trained on dataset \mathcal{A} but evaluated on \mathcal{B} .

Model	OLinear (Ours)	TimeMixer++ 2025a	TimeMixer 2024a	LLMTime 2023	DLinear 2023	PatchTST 2023	TimesNet 2023b	iTrans. 2024a	Crossfm. 2023	FEDformer 2022				
Metric	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE
ETTh1 → ETTh2	0.367	0.391	0.367	0.391	0.427	0.424	0.992	0.708	0.493	0.488	0.380	0.405	0.421	0.431
ETTh1 → ETTm2	0.311	0.352	0.301	0.357	0.361	0.397	1.867	0.869	0.415	0.452	0.314	0.360	0.327	0.361
ETTh2 → ETTh1	0.507	0.479	0.511	0.498	0.679	0.577	1.961	0.981	0.703	0.574	0.565	0.513	0.865	0.621
ETTm1 → ETTh2	0.413	0.425	0.417	0.422	0.452	0.441	0.992	0.708	0.464	0.475	0.439	0.438	0.457	0.454
ETTm1 → ETTm2	0.291	0.327	0.291	0.331	0.357	1.867	0.869	0.335	0.389	0.296	0.334	0.322	0.354	0.324
ETTm2 → ETTm1	0.480	0.445	0.427	0.448	0.554	0.478	1.933	0.984	0.649	0.537	0.568	0.492	0.769	0.567
	0.559	0.491	0.611	0.593	0.577	0.601								

Table 33: Performance comparison of OLinear and large time series models on fine-tuning and zero-shot scenarios. MSEs are reported, with a prediction length of 96. Better results are highlighted in **bold**. For OLinear’s zero-shot learning, we employ a cross-dataset strategy: training on ETTh1 and evaluating on ETTh2 (and vice versa), and similarly for ETTm1 and ETTm2. For Time-MoE [Shi et al., 2025], Moirai [Woo et al., 2024], and Chronos [Ansari et al., 2024], we use the largest model variants from their respective families. Baseline results are sourced from [Liu et al., 2024b] and [Liu et al., 2025].

Scenario	100% Target				20% Target				5% Target				0% Target (Zero-shot)			
	Model	OLinear (Ours)	Timer 2024b	OLinear (Ours)	Timer 2024b	OLinear (Ours)	Timer 2024b	OLinear (Ours)	Timer-28B 2024b	Time-MoE 2025	Moirai 2024	Moment 2024	TimesFM 2024	Chronos 2024		
Pre-trained	No	UTSD	No	UTSD	No	UTSD	ETT	UTSD+ LOTSA	Time-300B	LOTSA	TS Pile	O(100B) 13 Public Datasets				
ETTh1	0.363	0.358	0.396	0.359	0.413	0.362	0.416	0.393	0.349	0.381	0.674	0.414	0.441			
ETTh2	0.276	–	0.289	0.284	0.303	0.280	0.287	0.308	0.292	0.296	0.330	0.315	0.320			
ETTm1	0.275	–	0.300	0.321	0.321	0.321	0.355	0.420	0.281	0.380	0.670	0.361	0.457			
ETTm2	0.161	–	0.167	0.187	0.170	0.176	0.176	0.247	0.198	0.211	0.257	0.202	0.197			
ECL	0.123	0.136	0.131	0.134	0.153	0.132	–	0.147	–	0.153	0.744	–	–			
Traffic	0.338	0.351	0.371	0.352	0.419	0.361	–	0.414	–	0.399	1.293	–	–			
Weather	0.144	0.154	0.156	0.151	0.168	0.151	–	0.243	0.157	0.199	0.255	–	–			
PEMS03	0.103	0.118	0.116	0.116	0.226	0.125	–	–	–	–	–	–	–			
PEMS04	0.086	0.107	0.195	0.120	0.542	0.135	–	–	–	–	–	–	–			

Table 34: Performance with various embedding size d settings. In our main experiments, we set $d = 16$ without specific statements.

Emb. size		1		4		8		16		32		64	
	Metric	MSE	MAE										
ECL	96	0.133	0.223	<u>0.132</u>	0.221	<u>0.132</u>	0.221	0.131	0.221	0.131	0.221	0.131	0.221
	192	0.152	<u>0.240</u>	0.153	0.241	<u>0.152</u>	0.240	0.150	0.238	0.153	0.241	<u>0.152</u>	0.240
	336	<u>0.166</u>	0.256	0.169	0.257	0.165	<u>0.255</u>	0.165	0.254	0.165	<u>0.255</u>	0.168	0.258
	720	0.188	0.277	0.198	0.282	0.207	0.290	<u>0.191</u>	<u>0.279</u>	0.193	0.280	0.200	0.286
	Avg	0.159	<u>0.249</u>	0.163	0.250	0.164	0.252	0.159	0.248	0.161	<u>0.249</u>	0.163	0.251
Traffic	96	0.398	0.230	0.411	<u>0.227</u>	0.403	0.226	0.398	0.226	<u>0.402</u>	0.226	0.404	0.226
	192	0.425	0.243	<u>0.433</u>	0.241	0.443	0.241	0.439	0.241	0.436	0.241	0.435	<u>0.242</u>
	336	0.450	0.251	0.462	<u>0.250</u>	0.465	0.249	0.464	<u>0.250</u>	0.460	0.251	<u>0.459</u>	0.254
	720	0.496	0.269	0.518	0.269	0.520	<u>0.270</u>	<u>0.502</u>	<u>0.270</u>	0.507	0.274	0.515	0.275
	Avg	0.442	0.248	0.456	<u>0.247</u>	0.458	0.246	<u>0.451</u>	<u>0.247</u>	<u>0.451</u>	0.248	0.453	0.249
Solar-Energy	96	0.180	0.194	<u>0.179</u>	0.193	0.180	<u>0.191</u>	0.179	0.191	0.177	0.190	0.182	0.192
	192	0.212	0.216	<u>0.211</u>	0.214	0.209	0.213	0.209	0.213	0.212	0.213	<u>0.211</u>	<u>0.214</u>
	336	0.234	0.233	0.232	<u>0.230</u>	0.230	<u>0.230</u>	0.231	0.229	<u>0.231</u>	<u>0.230</u>	0.233	0.231
	720	0.244	0.238	0.241	<u>0.236</u>	0.241	0.235	0.241	<u>0.236</u>	<u>0.242</u>	<u>0.236</u>	0.243	0.238
	Avg	0.217	0.220	<u>0.216</u>	0.218	0.215	0.217	0.215	0.217	0.215	0.217	0.217	0.219
Weather	96	<u>0.152</u>	0.189	0.158	0.192	0.150	0.187	<u>0.152</u>	0.189	<u>0.152</u>	0.190	0.157	0.194
	192	0.203	<u>0.238</u>	0.200	0.234	0.207	0.242	0.207	0.243	<u>0.202</u>	<u>0.238</u>	0.205	0.239
	336	<u>0.261</u>	<u>0.282</u>	0.259	0.280	0.259	0.280	0.261	0.282	0.263	<u>0.282</u>	0.269	0.287
	720	0.342	<u>0.337</u>	<u>0.343</u>	0.338	<u>0.343</u>	<u>0.337</u>	0.344	<u>0.337</u>	0.350	0.339	0.344	0.336
	Avg	0.240	0.261	0.240	0.261	0.240	0.261	<u>0.241</u>	0.262	0.242	<u>0.262</u>	0.244	0.264
ILI	3	0.487	0.355	0.545	0.380	0.500	0.353	0.468	0.349	0.515	0.362	<u>0.486</u>	<u>0.352</u>
	6	0.915	<u>0.515</u>	1.021	0.533	0.890	0.506	0.923	0.516	<u>0.914</u>	<u>0.515</u>	1.059	0.558
	9	1.301	0.662	1.363	0.666	1.376	0.667	1.289	0.655	<u>1.265</u>	0.644	1.264	<u>0.645</u>
	12	<u>1.638</u>	0.783	1.693	0.788	1.810	0.799	1.698	0.791	1.619	0.769	1.747	<u>0.773</u>
	Avg	<u>1.085</u>	0.579	1.155	0.592	1.144	0.581	1.094	<u>0.578</u>	1.078	0.572	1.139	0.582
PEMS03	12	0.060	0.159	0.060	0.159	0.060	0.159	0.060	0.159	0.061	0.160	0.059	0.158
	24	<u>0.077</u>	0.179	<u>0.078</u>	0.179	<u>0.078</u>	0.180	<u>0.078</u>	0.179	0.077	0.179	0.077	0.179
	48	0.107	0.213	0.106	<u>0.211</u>	0.106	<u>0.211</u>	0.104	0.210	0.106	<u>0.211</u>	<u>0.105</u>	0.210
	96	0.146	0.253	0.143	<u>0.248</u>	<u>0.141</u>	0.248	0.140	0.247	0.140	<u>0.247</u>	0.143	<u>0.248</u>
	Avg	0.097	0.201	0.097	0.199	<u>0.096</u>	0.200	0.095	0.199	<u>0.096</u>	0.199	<u>0.096</u>	0.199
1 st Count		9	4	4	8	9	12	12	14	8	12	4	7

Table 35: Forecasting performance on the scale-free metrics: Coefficient of Determination (R^2), Pearson Correlation Coefficient (r), and MASE. The symbols ↓ and ↑ indicate that lower and higher values are better, respectively. The best and second-best results are highlighted in **bold** and underlined, respectively. The lookback length T is uniformly set as 96.

Model		OLinear (Ours)			Leddam 2024			CARD 2024b			iTrans. 2024a			TimeMixer 2024a		
		R^2	r (↑)	MASE (↓)	R^2	r (↑)	MASE (↓)	R^2	r (↑)	MASE (↓)	R^2	r (↑)	MASE (↓)	R^2	r (↑)	MASE (↓)
ECL	96	0.640	0.916	0.902	0.527	<u>0.911</u>	0.964	0.533	0.905	0.975	<u>0.550</u>	0.907	0.990	0.456	0.897	1.022
	192	0.638	0.908	0.983	<u>0.593</u>	<u>0.902</u>	1.041	0.547	0.899	<u>1.028</u>	0.590	0.900	1.057	0.575	0.893	1.076
	336	0.710	0.901	1.061	0.670	<u>0.893</u>	1.134	<u>0.678</u>	0.891	<u>1.113</u>	<u>0.678</u>	<u>0.893</u>	1.133	0.640	0.886	1.155
	720	0.701	0.889	1.174	<u>0.661</u>	<u>0.882</u>	<u>1.251</u>	0.653	0.877	1.273	0.638	0.872	1.315	0.634	0.871	1.314
	Avg	0.672	0.903	1.030	0.613	<u>0.897</u>	1.098	0.603	0.893	<u>1.097</u>	<u>0.614</u>	0.893	1.124	0.576	0.887	1.142
Traffic	96	0.737	0.899	0.729	0.681	0.885	0.909	<u>0.690</u>	0.882	<u>0.855</u>	<u>0.690</u>	<u>0.889</u>	0.860	0.640	0.888	0.979
	192	0.719	0.884	0.756	0.666	0.865	0.949	0.683	0.869	<u>0.856</u>	<u>0.689</u>	<u>0.877</u>	0.879	0.660	0.862	0.959
	336	0.730	0.875	0.758	0.707	<u>0.867</u>	0.897	0.702	0.862	<u>0.843</u>	<u>0.713</u>	<u>0.867</u>	0.883	0.685	0.852	0.960
	720	0.703	0.860	0.815	0.687	0.853	0.947	0.679	0.848	<u>0.894</u>	<u>0.693</u>	<u>0.857</u>	0.913	0.672	0.842	0.986
	Avg	0.722	0.879	0.764	0.685	0.868	0.926	0.689	0.865	<u>0.862</u>	<u>0.696</u>	<u>0.873</u>	0.884	0.664	0.861	0.971
Model		FilterNet 2024a			FITS 2024			DLinear 2023			PatchTST 2023			TimesNet 2023b		
		R^2	r (↑)	MASE (↓)	R^2	r (↑)	MASE (↓)	R^2	r (↑)	MASE (↓)	R^2	r (↑)	MASE (↓)	R^2	r (↑)	MASE (↓)
ECL	96	0.537	0.905	0.996	0.352	0.883	1.144	0.120	0.870	1.300	0.426	0.890	1.119	0.405	0.892	1.118
	192	0.576	0.898	1.052	0.446	0.880	1.153	0.215	0.867	1.312	0.509	0.887	1.140	0.503	0.885	1.195
	336	0.670	0.891	1.139	0.588	0.873	1.229	0.247	0.858	1.387	0.604	0.880	1.216	0.592	0.875	1.262
	720	0.628	0.875	1.328	0.593	0.858	1.376	0.148	0.844	1.536	0.605	0.864	1.371	0.603	0.864	1.388
	Avg	0.603	0.892	1.129	0.495	0.873	1.225	0.182	0.860	1.384	0.536	0.880	1.211	0.526	0.879	1.241
Traffic	96	0.637	0.876	0.987	0.458	0.811	1.284	0.441	0.811	1.346	0.668	0.880	0.903	0.625	0.877	1.023
	192	0.665	0.867	0.950	0.542	0.814	1.151	0.551	0.815	1.189	0.677	0.871	0.887	0.635	0.867	1.009
	336	0.696	0.861	0.930	0.582	0.810	1.119	0.582	0.810	1.167	0.700	0.863	0.874	0.654	0.852	1.079
	720	0.678	0.847	0.982	0.568	0.794	1.169	0.564	0.794	1.224	0.681	0.850	0.925	0.661	0.850	1.047
	Avg	0.669	0.862	0.962	0.537	0.807	1.181	0.535	0.808	1.232	0.682	0.866	0.897	0.644	0.862	1.040

Table 36: Forecasting performance when orthogonal matrices are computed from a subset of the training data ('Q_ratio'), while model training still uses the full trainset.

Q_ratio		100%		80%		60%		40%		20%		10%	
Metric		MSE	MAE										
ECL	96	0.131	0.221	0.131	0.221	0.131	0.221	0.131	0.221	0.130	0.221	0.131	0.221
	192	0.150	0.238	0.155	0.243	0.152	0.239	0.153	0.240	0.151	0.240	0.152	0.240
	336	0.165	0.254	0.169	0.257	0.164	0.253	0.166	0.255	0.165	0.254	0.166	0.255
	720	0.191	0.279	0.193	0.284	0.188	0.277	0.193	0.284	0.193	0.280	0.189	0.278
	Avg	0.159	0.248	0.162	0.251	0.159	0.247	0.161	0.250	0.160	0.248	0.159	0.248
Traffic	96	0.398	0.226	0.403	0.226	0.403	0.227	0.401	0.226	0.404	0.226	0.402	0.226
	192	0.439	0.241	0.436	0.240	0.434	0.240	0.432	0.240	0.434	0.240	0.431	0.240
	336	0.464	0.250	0.461	0.250	0.462	0.250	0.463	0.250	0.459	0.250	0.465	0.250
	720	0.502	0.270	0.499	0.271	0.499	0.271	0.501	0.271	0.502	0.271	0.512	0.271
	Avg	0.451	0.247	0.450	0.247	0.449	0.247	0.449	0.247	0.450	0.247	0.453	0.247
Weather	96	0.153	0.190	0.150	0.188	0.153	0.191	0.153	0.190	0.154	0.191	0.149	0.187
	192	0.200	0.235	0.202	0.237	0.204	0.239	0.203	0.238	0.202	0.237	0.203	0.239
	336	0.258	0.280	0.260	0.280	0.260	0.281	0.264	0.284	0.260	0.281	0.259	0.281
	720	0.337	0.333	0.339	0.333	0.338	0.332	0.337	0.332	0.335	0.330	0.339	0.332
	Avg	0.237	0.260	0.238	0.259	0.239	0.261	0.239	0.261	0.238	0.260	0.237	0.260
PEMS03	12	0.060	0.159	0.060	0.159	0.060	0.159	0.060	0.159	0.061	0.160	0.060	0.159
	24	0.078	0.179	0.077	0.179	0.077	0.180	0.077	0.179	0.077	0.180	0.077	0.180
	48	0.104	0.210	0.105	0.210	0.104	0.210	0.104	0.210	0.104	0.210	0.105	0.210
	96	0.140	0.247	0.141	0.248	0.142	0.248	0.141	0.248	0.143	0.249	0.140	0.249
	Avg	0.095	0.199	0.096	0.199	0.096	0.199	0.095	0.199	0.096	0.199	0.095	0.199
NASDAQ	3	0.036	0.092	0.036	0.092	0.036	0.092	0.036	0.093	0.036	0.093	0.036	0.093
	6	0.049	0.117	0.049	0.117	0.049	0.117	0.049	0.117	0.049	0.117	0.049	0.118
	9	0.062	0.137	0.062	0.137	0.062	0.137	0.062	0.137	0.062	0.137	0.062	0.137
	12	0.073	0.154	0.073	0.154	0.073	0.154	0.073	0.154	0.073	0.154	0.073	0.154
	Avg	0.055	0.125										
CarSales	24	0.320	0.302	0.319	0.302	0.319	0.301	0.319	0.302	0.320	0.302	0.320	0.302
	36	0.334	0.315	0.334	0.315	0.334	0.315	0.334	0.315	0.334	0.315	0.334	0.316
	48	0.347	0.327	0.347	0.327	0.347	0.327	0.347	0.327	0.347	0.327	0.347	0.328
	60	0.358	0.337	0.358	0.337	0.358	0.337	0.358	0.337	0.358	0.337	0.358	0.337
	Avg	0.340	0.320	0.339	0.320	0.339	0.320	0.340	0.320	0.340	0.320	0.340	0.320

Table 37: Ablation study on transformation functions and normalization methods in NormLin. For $S1$ and $S2$, the prediction lengths are $\{3, 6, 9, 12\}$ and $\{24, 36, 48, 60\}$, respectively.

Fun.	Norm	Hor.	ECL	Traffic	Solar	PEMS03	Weather	NASDAQ (S1)	ILI (S2)							
Softplus	L1	H1	0.131	0.221	0.398	0.226	0.179	0.191	0.060	0.159	0.153	0.190	0.036	0.092	1.737	0.800
		H2	0.150	0.238	0.439	0.241	0.209	0.213	0.078	0.179	0.200	0.235	0.049	0.117	1.714	0.795
		H3	0.165	0.254	0.464	0.250	0.231	0.229	0.104	0.210	0.258	0.280	0.062	0.137	1.821	0.804
		H4	0.191	0.279	0.502	0.270	0.241	0.236	0.140	0.247	0.337	0.333	0.073	0.154	1.785	0.810
		Avg	0.159	0.248	0.451	0.247	0.215	0.217	0.095	0.199	0.237	0.260	0.055	0.125	1.764	0.802
Identity	L1	H1	0.140	0.229	0.455	0.233	0.182	0.193	0.062	0.161	0.154	0.191	0.036	0.093	1.783	0.807
		H2	0.157	0.245	0.486	0.246	0.216	0.215	0.078	0.179	0.203	0.240	0.049	0.117	1.844	0.810
		H3	0.170	0.260	0.506	0.255	0.239	0.233	0.103	0.208	0.267	0.285	0.061	0.137	1.909	0.826
		H4	0.200	0.287	0.551	0.276	0.506	0.255	0.143	0.250	0.350	0.339	0.073	0.154	1.841	0.830
		Avg	0.167	0.255	0.499	0.252	0.285	0.224	0.096	0.199	0.243	0.264	0.055	0.125	1.844	0.818
Softmax	L1	H1	0.131	0.221	0.398	0.226	0.179	0.192	0.061	0.160	0.150	0.187	0.036	0.092	1.707	0.800
		H2	0.155	0.242	0.441	0.241	0.209	0.213	0.077	0.179	0.202	0.237	0.049	0.117	1.886	0.828
		H3	0.167	0.255	0.467	0.250	0.232	0.229	0.107	0.211	0.259	0.280	0.062	0.138	1.820	0.803
		H4	0.199	0.286	0.504	0.270	0.241	0.236	0.139	0.247	0.335	0.332	0.073	0.154	1.795	0.814
		Avg	0.163	0.251	0.452	0.247	0.215	0.217	0.096	0.199	0.236	0.259	0.055	0.125	1.802	0.811
Sigmoid	L1	H1	0.131	0.221	0.398	0.226	0.178	0.191	0.061	0.159	0.152	0.189	0.035	0.093	1.742	0.799
		H2	0.150	0.238	0.439	0.241	0.209	0.213	0.077	0.179	0.203	0.238	0.049	0.117	1.719	0.797
		H3	0.167	0.256	0.464	0.250	0.231	0.229	0.106	0.211	0.262	0.283	0.062	0.137	1.826	0.804
		H4	0.199	0.285	0.500	0.270	0.241	0.236	0.139	0.248	0.338	0.332	0.073	0.154	1.780	0.808
		Avg	0.161	0.250	0.450	0.247	0.215	0.217	0.096	0.199	0.239	0.260	0.055	0.125	1.767	0.802
ReLU	L1	H1	0.132	0.222	0.398	0.226	0.179	0.191	0.061	0.160	0.153	0.189	0.036	0.092	1.711	0.799
		H2	0.153	0.241	0.443	0.241	0.210	0.213	0.078	0.179	0.202	0.238	0.049	0.117	1.721	0.795
		H3	0.170	0.259	0.467	0.250	0.231	0.229	0.106	0.211	0.264	0.285	0.062	0.138	1.816	0.804
		H4	0.197	0.285	0.507	0.271	0.241	0.236	0.143	0.248	0.338	0.335	0.073	0.154	1.790	0.812
		Avg	0.163	0.251	0.454	0.247	0.215	0.217	0.097	0.199	0.239	0.262	0.055	0.125	1.759	0.802
Softplus	L2	H1	0.192	0.285	0.613	0.313	0.234	0.244	0.095	0.199	0.151	0.188	0.036	0.093	1.928	0.852
		H2	0.194	0.288	0.640	0.325	0.280	0.267	0.121	0.221	0.203	0.238	0.050	0.119	1.929	0.837
		H3	0.202	0.297	0.653	0.329	0.323	0.288	0.174	0.265	0.259	0.281	0.063	0.139	1.845	0.806
		H4	0.236	0.322	0.683	0.344	0.337	0.292	0.264	0.331	0.339	0.335	0.073	0.154	1.794	0.813
		Avg	0.206	0.298	0.647	0.328	0.293	0.273	0.163	0.254	0.238	0.260	0.055	0.126	1.874	0.827
Identity	L2	H1	0.139	0.230	0.456	0.238	0.187	0.198	0.063	0.163	0.151	0.189	0.036	0.092	1.718	0.795
		H2	0.159	0.249	0.516	0.251	0.221	0.219	0.080	0.184	0.201	0.238	0.050	0.118	1.855	0.812
		H3	0.171	0.261	0.529	0.261	0.239	0.234	0.112	0.217	0.266	0.284	0.062	0.137	1.947	0.831
		H4	0.215	0.298	0.573	0.288	0.248	0.239	0.153	0.258	0.347	0.339	0.073	0.154	1.832	0.829
		Avg	0.171	0.259	0.518	0.259	0.224	0.222	0.102	0.205	0.241	0.262	0.055	0.125	1.838	0.817
Softmax	L2	H1	0.183	0.278	0.612	0.312	0.233	0.242	0.096	0.199	0.153	0.190	0.036	0.093	1.876	0.837
		H2	0.192	0.286	0.641	0.325	0.266	0.256	0.118	0.219	0.199	0.235	0.050	0.119	1.875	0.825
		H3	0.202	0.297	0.653	0.330	0.313	0.280	0.174	0.265	0.259	0.282	0.063	0.139	1.823	0.803
		H4	0.226	0.316	0.686	0.346	0.336	0.292	0.268	0.334	0.337	0.333	0.073	0.154	1.798	0.814
		Avg	0.201	0.294	0.648	0.328	0.287	0.267	0.164	0.254	0.237	0.260	0.055	0.126	1.843	0.820
Sigmoid	L2	H1	0.193	0.285	0.613	0.313	0.234	0.244	0.096	0.200	0.149	0.187	0.036	0.094	1.993	0.861
		H2	0.206	0.296	0.640	0.324	0.279	0.267	0.119	0.220	0.204	0.239	0.050	0.119	1.762	0.807
		H3	0.201	0.296	0.653	0.329	0.324	0.288	0.175	0.267	0.262	0.284	0.063	0.139	1.866	0.811
		H4	0.256	0.334	0.684	0.345	0.336	0.291	0.265	0.332	0.337	0.333	0.073	0.155	1.797	0.813
		Avg	0.214	0.302	0.647	0.328	0.293	0.272	0.164	0.255	0.238	0.261	0.056	0.127	1.854	0.823
ReLU	L2	H1	0.181	0.277	0.612	0.312	0.231	0.241	0.095	0.197	0.153	0.190	0.036	0.093	1.860	0.838
		H2	0.190	0.284	0.640	0.325	0.265	0.255	0.118	0.219	0.201	0.237	0.050	0.119	1.900	0.833
		H3	0.200	0.295	0.653	0.329	0.309	0.278	0.176	0.268	0.260	0.282	0.063	0.139	1.803	0.800
		H4	0.232	0.320	0.686	0.346	0.335	0.291	0.263	0.330	0.336	0.332	0.073	0.154	1.798	0.813
		Avg	0.201	0.294	0.648	0.328	0.285	0.266	0.163	0.254	0.237	0.260	0.055	0.126	1.840	0.821

Table 38: Ablation study on pre- and post-linear layers in the NormLin module. ‘Hor.’ denotes prediction horizons. For $S1$ and $S2$, the prediction horizons are $\{3, 6, 9, 12\}$ and $\{24, 36, 48, 60\}$, respectively.

PreLin	PostLin	Hor.	ECL	Traffic	Solar	PEMS03	Weather	NASDAQ (S1)	ILI (S2)	Wiki (S1)	ETTm1			
			MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE		
\times	\times	H1	0.135	0.229	0.474	0.253	0.192	0.206	0.065	0.167	0.159	0.195	0.040	
		H2	0.153	0.246	0.492	0.264	0.225	0.228	0.085	0.191	0.207	0.241	0.053	
		H3	0.168	0.264	0.511	0.273	0.247	0.244	0.125	0.233	0.264	0.285	0.065	
		H4	0.204	0.297	0.555	0.294	0.255	0.247	0.178	0.281	0.342	0.336	0.077	
			Avg	<u>0.165</u>	0.259	0.508	0.271	0.230	0.231	0.113	0.218	<u>0.243</u>	<u>0.264</u>	<u>0.058</u>
\times	\checkmark	H1	0.131	0.222	0.403	0.228	0.179	0.192	0.061	0.161	0.155	0.192	0.036	
		H2	0.151	0.240	0.434	0.242	0.210	0.213	0.077	0.180	0.205	0.240	0.049	
		H3	0.166	0.256	0.457	0.251	0.231	0.230	0.107	0.212	0.262	0.282	0.062	
		H4	0.211	0.298	0.500	0.272	0.243	0.236	0.146	0.253	0.352	0.340	0.073	
			Avg	<u>0.165</u>	0.254	<u>0.448</u>	<u>0.248</u>	<u>0.216</u>	<u>0.218</u>	0.098	<u>0.201</u>	<u>0.243</u>	<u>0.263</u>	<u>0.055</u>
\checkmark	\times	H1	0.131	0.222	0.404	0.228	0.178	0.192	0.061	0.160	0.154	0.192	0.036	
		H2	0.155	0.243	0.434	0.242	0.211	0.214	0.077	0.180	0.204	0.239	0.049	
		H3	0.165	0.255	0.457	0.251	0.231	0.230	0.107	0.212	0.263	0.282	0.062	
		H4	0.209	0.293	0.499	0.272	0.242	0.236	0.144	0.252	0.344	0.339	0.073	
			Avg	<u>0.165</u>	<u>0.253</u>	<u>0.448</u>	<u>0.248</u>	<u>0.215</u>	<u>0.218</u>	<u>0.097</u>	<u>0.201</u>	<u>0.241</u>	<u>0.263</u>	<u>0.055</u>
\checkmark	\checkmark	H1	0.131	0.221	0.398	0.226	0.179	0.191	0.060	0.159	0.152	0.189	0.036	
		H2	0.150	0.238	0.439	0.241	0.209	0.213	0.078	0.179	0.207	0.243	0.049	
		H3	0.165	0.254	0.464	0.250	0.231	0.229	0.104	0.210	0.261	0.282	0.062	
		H4	0.191	0.279	0.502	0.270	0.241	0.236	0.140	0.247	0.344	0.337	0.073	
			Avg	<u>0.159</u>	<u>0.248</u>	<u>0.451</u>	<u>0.247</u>	<u>0.215</u>	<u>0.217</u>	<u>0.095</u>	<u>0.199</u>	<u>0.241</u>	<u>0.263</u>	<u>0.055</u>

Table 39: Ablation study on Q_i and Q_o . ‘Hor.’ denotes prediction horizons. For $S1$ and $S2$, the prediction lengths are $\{3, 6, 9, 12\}$ and $\{24, 36, 48, 60\}$, respectively.

Q_i	Q_o	Hor.	ECL	Solar-Energy	PEMS03	PEMS08	NASDAQ (S1)	ILI (S2)	COVID-19 (S2)	Exchange			
			MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE			
\times	\times	H1	0.134	0.224	<u>0.186</u>	0.199	0.063	0.163	0.070	<u>0.162</u>	<u>0.036</u>	<u>0.093</u>	
		H2	0.156	0.245	0.221	0.220	0.080	0.184	0.091	0.183	<u>0.050</u>	<u>0.119</u>	
		H3	<u>0.167</u>	<u>0.257</u>	0.244	<u>0.237</u>	0.117	0.222	0.128	<u>0.216</u>	<u>0.063</u>	<u>0.139</u>	
		H4	0.194	<u>0.282</u>	0.256	0.244	0.165	0.266	0.231	0.262	<u>0.074</u>	<u>0.156</u>	
			Avg	0.163	0.252	0.227	<u>0.225</u>	0.106	0.209	0.130	<u>0.206</u>	<u>0.056</u>	<u>0.127</u>
\times	\checkmark	H1	<u>0.132</u>	<u>0.223</u>	<u>0.186</u>	<u>0.198</u>	0.062	<u>0.162</u>	<u>0.069</u>	<u>0.162</u>	<u>0.036</u>	<u>0.093</u>	
		H2	<u>0.153</u>	0.243	0.219	0.219	0.079	0.183	0.092	0.183	<u>0.050</u>	<u>0.119</u>	
		H3	<u>0.167</u>	0.258	0.246	0.238	0.115	0.221	0.132	0.217	<u>0.063</u>	<u>0.139</u>	
		H4	0.194	0.285	<u>0.255</u>	0.244	0.164	0.267	0.216	0.261	<u>0.074</u>	<u>0.156</u>	
			Avg	<u>0.161</u>	0.252	<u>0.226</u>	<u>0.225</u>	0.105	<u>0.208</u>	0.127	<u>0.206</u>	<u>0.056</u>	<u>0.127</u>
\checkmark	\times	H1	<u>0.131</u>	<u>0.221</u>	<u>0.179</u>	<u>0.191</u>	0.061	<u>0.159</u>	<u>0.068</u>	<u>0.159</u>	<u>0.036</u>	<u>0.092</u>	
		H2	<u>0.153</u>	0.243	0.219	0.219	0.079	0.183	0.092	0.183	<u>0.050</u>	<u>0.119</u>	
		H3	<u>0.167</u>	0.258	0.246	0.238	0.115	0.221	0.132	0.217	<u>0.063</u>	<u>0.139</u>	
		H4	0.194	0.285	<u>0.255</u>	0.244	0.164	0.267	0.216	0.261	<u>0.074</u>	<u>0.156</u>	
			Avg	<u>0.161</u>	0.252	<u>0.226</u>	<u>0.225</u>	0.105	<u>0.208</u>	0.127	<u>0.206</u>	<u>0.056</u>	<u>0.127</u>
\checkmark	\checkmark	H1	<u>0.131</u>	<u>0.221</u>	<u>0.179</u>	<u>0.191</u>	0.061	<u>0.159</u>	<u>0.068</u>	<u>0.159</u>	<u>0.036</u>	<u>0.092</u>	
		H2	<u>0.154</u>	<u>0.241</u>	<u>0.210</u>	<u>0.212</u>	0.076	<u>0.178</u>	<u>0.087</u>	<u>0.177</u>	<u>0.049</u>	<u>0.117</u>	
		H3	<u>0.165</u>	<u>0.254</u>	<u>0.230</u>	<u>0.229</u>	0.105	<u>0.211</u>	<u>0.121</u>	<u>0.204</u>	<u>0.062</u>	<u>0.137</u>	
		H4	<u>0.188</u>	<u>0.279</u>	<u>0.241</u>	<u>0.235</u>	0.142	<u>0.249</u>	<u>0.182</u>	<u>0.238</u>	<u>0.073</u>	<u>0.154</u>	
			Avg	<u>0.159</u>	<u>0.249</u>	<u>0.215</u>	<u>0.217</u>	0.096	<u>0.199</u>	<u>0.114</u>	<u>0.194</u>	<u>0.055</u>	<u>0.125</u>
\checkmark	\checkmark	H1	<u>0.131</u>	<u>0.221</u>	<u>0.179</u>	<u>0.191</u>	0.060	<u>0.159</u>	<u>0.068</u>	<u>0.159</u>	<u>0.036</u>	<u>0.092</u>	
		H2	<u>0.150</u>	<u>0.238</u>	<u>0.209</u>	<u>0.213</u>	0.078	<u>0.179</u>	<u>0.089</u>	<u>0.178</u>	<u>0.049</u>	<u>0.117</u>	
		H3	<u>0.165</u>	<u>0.254</u>	<u>0.231</u>	<u>0.229</u>	0.104	<u>0.210</u>	<u>0.123</u>	<u>0.204</u>	<u>0.062</u>	<u>0.137</u>	
		H4	<u>0.191</u>	<u>0.279</u>	<u>0.241</u>	<u>0.236</u>	0.140	<u>0.247</u>	<u>0.173</u>	<u>0.236</u>	<u>0.073</u>	<u>0.154</u>	
			Avg	<u>0.159</u>	<u>0.248</u>	<u>0.215</u>	<u>0.217</u>	0.095	<u>0.199</u>	<u>0.113</u>	<u>0.194</u>	<u>0.055</u>	<u>0.125</u>

Table 40: Forecasting performance when both training and testing use Q matrices \mathbf{Q}_i and \mathbf{Q}_o from different datasets. $\mathcal{A} \rightarrow \mathcal{B}$ indicates that the Q matrices computed from dataset \mathcal{A} are applied in experiments on dataset \mathcal{B} .

Setting	ILI→ILI		Power→ILI		Web→ILI	
Metric	MSE	MAE	MSE	MAE	MSE	MAE
24	1.737	0.800	1.846	0.830	1.887	0.842
36	1.714	0.795	1.721	0.801	1.743	0.807
48	1.821	0.804	1.849	0.810	1.875	0.817
60	1.785	0.810	1.824	0.814	1.815	0.813
Avg	1.764	0.802	1.810	0.814	1.830	0.820
Setting	ILI→Power		Power→Power		Web→Power	
Metric	MSE	MAE	MSE	MAE	MSE	MAE
24	1.391	0.895	1.343	0.870	1.393	0.889
36	1.457	0.902	1.445	0.903	1.493	0.915
48	1.612	0.972	1.559	0.946	1.680	0.997
60	1.733	1.035	1.602	0.971	1.773	1.048
Avg	1.548	0.951	1.487	0.922	1.585	0.962
Setting	ILI→Web		Power→Web		Web→Web	
Metric	MSE	MAE	MSE	MAE	MSE	MAE
24	0.208	0.319	0.185	0.306	0.186	0.306
36	0.285	0.369	0.289	0.364	0.272	0.356
48	0.377	0.403	0.400	0.427	0.365	0.391
60	0.480	0.476	0.467	0.466	0.486	0.481
Avg	0.337	0.391	0.335	0.391	0.327	0.383

Table 41: Training and inference resource footprint for OLinear and baseline forecasters. ‘T.T.’, ‘T.M.’, ‘I.T.’ and ‘I.M.’ denote training time, training GPU memory usage, inference time, and inference GPU memory usage, respectively. All experiments use a lookback/prediction of 96 and a batch size of 16. To ensure a fair comparison, the number of layers (or blocks) is fixed at 2, except for DLinear, which employs only 1 linear layer. ‘OL’ refers to OLinear. ‘OOM’ indicates that the experiment runs out of memory on a 24 GB GPU.

	Model	OL (Ours)	OL-C (Ours)	Leddam 2024	CARD 2024b	Fredformer 2024	iTrans. 2024a	TimeMix.++ 2025a	TimeMix. 2024a	Dlinear 2023	TimesNet 2023b	PatchTST 2023	FreTS 2023
ETT	Params(M)	4.52	4.52	8.55	0.03	8.59	4.83	1.19	0.08	0.02	299.94	3.76	0.42
	FLOPs(M)	33.74	33.74	111.54	1.41	135.01	33.99	1396.90	10.72	0.13	289708	272.20	2.94
	T.T. (ms/iter)	8.09	7.26	25.02	29.10	26.48	10.11	74.74	27.99	1.42	501.98	8.37	3.63
	T.M.(GB)	0.20	0.2	0.18	0.03	0.24	0.21	0.27	0.05	0.02	5.80	0.14	0.03
	I.T.(ms/iter)	1.19	1.09	2.62	4.05	3.73	1.72	28.15	3.86	0.21	155.97	1.38	0.49
ECL	I.M. (MB)	156.41	156.41	51.04	11.02	78.04	156.00	149.58	18.06	8.48	1396.24	51.83	14.47
	Params(M)	4.79	4.58	8.56	1.39	12.12	4.83	1.19	0.12	0.02	300.58	3.76	0.42
	FLOPs(G)	1.65	1.65	5.32	5.07	5.55	1.87	64.02	1.74	0.01	293.98	12.48	0.13
	T.T. (ms/iter)	7.75	6.94	65.19	78.25	26.59	10.66	106.06	57.09	1.37	509.17	65.83	3.78
	T.M.(GB)	0.45	0.42	1.44	5.08	1.24	0.70	1.02	4.24	0.04	5.82	3.02	0.30
Exchange	I.T.(ms/iter)	2.11	2.09	14.78	28.21	7.08	3.46	34.10	22.79	0.20	157.13	20.52	0.48
	I.M. (GB)	0.17	0.17	0.49	0.63	0.26	0.23	0.36	0.82	0.02	1.41	0.93	0.23
	Params(M)	1.74	1.74	8.56	1.39	8.59	4.83	1.19	0.08	0.02	299.94	3.76	0.42
	FLOPs(M)	16.27	16.27	127.49	114.83	135.01	38.88	1595.72	12.25	0.15	287909	311.08	3.36
	T.T. (ms/iter)	7.33	7.53	24.81	35.57	28.56	10.14	97.05	27.46	1.47	505.39	9.75	3.51
Traffic	T.M.(GB)	0.17	0.17	0.18	0.17	0.24	0.21	0.28	0.06	0.02	5.80	0.15	0.03
	I.T.(ms/iter)	1.18	1.89	2.62	6.02	6.86	1.70	28.46	3.81	0.21	156.33	2.47	0.48
	I.M. (GB)	0.02	0.02	0.05	0.03	0.08	0.16	0.15	0.02	0.01	1.40	0.05	0.02
	Params(M)	6.17	4.69	8.56	0.98	11.09	4.83	4.73	0.12	0.02	301.69	3.76	0.42
	FLOPs(G)	4.91	4.91	15.24	10.47	14.78	6.45	588.47	4.66	0.02	288.72	33.52	0.36
Weather	T.T. (s/iter)	0.02	0.02	0.19	0.18	0.10	0.04	OOM	0.17	0.001	0.50	0.17	0.01
	T.M.(GB)	1.01	0.95	3.74	9.50	5.76	2.95	OOM	11.33	0.07	5.86	8.01	0.78
	I.T.(ms/iter)	5.71	5.67	45.68	61.21	31.60	13.45	1610.66	67.21	0.20	15.64	55.10	3.48
	I.M. (GB)	0.43	0.43	1.25	1.69	1.84	1.27	9.35	2.20	0.04	1.42	2.44	0.59
	Params(M)	4.52	4.52	8.56	0.03	0.50	4.83	2.37	0.10	0.02	299.97	3.76	0.42
Solar-Energy	FLOPs(G)	0.10	0.10	0.34	0.004	0.01	0.10	8.36	0.05	0.000	291.21	0.82	0.01
	T.T. (ms/iter)	7.47	6.90	27.55	28.69	37.85	9.45	102.04	33.30	1.38	505.88	9.89	3.53
	T.M.(GB)	0.21	0.21	0.21	0.05	0.05	0.22	0.86	0.16	0.02	5.80	0.27	0.04
	I.T.(ms/iter)	1.18	1.05	3.15	4.11	7.19	1.66	30.18	5.10	0.20	156.27	2.96	0.48
	I.M. (GB)	39.26	39.26	70.06	14.45	18.28	34.88	316.49	36.03	9.06	1397.18	90.78	24.37
PEMS07	Params(M)	4.58	4.54	8.56	1.39	4.61	4.83	2.37	0.12	0.02	300.21	3.76	0.42
	FLOPs(G)	0.68	0.68	2.22	2.05	1.20	0.72	46.81	0.74	0.003	290.64	5.33	0.06
	T.T. (ms/iter)	9.12	7.46	48.64	36.28	31.77	9.97	515.44	35.90	1.44	506.58	28.89	3.60
	T.M.(GB)	0.30	0.30	0.68	2.22	0.61	0.37	5.11	1.83	0.03	5.81	1.34	0.15
	I.T.(ms/iter)	1.20	1.08	7.65	10.44	5.16	1.96	189.20	8.54	0.20	156.66	8.49	0.48
PEMS07	I.M. (GB)	0.21	0.21	0.23	0.28	0.16	0.18	1.94	0.36	0.01	1.40	0.42	0.11
	Params(M)	4.75	3.19	8.56	1.39	13.44	4.83	2.38	0.12	0.02	301.73	3.76	0.42
	FLOPs(G)	3.67	3.67	15.65	16.10	20.53	6.66	301.92	4.77	0.02	290.97	34.34	0.37
	T.T. (s/iter)	0.01	0.01	0.19	0.28	0.11	0.04	OOM	0.17	0.002	0.51	0.19	0.01
	T.M.(GB)	0.83	0.76	3.85	13.99	6.42	3.07	OOM	11.62	0.07	5.86	9.18	0.79
PEMS07	I.T.(ms/iter)	1.21	1.10	47.06	94.01	36.54	13.92	1370.91	68.96	0.20	156.78	56.49	0.50
	I.M. (GB)	0.43	0.43	1.28	1.74	2.03	1.32	12.37	2.25	0.04	1.43	2.50	0.61

Table 42: Efficiency comparison of Leddam, iTransformer, PatchTST, and Fredformer before and after replacing multi-head self-attention with the NormLin module. In Leddam, only the *cross-channel attention* module is updated. Metrics include number of learnable parameters (Params), FLOPs, training time (T.T.), training memory (T.M.), inference time (I.T.), and inference memory (I.M.). All experiments use a lookback/prediction length of 96 and a batch size of 16. Notably, in its official implementation, Fredformer employs Nystromformer [Xiong et al., 2021]—an approximate yet more efficient self-attention mechanism—for the ETTm1, ECL, and Weather datasets.

Model	Leddam 2024		iTrans. 2024a		PatchTST 2023		Fredformer 2024	
	Attn.	NormLin	Attn.	NormLin	Attn.	NormLin	Attn.	NormLin
ETTm1	Params(M)	8.55	7.51	4.83	2.21	3.76	2.71	8.59
	FLOPs(M)	111.54	104.15	33.99	15.51	272.20	182.20	135.01
	T.T. (ms/iter)	25.02	23.67	10.11	7.10	8.37	6.82	26.48
	T.M.(GB)	0.18	0.16	0.21	0.17	0.14	0.12	0.24
	I.T.(ms/iter)	2.62	2.50	1.72	1.25	1.38	1.14	3.73
	I.M. (MB)	51.04	47.03	156.00	146.20	51.83	50.45	78.04
ECL	Params(M)	8.56	7.71	4.83	2.41	3.76	2.71	12.12
	FLOPs(G)	5.32	4.88	1.87	0.81	12.48	8.35	5.55
	T.T. (ms/iter)	65.19	64.93	10.66	7.39	65.83	50.82	26.59
	T.M.(GB)	1.44	1.41	0.70	0.41	3.02	2.78	1.24
	I.T.(ms/iter)	14.78	14.52	3.46	1.01	20.52	14.08	7.08
	I.M. (GB)	0.49	0.48	0.23	0.10	0.93	0.99	0.26
Weather	Params(M)	8.56	7.52	4.83	2.21	3.76	2.71	0.50
	FLOPs(G)	0.34	0.31	0.10	0.05	0.82	0.55	0.01
	T.T. (ms/iter)	27.55	27.21	9.45	7.49	9.89	7.47	37.85
	T.M.(GB)	0.21	0.19	0.22	0.18	0.27	0.24	0.05
	I.T.(ms/iter)	3.15	3.08	1.66	0.99	2.96	1.12	7.19
	I.M. (MB)	70.06	66.05	34.88	25.51	90.78	92.69	18.28
Solar-Energy	Params(M)	8.56	7.54	4.83	2.24	3.76	2.71	4.61
	FLOPs(G)	2.22	2.06	0.72	0.32	5.33	3.56	1.20
	T.T. (ms/iter)	48.64	48.12	9.97	7.40	28.89	22.31	31.77
	T.M.(GB)	0.68	0.67	0.37	0.26	1.34	1.23	0.61
	I.T.(ms/iter)	7.65	7.59	1.96	1.01	8.49	5.83	5.16
	I.M. (GB)	0.23	0.22	0.18	0.18	0.42	0.44	0.16
PEMS07	Params(M)	8.56	9.07	4.83	3.77	3.76	2.71	13.44
	FLOPs(G)	15.65	13.93	6.66	2.75	34.34	22.96	20.53
	T.T. (s/iter)	0.19	0.18	0.04	0.01	0.19	0.14	0.11
	T.M.(GB)	3.85	3.74	3.07	0.82	9.18	7.57	6.42
	I.T.(ms/iter)	47.06	45.16	13.92	3.02	56.49	38.63	36.54
	I.M. (GB)	1.28	1.28	1.32	0.27	2.50	2.70	2.03